

Project Summary

Directory Structure

```
measurement-free-quantum-classifier/
    configs/
        paths.yaml
    src/
        __init__.py
        IQL/
            __init__.py
            models/
                fixed_memory_iqc.py
                static_isdo_model.py
                adaptive_iqc.py
                __init__.py
            learning/
                class_state.py
                update.py
                metrics.py
                memory_bank.py
                calculate_prototype.py
                __init__.py
            backends/
                base.py
                prime_b.py
                hadamard.py
                transition.py
                exact.py
                __init__.py
            regimes/
                regime3c_adaptive.py
                regime3a_wta.py
                regime2_online.py
                __init__.py
        inference/
            weighted_vote_classifier.py
            __init__.py
        encoding/
            embedding_to_state.py
            __init__.py
        baselines/
            static_isdo_classifier.py
    utils/
        common_backup.py
        common.py
        paths.py
        seed.py
        label_utils.py
        __init__.py
```

```
data/
    pcam_loader.py
    transforms.py
    __init__.py
quantum/
    compute_qsvm_kernel.py
    __init__.py
training/
    test_adaptive_iqc.py
    verify_consistency.py
    run_final_comparison.py
    compare_best_iqc_vs_classical.py
    test_fixed_memory_iqc.py
    validate_backends.py
    test_static_isdo_model.py
    test_core_functionality.py
    compare_iqc_algorithms.py
protocol_online/
    train_perceptron.py
protocol_static/
    evaluate_isdo_k_sweep.py
    evaluate_static_isdo.py
classical/
    make_embedding_split.py
    train_embedding_models.py
    extract_embeddings.py
    visualize_embeddings.py
    train_cnn.py
    verify_embbeings.py
    visualize_pcam.py
protocol_adaptive/
    consolidate_memory.py
    train_adaptive_memory.py
classical/
    cnn.py
    __init__.py
experiments/
results/
    artifacts/
        regime3c_memory.pkl
    checkpoints/
        pcam_cnn_final.pt
        pcam_cnn_best.pt
    embeddings/
        val_labels.npy
        val_labels_polar.npy
        val_embeddings.npy
        split_test_idx.npy
        split_train_idx.npy
        class_states_meta.json
    class_prototypes/
        K7/
            class0_proto5.npy
            class1_proto6.npy
```

```
class1_proto0.npy  
class1_proto1.npy  
class1_proto2.npy  
class0_proto0.npy  
class0_proto3.npy  
class0_proto1.npy  
class0_proto4.npy  
class1_proto4.npy  
class1_proto3.npy  
class0_proto2.npy  
class0_proto6.npy  
class1_proto5.npy  
K17/  
class0_proto12.npy  
class0_proto5.npy  
class1_proto6.npy  
class1_proto10.npy  
class1_proto0.npy  
class1_proto7.npy  
class1_proto1.npy  
class1_proto9.npy  
class1_proto2.npy  
class1_proto14.npy  
class1_proto8.npy  
class1_proto11.npy  
class1_proto13.npy  
class0_proto0.npy  
class0_proto7.npy  
class0_proto3.npy  
class0_proto14.npy  
class0_proto1.npy  
class0_proto4.npy  
class0_proto9.npy  
class0_proto11.npy  
class0_proto16.npy  
class0_proto8.npy  
class0_proto15.npy  
class1_proto12.npy  
class0_proto10.npy  
class1_proto15.npy  
class1_proto16.npy  
class1_proto4.npy  
class1_proto3.npy  
class0_proto2.npy  
class0_proto6.npy  
class0_proto13.npy  
class1_proto5.npy  
K1/  
class1_proto0.npy  
class0_proto0.npy  
K13/  
class0_proto12.npy  
class0_proto5.npy  
class1_proto6.npy
```

```
class1_proto10.npy  
class1_proto0.npy  
class1_proto7.npy  
class1_proto1.npy  
class1_proto9.npy  
class1_proto2.npy  
class1_proto8.npy  
class1_proto11.npy  
class0_proto0.npy  
class0_proto7.npy  
class0_proto3.npy  
class0_proto1.npy  
class0_proto4.npy  
class0_proto9.npy  
class0_proto11.npy  
class0_proto8.npy  
class1_proto12.npy  
class0_proto10.npy  
class1_proto4.npy  
class1_proto3.npy  
class0_proto2.npy  
class0_proto6.npy  
class1_proto5.npy
```

K23/

```
class1_proto18.npy  
class0_proto12.npy  
class0_proto5.npy  
class1_proto6.npy  
class0_proto18.npy  
class1_proto10.npy  
class1_proto0.npy  
class0_proto19.npy  
class0_proto21.npy  
class1_proto7.npy  
class1_proto1.npy  
class1_proto9.npy  
class1_proto2.npy  
class1_proto14.npy  
class1_proto8.npy  
class1_proto11.npy  
class1_proto13.npy  
class0_proto0.npy  
class0_proto7.npy  
class0_proto17.npy  
class0_proto20.npy  
class1_proto22.npy  
class0_proto3.npy  
class0_proto14.npy  
class0_proto1.npy  
class0_proto4.npy  
class0_proto9.npy  
class0_proto11.npy  
class0_proto16.npy  
class0_proto8.npy
```

```
class0_proto15.npy  
class1_proto12.npy  
class1_proto21.npy  
class0_proto10.npy  
class0_proto22.npy  
class1_proto19.npy  
class1_proto15.npy  
class1_proto16.npy  
class1_proto4.npy  
class1_proto20.npy  
class1_proto3.npy  
class0_proto2.npy  
class0_proto6.npy  
class0_proto13.npy  
class1_proto17.npy  
class1_proto5.npy  
K5/  
class1_proto0.npy  
class1_proto1.npy  
class1_proto2.npy  
class0_proto0.npy  
class0_proto3.npy  
class0_proto1.npy  
class0_proto4.npy  
class1_proto4.npy  
class1_proto3.npy  
class0_proto2.npy  
K11/  
class0_proto5.npy  
class1_proto6.npy  
class1_proto10.npy  
class1_proto0.npy  
class1_proto7.npy  
class1_proto1.npy  
class1_proto9.npy  
class1_proto2.npy  
class1_proto8.npy  
class0_proto0.npy  
class0_proto7.npy  
class0_proto3.npy  
class0_proto1.npy  
class0_proto4.npy  
class0_proto9.npy  
class0_proto8.npy  
class0_proto10.npy  
class1_proto4.npy  
class1_proto3.npy  
class0_proto2.npy  
class0_proto6.npy  
class1_proto5.npy  
K19/  
class1_proto18.npy  
class0_proto12.npy  
class0_proto5.npy
```

```
class1_proto6.npy  
class0_proto18.npy  
class1_proto10.npy  
class1_proto0.npy  
class1_proto7.npy  
class1_proto1.npy  
class1_proto9.npy  
class1_proto2.npy  
class1_proto14.npy  
class1_proto8.npy  
class1_proto11.npy  
class1_proto13.npy  
class0_proto0.npy  
class0_proto7.npy  
class0_proto17.npy  
class0_proto3.npy  
class0_proto14.npy  
class0_proto1.npy  
class0_proto4.npy  
class0_proto9.npy  
class0_proto11.npy  
class0_proto16.npy  
class0_proto8.npy  
class0_proto15.npy  
class1_proto12.npy  
class0_proto10.npy  
class1_proto15.npy  
class1_proto16.npy  
class1_proto4.npy  
class1_proto3.npy  
class0_proto2.npy  
class0_proto6.npy  
class0_proto13.npy  
class1_proto17.npy  
class1_proto5.npy  
K2/  
    class1_proto0.npy  
    class1_proto1.npy  
    class0_proto0.npy  
    class0_proto1.npy  
K3/  
    class1_proto0.npy  
    class1_proto1.npy  
    class1_proto2.npy  
    class0_proto0.npy  
    class0_proto1.npy  
    class0_proto2.npy  
logs/  
    train_history.json  
    embedding_baseline_results.json  
figures/
```

File: configs/paths.yaml

```
base_root: "/home/tarakesh/Work/Repo/measurement-free-quantum-classifier"

paths:
    dataset: "dataset"
    checkpoints: "results/checkpoints"
    embeddings: "results/embeddings"
    figures: "results/figures"
    logs: "results/logs"
    class_prototypes: "results/embeddings/class_prototypes"
    artifacts: "results/artifacts"

class_count:
    K: 3
    K_values: [1, 2, 3, 5, 7, 11, 13, 17, 19, 23]
```

File: src/**init**.py

```
"""
```

File: src/IQL/**init**.py

```
"""
```

File: src/IQL/models/fixed_memory_iqc.py

```
# src/IQL/models/fixed_memory_iqc.py

import os
import numpy as np

from src.utils.paths import load_paths
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.regimes.regime3a_wta import WinnerTakeAll
from src.IQL.inference.weighted_vote_classifier import
WeightedVoteClassifier
from src.IQL.backends.exact import ExactBackend
from src.IQL.learning.calculate_prototype import generate_prototypes
from src.utils.label_utils import ensure_binary

class FixedMemoryIQC:
    """
```

Fixed-Memory Interference Quantum Classifier (IQC)

Training pipeline:

1. Generate K prototypes per class (if missing)
 2. Initialize Kx2 quantum memory states
 3. Train with Winner-Take-All (Regime-3A)
 4. Freeze memory
- """

```
def __init__(self, K: int, eta: float = 0.1, backend=None):  
    self.K = K  
    self.eta = eta  
    self.backend = backend or ExactBackend()  
  
    self.memory_bank = None  
    self.trainer = None  
    self.classifier = None  
  
def _ensure_prototypes(self, X, y):  
    """  
    Generate prototypes if they do not already exist.  
    """  
    _, PATHS = load_paths()  
    proto_base = PATHS["class_prototypes"]  
    proto_dir = os.path.join(proto_base, f"K{self.K}")  
  
    os.makedirs(proto_dir, exist_ok=True)  
    y_binary = ensure_binary(y)  
    generate_prototypes(  
        X=X,  
        y=y_binary,  
        K=self.K,  
        output_dir=proto_dir  
    )  
  
    def _load_prototypes(self):  
        _, PATHS = load_paths()  
        proto_dir = PATHS["class_prototypes"]  
  
        vectors = []  
        for cls in [0, 1]:  
            for i in range(self.K):  
                path = os.path.join(  
                    proto_dir, f"K{self.K}", f"class{cls}_proto{i}.npy"  
                )  
                vectors.append(np.load(path))  
        return vectors  
  
    def fit(self, X, y):  
        # -----  
        # Step 1: ensure prototypes exist  
        # -----  
        self._ensure_prototypes(X, y)
```

```

# -----
# Step 2: initialize memory bank
# -----
proto_vectors = self._load_prototypes()
class_states = [
    ClassState(v, backend=self.backend)
    for v in proto_vectors
]
self.memory_bank = MemoryBank(class_states)

# -----
# Step 3: Regime-3A training
# -----
self.trainer = WinnerTakeAll(
    memory_bank=self.memory_bank,
    eta=self.eta,
    backend=self.backend
)
self.trainer.fit(X, y)

# -----
# Step 4: freeze → inference
# -----
self.classifier = WeightedVoteClassifier(self.memory_bank)
return self

def predict(self, X):
    if self.classifier is None:
        raise RuntimeError("Model not trained. Call fit() first.")
    return [self.classifier.predict(x) for x in X]

```

File: src/IQL/models/static_isdo_model.py

```

# src/IQL/models/static_isdo_model.py

from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier
from src.utils.paths import load_paths
from src.IQL.learning.calculate_prototype import generate_prototypes
import os

class StaticISDOModel:
    """
    Static ISDO Model (Baseline)

    - K prototypes per class
    - No learning
    - Fixed interference reference state |chi>
    """
    def __init__(self, K: int):

```

```

    _, PATHS = load_paths()
    self.proto_dir = PATHS["class_prototypes"]
    self.K = K
    self.classifier = None

    def _ensure_prototypes(self, X, y):
        """
        Generate prototypes if they do not already exist.
        """
        _, PATHS = load_paths()
        proto_base = PATHS["class_prototypes"]
        proto_dir = os.path.join(proto_base, f"K{self.K}")
        os.makedirs(proto_dir, exist_ok=True)
        generate_prototypes(
            X=X,
            y=y,
            K=self.K,
            output_dir=proto_dir,
            seed = 42
        )

    def fit(self,X,y):
        """
        Offline preparation only.
        Loads precomputed prototypes and builds classifier.
        """
        self._ensure_prototypes(X,y)
        self.classifier = StaticISDOClassifier(
            proto_dir=self.proto_dir,
            K=self.K
        )
        return self

    def predict(self, X):
        if self.classifier is None:
            raise RuntimeError("Model not fitted. Call fit() first.")
        return self.classifier.predict(X)

```

File: src/IQL/models/adaptive_iqc.py

```

import numpy as np

from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.regimes.regime3c_adaptive import AdaptiveMemory
from src.IQL.regimes.regime3a_wta import WinnerTakeAll
from src.IQL.inference.weighted_vote_classifier import
WeightedVoteClassifier
from src.IQL.backends.exact import ExactBackend
from src.IQL.learning.calculate_prototype import generate_prototypes

```

```
from src.utils.paths import load_paths
from src.utils.label_utils import ensure_polar, ensure_binary

import os

class AdaptiveIQC:
    """
    Final Adaptive Interference Quantum Classifier (IQC)

    Pipeline:
    1. Prototype generation (offline, classical)
    2. Adaptive growth (Regime-3C)
    3. Consolidation (Regime-3A)
    4. Inference-only classifier
    """

    def __init__(
        self,
        K_init=3,
        eta=0.1,
        percentile=5,
        backend=None,
        consolidate=True,
    ):
        self.K_init = K_init
        self.eta = eta
        self.percentile = percentile
        self.backend = backend or ExactBackend()
        self.consolidate = consolidate

        self.memory_bank = None
        self.regime3c = None
        self.classifier = None

    def _initialize_memory(self, X, y):
        _, PATHS = load_paths()
        proto_dir = os.path.join(PATHS["class_prototypes"],
f"K{self.K_init}")

        y_binary = ensure_binary(y)
        generate_prototypes(
            X=X,
            y=y_binary,
            K=self.K_init,
            output_dir=proto_dir,
            seed =42,
        )

        class_states = []
        for cls in [0, 1]:
            for i in range(self.K_init):
                vec = np.load(
                    os.path.join(proto_dir, f"class{cls}_proto{i}.npy"))
                /
```

```
)  
    class_states.append(  
        ClassState(vec, backend=self.backend)  
)  
  
self.memory_bank = MemoryBank(class_states)  
  
def fit(self, X, y):  
    # -----  
    # Step 1 – initialize memory  
    # -----  
    y = ensure_polar(y)  
    self._initialize_memory(X, y)  
  
    # -----  
    # Step 2 – adaptive growth + pruning (Regime-3C)  
    # -----  
    self.regime3c = AdaptiveMemory(  
        memory_bank=self.memory_bank,  
        eta=self.eta,  
        percentile=self.percentile,  
        backend=self.backend,  
)  
    self.regime3c.fit(X, y)  
  
    # -----  
    # Step 3 – optional consolidation (Regime-3A)  
    # -----  
    if self.consolidate:  
        consolidator = WinnerTakeAll(  
            memory_bank=self.memory_bank,  
            eta=self.eta,  
            backend=self.backend,  
)  
        consolidator.fit(X, y)  
  
    # -----  
    # Step 4 – freeze & inference  
    # -----  
    self.classifier = WeightedVoteClassifier(self.memory_bank)  
return self  
  
def predict(self, X):  
    if self.classifier is None:  
        raise RuntimeError("Model not trained.")  
    return [self.classifier.predict(x) for x in X]
```

File: src/IQL/models/**init**.py

File: src/IQL/learning/class_state.py

```
import numpy as np
from src.IQL.backends.base import InterferenceBackend

def normalize(v: np.ndarray) -> np.ndarray:
    norm = np.linalg.norm(v)
    if norm == 0:
        raise ValueError("Zero-norm vector cannot be normalized")
    return v / norm

class ClassState:
    """
    Represents the quantum class memory |chi>.
    Invariant: ||chi|| = 1 always.
    """

    def __init__(self, vector: np.ndarray, backend: InterferenceBackend):
        self.vector = normalize(vector.astype(np.complex128))
        self.backend = backend

    def score(self, psi: np.ndarray) -> float:
        """
        ISDO score: Re <chi | psi>
        """
        return self.backend.score(self.vector, psi)

    def update(self, delta: np.ndarray):
        """
        Update |chi> <- normalize(|chi> + delta)
        """
        self.vector = normalize(self.vector + delta)
```

File: src/IQL/learning/update.py

```
import numpy as np
from src.IQL.backends.base import InterferenceBackend

def update(
    chi: np.ndarray,
    psi: np.ndarray,
    y: int,
    eta: float,
    backend: InterferenceBackend,
```

```

):

"""
Regime-2 update rule (quantum perceptron):

If  $y * \text{Re}[\chi|\psi] \geq 0$ :
    no update
else:
     $\chi \leftarrow \text{normalize}(\chi + \eta * y * \psi)$ 
"""

s = backend.score(chi, psi)

if  $y * s \geq 0$ :
    return chi, False # correct classification

delta = eta * y * psi
chi_new = chi + delta
chi_new = chi_new / np.linalg.norm(chi_new)

return chi_new, True

```

File: src/IQL/learning/metrics.py

```

import numpy as np

def summarize_training(history: dict):
    margins = np.array(history["margins"])
    updates = np.array(history["updates"])

    return {
        "mean_margin": float(margins.mean()),
        "min_margin": float(margins.min()),
        "num_updates": int(updates.sum()),
        "update_rate": float(updates.mean()),
    }

```

File: src/IQL/learning/memory_bank.py

```

from src.IQL.learning.class_state import ClassState

class MemoryBank:
    def __init__(self, class_states):
        self.class_states = class_states

    def scores(self, psi):
        return [
            cs.score(psi)
            for cs in self.class_states
        ]

```

```

    ]

def winner(self, psi):
    scores = self.scores(psi)
    idx = int(max(range(len(scores)), key=lambda i: abs(scores[i])))
    #idx = int(max(range(len(scores)), key=lambda i: scores[i])) ## causes lower score ??
    return idx, scores[idx]

def add_memory(self, chi_vector, backend):
    self.class_states.append(ClassState(chi_vector, backend=backend))

```

File: src/IQL/learning/calculate_prototype.py

```

import os
import numpy as np
from sklearn.cluster import KMeans

from src.utils.seed import set_seed


# -----
# Helper: quantum-safe normalization
# -----
def to_quantum_state(x):
    x = np.asarray(x, dtype=np.float64).reshape(-1)
    x = x / np.sqrt(np.sum(x ** 2))
    assert np.isclose(np.sum(x ** 2), 1.0, atol=1e-12)
    return x


# -----
# Core function (IMPORTABLE)
# -----
def generate_prototypes(X, y, K, output_dir, seed=42):
    """
    Generate K prototypes per class using KMeans clustering.

    Args:
        X (np.ndarray): embeddings, shape (N, D)
        y (np.ndarray): labels in {0,1}
        K (int): number of prototypes per class
        output_dir (str): directory to save prototypes
        seed (int): random seed
    """
    set_seed(seed)
    os.makedirs(output_dir, exist_ok=True)

    for cls in [0, 1]:
        X_cls = X[y == cls].astype(np.float64)

```

```
if len(X_cls) < K:
    raise ValueError(
        f"Not enough samples for class {cls}: "
        f"{len(X_cls)} < K={K}"
    )

kmeans = KMeans(
    n_clusters=K,
    random_state=seed,
    n_init=10
)
kmeans.fit(X_cls)

centers = kmeans.cluster_centers_

for i in range(K):
    proto = to_quantum_state(centers[i])
    path = os.path.join(output_dir, f"class{cls}_proto{i}.npy")
    np.save(path, proto)

# -----
# Script mode (EXPERIMENTS ONLY)
# -----
if __name__ == "__main__":
    from src.utils.paths import load_paths

    # Reproducibility
    set_seed(42)

    # Load paths
    _, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]
    PROTO_BASE = PATHS["class_prototypes"]

    os.makedirs(EMBED_DIR, exist_ok=True)
    os.makedirs(PROTO_BASE, exist_ok=True)

    # Load embeddings (TRAIN ONLY)
    X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
    y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))
    train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))

    X_train = X[train_idx]
    y_train = y[train_idx]

    print("Loaded train embeddings:", X_train.shape)

    K_VALUES = PATHS["class_count"]["K_values"]

    for K in K_VALUES:
        print(f"\n==== Computing prototypes for K={K} ====")
        CLASS_DIR = os.path.join(PROTO_BASE, f"K{K}")
        /
```

```
    generate_prototypes(
        X=X_train,
        y=y_train,
        K=K,
        output_dir=CLASS_DIR,
        seed=42
    )
    print(f"Saved prototypes to {CLASS_DIR}")
```

File: src/IQL/learning/**init**.py

File: src/IQL/backends/base.py

```
from abc import ABC, abstractmethod

class InterferenceBackend(ABC):
    """
    Abstract interface for computing interference scores.
    """

    @abstractmethod
    def score(self, chi, psi) -> float:
        """
        Return Re<chi | psi> as a real scalar.
        """
        pass
```

File: src/IQL/backends/prime_b.py

```
import numpy as np
from qiskit import QuantumCircuit
from qiskit.quantum_info import Statevector, Pauli
from qiskit.circuit.library import StatePreparation

from .base import InterferenceBackend

class PrimeBBackend(InterferenceBackend):
    """
    PrimeB (ISDO-B') Backend
    -----
```

Observable-engineered, decision-sufficient implementation of ISDO.

Computes:

$$S(\psi; X) = \langle \psi | U_X^\dagger Z^{\{\otimes n\}} U_X | \psi \rangle$$

Properties:

- No ancilla qubit
- No controlled unitaries
- X appears only as a basis rotation
- Fixed, hardware-native observable
- Preserves sign + ordering (not exact inner product)

Intended role:

- Fast inference
- NISQ-friendly deployment backend

"""

@staticmethod

def _statevector_to_unitary(state: np.ndarray) -> np.ndarray:

"""

Construct a unitary U such that:

$$U |0\dots0\rangle = |\text{state}\rangle$$

Uses Gram-Schmidt completion.

"""

state = np.asarray(state, dtype=np.complex128)
state = state / np.linalg.norm(state)

dim = len(state)

U = np.zeros((dim, dim), dtype=np.complex128)
U[:, 0] = state

for i in range(1, dim):

v = np.zeros(dim, dtype=np.complex128)
v[i] = 1.0

for j in range(i):

v -= np.vdot(U[:, j], v) * U[:, j]

norm = np.linalg.norm(v)

if norm < 1e-12:

v = np.random.randn(dim) + 1j * np.random.randn(dim)
for j in range(i):

v -= np.vdot(U[:, j], v) * U[:, j]
v /= np.linalg.norm(v)

else:

v /= norm

U[:, i] = v

return U

def score(self, chi: np.ndarray, psi: np.ndarray) -> float:
"""

```

Compute PrimeB interference score.

Args:
    chi : np.ndarray
        Class memory state  $|x\rangle$ 
    psi : np.ndarray
        Input state  $|\psi\rangle$ 

Returns:
    float
        Decision-sufficient interference score
"""

chi = np.asarray(chi, dtype=np.complex128)
psi = np.asarray(psi, dtype=np.complex128)

chi /= np.linalg.norm(chi)
psi /= np.linalg.norm(psi)

dim = len(psi)
n = int(np.log2(dim))
if 2 ** n != dim:
    raise ValueError("State dimension must be a power of 2")

# Build circuit
qc = QuantumCircuit(n)

# Prepare  $|\psi\rangle$ 
qc.append(StatePreparation(psi), range(n))

# Apply  $U_X$ 
U_chi = self._statevector_to_unitary(chi)
qc.unitary(U_chi, range(n), label="U_chi")

# Evaluate  $\langle Z^{\otimes n} \rangle$ 
sv = Statevector.from_instruction(qc)
observable = Pauli("Z"+"I" * (n-1))

    return float(sv.expectation_value(observable).real)
"""
pass

```

File: src/IQL/backends/hadamard.py

```

import numpy as np
from qiskit import QuantumCircuit
from qiskit.quantum_info import Statevector, Pauli
from qiskit.circuit.library import StatePreparation # ✓ Correct import
from .base import InterferenceBackend

# If you also want the conceptual/oracle version:
class HadamardBackend(InterferenceBackend):

```

```
"""
CONCEPTUAL Hadamard-test using oracle state preparation.

WARNING: This uses non-unitary StatePreparation and is NOT
physically realizable. Use only for conceptual understanding.
For actual implementation, use TransitionInterferenceBackend.

Computes Re<chi | psi> in oracle model.
"""

def score(self, chi, psi) -> float:
    chi = np.asarray(chi, dtype=np.complex128)
    psi = np.asarray(psi, dtype=np.complex128)

    # Normalize
    chi = chi / np.linalg.norm(chi)
    psi = psi / np.linalg.norm(psi)

    assert chi.shape == psi.shape
    n = int(np.log2(len(psi)))
    assert 2**n == len(psi)

    qc = QuantumCircuit(1 + n)
    anc = 0
    data = list(range(1, 1 + n))

    # Hadamard on ancilla
    qc.h(anc)

    # Controlled state preparation (ORACLE ASSUMPTION)
    # When anc=0: prepare |psi>
    state_prep_psi = StatePreparation(psi)
    qc.append(state_prep_psi.control(1), [anc] + data)

    # Flip ancilla
    qc.x(anc)

    # When anc=1 (after flip, so anc=0): prepare |chi>
    state_prep_chi = StatePreparation(chi)
    qc.append(state_prep_chi.control(1), [anc] + data)

    # Flip back
    qc.x(anc)

    # Final Hadamard
    qc.h(anc)

    # Get statevector and measure Z on ancilla
    sv = Statevector.from_instruction(qc)
    z_exp = sv.expectation_value(Pauli('Z')), [anc]).real

    return float(z_exp)
```

File: src/IQL/backends/transition.py

```
import numpy as np
from qiskit import QuantumCircuit
from qiskit.quantum_info import Statevector, Pauli
from qiskit.circuit.library import UnitaryGate, StatePreparation # ✓
Correct import
from .base import InterferenceBackend

class TransitionBackend(InterferenceBackend):
    """
    CORRECT physical Hadamard-test using transition unitary.

    This is the physically realizable ISDO implementation.
    Computes Re<chi | psi> using U_chi_psi = U_chi @ U_psi^dagger

    This should be used for all hardware experiments and claims.
    """

    @staticmethod
    def _statevector_to_unitary(vec):
        """Build unitary that prepares vec from |0...0>"""
        vec = np.asarray(vec, dtype=np.complex128)
        vec = vec / np.linalg.norm(vec)
        dim = len(vec)

        U = np.zeros((dim, dim), dtype=complex)
        U[:, 0] = vec

        # Gram-Schmidt to complete the unitary
        for i in range(1, dim):
            v = np.zeros(dim, dtype=complex)
            v[i] = 1.0

            for j in range(i):
                v -= np.vdot(U[:, j], v) * U[:, j]

            v_norm = np.linalg.norm(v)
            if v_norm > 1e-10:
                U[:, i] = v / v_norm
            else:
                v = np.random.randn(dim) + 1j * np.random.randn(dim)
                for j in range(i):
                    v -= np.vdot(U[:, j], v) * U[:, j]
                U[:, i] = v / np.linalg.norm(v)

        return U

    @staticmethod
    def _build_transition_unitary(psi, chi):
        """Build U_chi_psi = U_chi @ U_psi^dagger"""

    
```

```

U_psi = TransitionBackend._statevector_to_unitary(psi)
U_chi = TransitionBackend._statevector_to_unitary(chi)

# Transition unitary
U_chi_psi = U_chi @ U_psi.conj().T

return UnitaryGate(U_chi_psi)

def score(self, chi, psi) -> float:
    chi = np.asarray(chi, dtype=np.complex128)
    psi = np.asarray(psi, dtype=np.complex128)

    # Normalize
    chi = chi / np.linalg.norm(chi)
    psi = psi / np.linalg.norm(psi)

    assert chi.shape == psi.shape
    n = int(np.log2(len(psi)))
    assert 2**n == len(psi)

    qc = QuantumCircuit(1 + n)
    anc = 0
    data = list(range(1, 1 + n))

    # Prepare |psi> on data qubits
    qc.append(StatePreparation(psi), data)

    # Hadamard on ancilla
    qc.h(anc)

    # Controlled transition unitary
    U_chi_psi = self._build_transition_unitary(psi, chi)
    qc.append(U_chi_psi.control(1), [anc] + data)

    # Final Hadamard
    qc.h(anc)

    # Get statevector and measure Z on ancilla
    sv = Statevector.from_instruction(qc)
    z_exp = sv.expectation_value(Pauli('Z'), [anc]).real

return float(z_exp)

```

File: src/IQL/backends/exact.py

```

import numpy as np
from .base import InterferenceBackend

class ExactBackend(InterferenceBackend):
    """
    Numpy-based interference backend.

```

```
This reproduces existing behavior exactly.  
"""  
  
def score(self, chi, psi) -> float:  
    return float(np.real(np.vdot(chi, psi)))
```

File: src/IQL/backends/**init**.py

File: src/IQL/regimes/regime3c_adaptive.py

```
import numpy as np  
from collections import deque  
from src.IQL.learning.update import update  
from src.IQL.backends.exact import ExactBackend  
import pickle  
  
class AdaptiveMemory:  
    """  
        Regime 3-C: Dynamic Memory Growth with Percentile-based τ  
    """  
  
    def __init__(  
        self,  
        memory_bank,  
        eta=0.1,  
        percentile=5,  
        tau_abs = -0.4,  
        margin_window=500,  
        backend=ExactBackend()  
    ):  
        self.memory_bank = memory_bank  
        self.eta = eta  
        self.percentile = percentile  
        self.tau_abs = tau_abs  
        self.backend = backend  
  
        # store recent margins  
        self.margins = deque(maxlen=margin_window)  
  
        self.num_updates = 0  
        self.num_spawns = 0  
  
        self.history = {  
            "margin": [],  
            "spawned": [],  
            "num_memories": []},
```

```
}

def aggregated_score(self, psi):
    scores = self.memory_bank.scores(psi)
    return sum(scores) / len(scores)

def step(self, psi, y):
    S = self.aggregated_score(psi)
    margin = y * S

    # collect negative margins only
    neg_margins = [m for m in self.margins if m < 0]

    spawned = False

    # compute percentile only if we have enough negative history
    if len(neg_margins) >= 20:
        tau = np.percentile(neg_margins, self.percentile)

        if margin < tau:
            # 🔥 spawn new memory
            chi_new = y * psi
            chi_new = chi_new / np.linalg.norm(chi_new)
            self.memory_bank.add_memory(chi_new, self.backend)
            self.num_spawns += 1
            spawned = True

    # otherwise, normal Regime-2 update on winner
    if not spawned and margin < 0:
        idx, _ = self.memory_bank.winner(psi)
        cs = self.memory_bank.class_states[idx]

        chi_new, updated = update(
            cs.vector, psi, y, self.eta, self.backend
        )

        if updated:
            cs.vector = chi_new
            self.num_updates += 1

    # logging
    self.margins.append(margin)
    self.history["margin"].append(margin)
    self.history["spawned"].append(spawned)

    self.history["num_memories"].append(len(self.memory_bank.class_states))

    return margin, spawned

def memory_size(self):
    return len(self.memory_bank.class_states)

def fit(self, X, y):
    for psi, y in zip(X, y):
```

```
        self.step(psi, y)

    def predict_one(self, X):
        _, score = self.memory_bank.winner(X)
        return 1 if score >= 0 else -1

    def predict(self, X):
        return [self.predict_one(x) for x in X]

    def save(self, path):
        """
        Save trained memory + training history.
        """
        payload = {
            "memory_bank": self.memory_bank,
            "eta": self.eta,
            "percentile": self.percentile,
            "tau_abs": self.tau_abs,
            "margins": list(self.margins),
            "num_updates": self.num_updates,
            "num_spawns": self.num_spawns,
            "history": self.history,
            "backend": self.backend,
        }

        with open(path, "wb") as f:
            pickle.dump(payload, f)

    @classmethod
    def load(cls, path):
        """
        Load a previously trained Regime-3C model.
        """
        with open(path, "rb") as f:
            payload = pickle.load(f)

        obj = cls(
            memory_bank=payload["memory_bank"],
            eta=payload["eta"],
            percentile=payload["percentile"],
            tau_abs=payload["tau_abs"],
            margin_window=len(payload["margins"]),
            backend=payload["backend"],
        )

        # restore training state
        from collections import deque
        obj.margins = deque(payload["margins"],
maxlen=len(payload["margins"]))
        obj.num_updates = payload["num_updates"]
        obj.num_spawns = payload["num_spawns"]
        obj.history = payload["history"]

        return obj
```

File: src/IQL/regimes/regime3a_wta.py

```
from src.IQL.learning.update import update
from src.IQL.backends.exact import ExactBackend
import pickle

class WinnerTakeAll:
    """
    Regime 3-A: Winner-Takes-All IQC
    Only the winning memory is updated.
    """

    def __init__(self, memory_bank, eta, backend = ExactBackend()):
        self.memory_bank = memory_bank
        self.eta = eta
        self.backend = backend
        self.num_updates = 0

        self.history = {
            "winner_idx": [],
            "scores": [],
            "updates": [],
        }

    def step(self, psi, y):
        idx, score = self.memory_bank.winner(psi)
        cs = self.memory_bank.class_states[idx]

        chi_new, updated = update(
            cs.vector, psi, y, self.eta, self.backend
        )

        if updated:
            cs.vector = chi_new
            self.num_updates += 1

        y_hat = 1 if score >= 0 else -1

        # logging
        self.history["winner_idx"].append(idx)
        self.history["scores"].append(score)
        self.history["updates"].append(updated)

        return y_hat, idx, updated

    def fit(self, X, y):
        correct = 0
        for x, y in zip(X, y):
            y_hat, _, _ = self.step(x, y)
            if y_hat == y:
                correct += 1
```

```

        correct += 1
    return correct / len(X)

def predict_one(self, X):
    _, score = self.memory_bank.winner(X)
    return 1 if score >= 0 else -1

def predict(self, X):
    return [self.predict_one(x) for x in X]

def save(self, path):
    """
    Save trained memory bank and history.
    """
    payload = {
        "memory_bank": self.memory_bank,
        "eta": self.eta,
        "num_updates": self.num_updates,
        "history": self.history,
        "backend": self.backend,
    }

    with open(path, "wb") as f:
        pickle.dump(payload, f)

@classmethod
def load(cls, path):
    """
    Load a trained Winner-Take-All model.
    """
    with open(path, "rb") as f:
        payload = pickle.load(f)

    obj = cls(
        memory_bank=payload["memory_bank"],
        eta=payload["eta"],
        backend=payload["backend"],
    )

    # restore training statistics
    obj.num_updates = payload["num_updates"]
    obj.history = payload["history"]

    return obj

```

File: src/IQL/regimes/regime2_online.py

```

import numpy as np
from src.IQL.learning.update import update
import pickle

```

```
class OnlinePerceptron:  
    """  
        Online Interference Quantum Classifier (Regime 2)  
  
        Fixed circuit.  
        Trainable object: |chi>  
    """  
  
    def __init__(self, class_state, eta: float):  
        self.class_state = class_state  
        self.eta = eta  
        # logs  
        self.num_updates = 0  
        self.history = {  
            "scores": [],  
            "margins": [],  
            "updates": [],  
        }  
  
    def step(self, psi: np.ndarray, y: int):  
        """  
            Process a single training example.  
        """  
        s = self.class_state.score(psi)  
        margin = y * s  
        y_hat = 1 if s >= 0 else -1  
  
        chi_new, updated = update(  
            self.class_state.vector, psi, y, self.eta,  
            self.class_state.backend  
        )  
  
        if updated:  
            self.class_state.vector = chi_new  
            self.num_updates += 1  
  
        # logging  
        self.history["scores"].append(s)  
        self.history["margins"].append(margin)  
        self.history["updates"].append(updated)  
  
        return y_hat, s, updated  
  
    def fit(self, X, y):  
        """  
            Single-pass online training.  
            dataset: iterable of (psi, y)  
        """  
        correct = 0  
  
        for i in range(len(X)):  
            y_hat, _, _ = self.step(X[i], y[i])  
            if y_hat == y[i]:  
                /
```

```
        correct += 1

    accuracy = correct / len(X)
    return accuracy

def predict_one(self, X):
    s = self.class_state.score(X)
    return 1 if s >= 0 else -1

def predict(self, X):
    return [self.predict_one(x) for x in X]

def save(self, path):
    """
    Save trained perceptron state and history.
    """
    payload = {
        "class_state": self.class_state,
        "eta": self.eta,
        "num_updates": self.num_updates,
        "history": self.history,
        "backend": self.class_state.backend,
    }

    with open(path, "wb") as f:
        pickle.dump(payload, f)

@classmethod
def load(cls, path):
    """
    Load a trained perceptron model.
    """
    with open(path, "rb") as f:
        payload = pickle.load(f)

    obj = cls(
        class_state=payload["class_state"],
        eta=payload["eta"],
    )

    # restore training statistics
    obj.num_updates = payload["num_updates"]
    obj.num_mistakes = payload["num_mistakes"]
    obj.margin_history = payload["margin_history"]
    obj.history = payload["history"]

    return obj
```

File: src/IQL/regimes/**init**.py

File: src/IQL/inference/weighted_vote_classifier.py

```
class WeightedVoteClassifier:  
    def __init__(self, memory_bank, weights=None):  
        self.memory_bank = memory_bank  
        self.M = len(memory_bank.class_states)  
  
        if weights is None:  
            self.weights = [1.0 / self.M] * self.M  
        else:  
            s = sum(weights)  
            self.weights = [w / s for w in weights]  
  
    def score(self, psi):  
        scores = self.memory_bank.scores(psi)  
        return sum(w * s for w, s in zip(self.weights, scores))  
  
    def predict(self, psi):  
        return 1 if self.score(psi) >= 0 else -1  
  
    def save(self, path):  
        import pickle  
        payload = {  
            "memory_bank": self.memory_bank,  
            "weights": self.weights,  
        }  
        with open(path, "wb") as f:  
            pickle.dump(payload, f)  
  
    @classmethod  
    def load(cls, path):  
        import pickle  
        with open(path, "rb") as f:  
            payload = pickle.load(f)  
        obj = cls(payload["memory_bank"], payload["weights"])  
        return obj
```

File: src/IQL/inference/init.py

File: src/IQL/encoding/embedding_to_state.py

```
import numpy as np
```

```
def embedding_to_state(x: np.ndarray) -> np.ndarray:
    """
    Maps a real embedding  $x \in \mathbb{R}^d$  to a quantum state  $|\psi\rangle$ .
    This is a purely geometric normalization.
    """
    x = x.astype(np.complex128)
    norm = np.linalg.norm(x)
    if norm == 0:
        raise ValueError("Zero embedding encountered")
    return x / norm
```

File: src/IQL/encoding/**init**.py

File: src/IQL/baselines/static_isdo_classifier.py

```
import os
import numpy as np
from tqdm import tqdm
from src.IQL.backends.exact import ExactBackend

class StaticISDOClassifier:
    def __init__(self, proto_dir, K):
        self.proto_dir = proto_dir
        self.K = K
        self.exact = ExactBackend()
        self.prototypes = {
            0: [np.load(os.path.join(proto_dir,
f"K{K}/class0_proto{i}.npy")) for i in range(K)],
            1: [np.load(os.path.join(proto_dir,
f"K{K}/class1_proto{i}.npy")) for i in range(K)],
        }

    def predict_one(self, psi):
        #A0 = sum(np.vdot(p, psi) for p in self.prototypes[0])
        #A1 = sum(np.vdot(p, psi) for p in self.prototypes[1])
        #return 1 if np.real(A0 - A1) < 0 else 0
        chi = sum(self.prototypes[0]) - sum(self.prototypes[1])
        chi /= np.linalg.norm(chi)
        return 1 if self.exact.score(chi, psi) < 0 else 0

    def predict(self, X):
        return np.array([self.predict_one(x) for x in tqdm(X, desc="ISDO Prediction", leave=False)])
```

File: src/utils/common_backup.py

```
import numpy as np
from qiskit import QuantumCircuit
from qiskit.circuit.library import StatePreparation, UnitaryGate

def load_statevector(vec):
    """
    Create a Qiskit StatePreparation gate from a normalized vector.

    NOTE: This is for CONCEPTUAL/ORACLE model only (Circuit A)
    For physical implementation, use build_transition_unitary instead
    """
    vec = np.asarray(vec, dtype=np.complex128)
    norm = np.linalg.norm(vec)
    if not np.isclose(norm, 1.0, atol=1e-12):
        raise ValueError("Statevector must be normalized")
    return StatePreparation(vec)

def statevector_to_unitary(psi):
    """
    Convert a statevector to a unitary operator that creates it from
    |0...0>
    Uses Gram-Schmidt to complete the unitary matrix.

    This creates U_psi such that U_psi |0...0> = |psi>

    Used for building transition unitaries in Circuit B'.
    """
    psi = np.asarray(psi, dtype=np.complex128)
    dim = len(psi)

    # Normalize
    psi = psi / np.linalg.norm(psi)

    # Create unitary matrix where first column is psi
    U = np.zeros((dim, dim), dtype=complex)
    U[:, 0] = psi

    # Complete to full unitary using Gram-Schmidt orthogonalization
    for i in range(1, dim):
        # Start with standard basis vector
        v = np.zeros(dim, dtype=complex)
        v[i] = 1.0

        # Orthogonalize against all previous columns
        for j in range(i):
            v -= np.vdot(U[:, j], v) * U[:, j]

        # Normalize and store
        v /= np.linalg.norm(v)
        U[:, i] = v
```

```
v_norm = np.linalg.norm(v)
if v_norm > 1e-10:
    U[:, i] = v / v_norm
else:
    # Use random vector if degenerate
    v = np.random.randn(dim) + 1j * np.random.randn(dim)
    for j in range(i):
        v -= np.vdot(U[:, j], v) * U[:, j]
    U[:, i] = v / np.linalg.norm(v)

return U

def build_transition_unitary(psi, chi):
    """
    Build the transition unitary  $U_{\chi\psi} = U_\chi @ U_\psi^\dagger$ 

    This is the KEY OPERATION for physically realizable ISDO (Circuit B').

    This unitary satisfies:  $U_{\chi\psi} |\psi\rangle = |\chi\rangle$ 

    Args:
        psi: Source statevector
        chi: Target statevector

    Returns:
        UnitaryGate that implements the transition
    """
    # Build unitaries that prepare each state from |0...0>
    U_psi = statevector_to_unitary(psi)
    U_chi = statevector_to_unitary(chi)

    # Transition unitary:  $U_{\chi\psi} = U_\chi @ U_\psi^\dagger$ 
    U_chi_psi = U_chi @ U_psi.conj().T

    # Verify it works
    psi_normalized = np.asarray(psi, dtype=np.complex128)
    psi_normalized = psi_normalized / np.linalg.norm(psi_normalized)
    chi_normalized = np.asarray(chi, dtype=np.complex128)
    chi_normalized = chi_normalized / np.linalg.norm(chi_normalized)

    result = U_chi_psi @ psi_normalized
    if not np.allclose(result, chi_normalized, atol=1e-10):
        raise ValueError("Transition unitary does not correctly map |\psi> to |\chi>")

    return UnitaryGate(U_chi_psi)

def build_chi_state(class0_protos, class1_protos):
    """
    Build  $|\chi\rangle = \sum_k |\phi_k^0\rangle - \sum_k |\phi_k^1\rangle$ , normalized

    This constructs the reference state for ISDO classification.
    """
    pass
```

```

"""
chi = np.zeros_like(class0_protos[0], dtype=np.float64)

for p in class0_protos:
    chi += p
for p in class1_protos:
    chi -= p

chi /= np.linalg.norm(chi)
return chi

```

File: src/utils/common.py

```

import numpy as np
from qiskit import QuantumCircuit
from qiskit.circuit.library import StatePreparation, UnitaryGate


def load_statevector(vec):
    """
    Create a Qiskit StatePreparation gate from a normalized vector.

    NOTE: This is for CONCEPTUAL/ORACLE model only (Circuit A)
    For physical implementation, use build_transition_unitary instead
    """
    vec = np.asarray(vec, dtype=np.complex128)
    norm = np.linalg.norm(vec)
    if not np.isclose(norm, 1.0, atol=1e-12):
        raise ValueError("Statevector must be normalized")
    return StatePreparation(vec)


def statevector_to_unitary(psi):
    """
    Convert a statevector to a unitary operator using Householder
    efficiency.

    Construct a Householder reflection U such that U |e1> = |\psi>
    where e1 = [1, 0, ..., 0]^T.

    This is O(D^2) to build the matrix, compared to O(D^3) for Gram-
    Schmidt.
    """
    psi = np.asarray(psi, dtype=np.complex128)
    norm = np.linalg.norm(psi)
    if norm > 1e-15:
        psi = psi / norm

    dim = len(psi)
    e1 = np.zeros(dim, dtype=np.complex128)
    e1[0] = 1.0

```

```

# Adjust phase to avoid numerical instability (choose phase to make w
large)
# We want to map phase * e1 to psi where phase has same angle as psi[0]
# This ensures w = phase * e1 - psi is stable.
angle = np.angle(psi[0]) if np.abs(psi[0]) > 1e-10 else 0.0
phase = np.exp(1j * angle)

target = phase * e1
w = target - psi
w_norm = np.linalg.norm(w)

if w_norm < 1e-12:
    # psi is already phase * e1, so just return identity * phase
    return np.eye(dim, dtype=np.complex128) * phase

v = w / w_norm
# R = I - 2vv* maps target (phase * e1) to psi
# R * phase * e1 = psi => R * e1 = psi * phase*
# To get U * e1 = psi, we need U = R * phase
H = (np.eye(dim, dtype=np.complex128) - 2.0 * np.outer(v, v.conj())) *
phase
return H

```

```

def build_transition_unitary(psi, chi):
    """
    Build the transition unitary U_chi_psi = U_chi @ U_psi^dagger

    This is the KEY OPERATION for physically realizable ISDO (Circuit B').

    This unitary satisfies: U_chi_psi |psi> = |chi>

    Args:
        psi: Source statevector
        chi: Target statevector

    Returns:
        UnitaryGate that implements the transition
    """
    # Build unitaries that prepare each state from |0...0>
    U_psi = statevector_to_unitary(psi)
    U_chi = statevector_to_unitary(chi)

    # Transition unitary: U_chi @ U_psi^dagger
    U_chi_psi = U_chi @ U_psi.conj().T

    # Verify it works
    psi_normalized = np.asarray(psi, dtype=np.complex128)
    psi_normalized = psi_normalized / np.linalg.norm(psi_normalized)
    chi_normalized = np.asarray(chi, dtype=np.complex128)
    chi_normalized = chi_normalized / np.linalg.norm(chi_normalized)

    result = U_chi_psi @ psi_normalized

```

```

    if not np.allclose(result, chi_normalized, atol=1e-10):
        raise ValueError("Transition unitary does not correctly map |psi>
to |chi>")

    return UnitaryGate(U_chi_psi)

def build_chi_state(class0_protos, class1_protos):
    """
    Build |chi> = sum_k |phi_k^0> - sum_k |phi_k^1>, normalized

    This constructs the reference state for ISDO classification.
    """
    chi = np.zeros_like(class0_protos[0], dtype=np.float64)

    for p in class0_protos:
        chi += p
    for p in class1_protos:
        chi -= p

    chi /= np.linalg.norm(chi)
    return chi

```

File: src/utils/paths.py

```

import yaml
import os

def load_paths(config_path="configs/paths.yaml"):
    with open(config_path, "r") as f:
        cfg = yaml.safe_load(f)

    base_root = cfg["base_root"]
    paths = {
        k: os.path.join(base_root, v)
        for k, v in cfg["paths"].items()
    }
    paths["class_count"] = cfg["class_count"]
    return base_root, paths

```

File: src/utils/seed.py

```

import random
import numpy as np
import torch
import os

def set_seed(seed: int = 42):
    /

```

```
# Python
random.seed(seed)

# NumPy
np.random.seed(seed)

# PyTorch
torch.manual_seed(seed)
torch.cuda.manual_seed(seed)
torch.cuda.manual_seed_all(seed)

# cuDNN (important)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False

# Extra safety (hash-based ops)
os.environ["PYTHONHASHSEED"] = str(seed)

print(f"🌱 Global seed set to {seed}")
```

File: src/utils/label_utils.py

```
# src/utils/label_utils.py
"""
Unified label conversion utilities for quantum classifier.

Standard convention:
- Binary: {0, 1} for storage and classical models
- Polar: {-1, +1} for quantum interference calculations
"""

import numpy as np

def binary_to_polar(labels):
    """
    Convert binary labels {0, 1} to polar {-1, +1}.

    Args:
        labels: array-like with values in {0, 1}

    Returns:
        numpy array with values in {-1, +1}
    """
    labels = np.asarray(labels)
    return 2 * labels - 1

def polar_to_binary(labels):
    """
    Convert polar labels {-1, +1} to binary {0, 1}.

    Args:

```

```
    labels: array-like with values in {-1, +1}

>Returns:
    numpy array with values in {0, 1}
"""
labels = np.asarray(labels)
return (labels + 1) // 2

def ensure_polar(labels):
    """
    Ensure labels are in polar format {-1, +1}.
    Automatically detects format and converts if needed.
    """
    labels = np.asarray(labels)
    unique_vals = np.unique(labels)

    if set(unique_vals).issubset({0, 1}):
        return binary_to_polar(labels)
    elif set(unique_vals).issubset({-1, 1}):
        return labels
    else:
        raise ValueError(f"Labels must be binary {{0,1}} or polar {{-1,+1}}. Got: {unique_vals}")

def ensure_binary(labels):
    """
    Ensure labels are in binary format {0, 1}.
    Automatically detects format and converts if needed.
    """
    labels = np.asarray(labels)
    unique_vals = np.unique(labels)

    if set(unique_vals).issubset({0, 1}):
        return labels
    elif set(unique_vals).issubset({-1, 1}):
        return polar_to_binary(labels)
    else:
        raise ValueError(f"Labels must be binary {{0,1}} or polar {{-1,+1}}. Got: {unique_vals}")
```

File: src/utils/**init**.py

File: src/data/pcam_loader.py

```
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
```

```
def get_pciam_dataset(data_dir='/home/tarakesh/Work/Repo/measurement-free-quantum-classifier/dataset', split='train', download=True, transform=None):
    """
    Wrapper for torchvision's built-in PCAM dataset.
    Automatically handles downloading and formatting.
    """
    if transform is None:
        # Default transformation for the hybrid model
        transform = transforms.Compose([
            transforms.ToTensor(), # Scales [0, 255] to [0.0, 1.0] and HWC to CHW
        ])
    dataset = datasets.PCAM(
        root=data_dir,
        split=split,
        download=download,
        transform=transform
    )
    return dataset

if __name__ == "__main__":
    print("PCAM Loader (using torchvision) initialized.")
```

File: src/data/transforms.py

```
from torchvision import transforms

def get_train_transforms():
    """
    Minimal, label-preserving augmentations for CNN training only.
    """
    return transforms.Compose([
        transforms.RandomHorizontalFlip(),
        transforms.RandomVerticalFlip(),
        transforms.ColorJitter(
            brightness=0.1,
            contrast=0.1,
            saturation=0.05,
        ),
        transforms.ToTensor(),
        transforms.Normalize(
            mean=[0.5, 0.5, 0.5],
            std=[0.5, 0.5, 0.5],
        ),
    ])

def get_eval_transforms():
```

```
"""
Deterministic transforms for validation, testing, and embedding
extraction.

"""
return transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.5, 0.5, 0.5],
        std=[0.5, 0.5, 0.5],
    ),
])
```

File: src/data/init.py

File: src/quantum/compute_qsvm_kernel.py

```
import os
import json
import numpy as np
from tqdm import tqdm

from qiskit_aer.primitives import SamplerV2
from qiskit.circuit.library import ZZFeatureMap
from qiskit_machine_learning.kernels import FidelityQuantumKernel
from qiskit_algorithms.state_fidelities import ComputeUncompute

from src.utils.paths import load_paths
from src.utils.seed import set_seed

# -----
# Reproducibility
# -----
set_seed(42)

# -----
# Load paths and data
# -----
BASE_ROOT, PATHS = load_paths()

EMBED_DIR = PATHS["embeddings"]
OUT_DIR = os.path.join(BASE_ROOT, "results", "qsvm_cache")
os.makedirs(OUT_DIR, exist_ok=True)

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))
```

```
train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_train = X[train_idx]
y_train = y[train_idx]

X_test = X[test_idx]
y_test = y[test_idx]

# -----
# SUBSAMPLING for Baseline Efficiency
# -----
# Limiting to 500 samples because O(N^2) kernel computation
# for 3500 samples would take ~17 hours on GPU.
MAX_TRAIN = 500000
MAX_TEST = 200000

if len(X_train) > MAX_TRAIN:
    print(f"Subsampling train set from {len(X_train)} to {MAX_TRAIN}...")
    rng = np.random.default_rng(42)
    indices = rng.choice(len(X_train), MAX_TRAIN, replace=False)
    X_train = X_train[indices]
    y_train = y_train[indices]

if len(X_test) > MAX_TEST:
    print(f"Subsampling test set from {len(X_test)} to {MAX_TEST}...")
    rng = np.random.default_rng(42)
    indices = rng.choice(len(X_test), MAX_TEST, replace=False)
    X_test = X_test[indices]
    y_test = y_test[indices]

# -----
# Normalize embeddings
# -----
X_train = X_train / np.linalg.norm(X_train, axis=1, keepdims=True)
X_test = X_test / np.linalg.norm(X_test, axis=1, keepdims=True)

# Infer number of qubits
dim = X_train.shape[1]
num_qubits = int(np.log2(dim))
assert 2 ** num_qubits == dim, "Embedding dimension must be 2^n"

# -----
# Define FIXED quantum feature map
# -----
feature_map = ZZFeatureMap(
    feature_dimension=num_qubits,
    reps=1,
    entanglement="linear"
)

# -----
# GPU Accelerated Backend (Aer SamplerV2)
# -----
```

```
sampler = SamplerV2(  
    options={"backend_options": {"method": "statevector", "device": "GPU"} }  
)  
fidelity = ComputeUncompute(sampler=sampler)  
  
quantum_kernel = FidelityQuantumKernel(  
    feature_map=feature_map,  
    fidelity=fidelity  
)  
  
# -----  
# Compute and save TRAIN kernel  
# -----  
print(f"Computing QSVM TRAIN kernel ({len(X_train)}x{len(X_train)})...")  
K_train = quantum_kernel.evaluate(X_train, X_train)  
np.save(os.path.join(OUT_DIR, "qsvm_kernel_train.npy"), K_train)  
  
# -----  
# Compute and save TEST kernel  
# -----  
print(f"Computing QSVM TEST kernel ({len(X_test)}x{len(X_train)})...")  
K_test = quantum_kernel.evaluate(X_test, X_train)  
np.save(os.path.join(OUT_DIR, "qsvm_kernel_test.npy"), K_test)  
  
# -----  
# Save Labels for verification  
# -----  
np.save(os.path.join(OUT_DIR, "y_train_sub.npy"), y_train)  
np.save(os.path.join(OUT_DIR, "y_test_sub.npy"), y_test)  
  
# -----  
# Save metadata  
# -----  
meta = {  
    "model": "QSVM",  
    "num_qubits": num_qubits,  
    "num_train": int(X_train.shape[0]),  
    "num_test": int(X_test.shape[0]),  
    "embedding_dimension": int(dim),  
    "subsampling": True  
}  
  
with open(os.path.join(OUT_DIR, "qsvm_kernel_meta.json"), "w") as f:  
    json.dump(meta, f, indent=2)  
  
print("QSVM kernel computation complete.")
```

File: src/quantum/**init**.py

File: src/training/test_adaptive_iqc.py

```
import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.IQL.models.adaptive_iqc import AdaptiveIQC
from src.utils.paths import load_paths


def main():
    # -----
    # Load data
    # -----
    _, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]

    X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
    y = np.load(os.path.join(EMBED_DIR, "val_labels.npy")) # ±1

    train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
    test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

    X_train, y_train = X[train_idx], y[train_idx]
    X_test, y_test = X[test_idx], y[test_idx]

    # quantum-safe normalization
    X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
    X_test /= np.linalg.norm(X_test, axis=1, keepdims=True)

    # -----
    # Train AdaptiveIQC
    # -----
    model = AdaptiveIQC(
        K_init=3,
        eta=0.1,
        percentile=5,
        consolidate=True,
    )

    model.fit(X_train, y_train)

    # -----
    # Evaluate
    # -----
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)

    print("✓ AdaptiveIQC Test Accuracy:", round(acc, 4))
```

```
if __name__ == "__main__":
    main()
```

File: src/training/verify_consistency.py

```
import numpy as np
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime2_online import OnlinePerceptron
from src.IQL.regimes.regime3a_wta import WinnerTakeAll
from src.IQL.regimes.regime3c_adaptive import AdaptiveMemory

def test_consistency():
    print("Running consistency tests...")

    # 1. Backend
    backend = ExactBackend()

    # 2. ClassState
    vec = np.array([1, 0, 0, 0], dtype=np.complex128)
    cs = ClassState(vec, backend)
    print("ClassState initialized.")

    psi = np.array([1, 0, 0, 0], dtype=np.complex128)
    score = cs.score(psi)
    print(f"ClassState score: {score}")
    assert np.isclose(score, 1.0)

    # 3. MemoryBank
    mb = MemoryBank([cs])
    print("MemoryBank initialized.")
    scores = mb.scores(psi)
    print(f"MemoryBank scores: {scores}")
    assert np.isclose(scores[0], 1.0)

    # 4. Models
    # OnlinePerceptron
    op = OnlinePerceptron(cs, eta=0.1)
    y_hat, s, updated = op.step(psi, 1)
    print(f"OnlinePerceptron step: y_hat={y_hat}, s={s}, updated={updated}")

    # WinnerTakeAll
    wta = WinnerTakeAll(mb, eta=0.1, backend=backend)
    y_hat, idx, updated = wta.step(psi, 1)
    print(f"WinnerTakeAll step: y_hat={y_hat}, idx={idx}, updated={updated}")

    # AdaptiveMemory
```

```
am = AdaptiveMemory(mb, eta=0.1, backend=backend)
margin, spawned = am.step(psi, 1)
print(f"AdaptiveMemory step: margin={margin}, spawned={spawned}")

print("All basic consistency tests passed!")

if __name__ == "__main__":
    test_consistency()
```

File: src/training/run_final_comparison.py

```
import os
import json
import numpy as np
from tqdm import tqdm
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC

from src.utils.paths import load_paths
from src.IQL.interference.exact_backend import ExactBackend
from src.IQL.interference.transition_backend import TransitionBackend
from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier


# -----
# Config
# -----
INCLUDE_QSVM = False
K_ISDO = 3 # chosen from K-sweep (best)

# -----
# Load paths and data
# -----
_, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]
PROTO_DIR = PATHS["class_prototypes"]
LOG_DIR = PATHS["logs"]
QSVM_DIR = os.path.join(PATHS["artifacts"], "qsvm_cache")

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy"))

test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))
X_test = X[test_idx]
y_test = y[test_idx]

# quantum-safe normalization (already true, but explicit)
X_test = X_test / np.linalg.norm(X_test, axis=1, keepdims=True)

# Load base prototype once to avoid disk I/O in loops
```

```
chi_single = np.load(os.path.join(PROTO_DIR, "K1/class1_proto0.npy"))

results = {}

# =====
# IQC - Exact (measurement-free)
# =====
exact_backend = ExactBackend()

print("Evaluating IQC-Exact...")
y_pred_exact = []
for psi in tqdm(X_test, desc="IQC Exact"):
    s = exact_backend.score(chi=chi_single, psi=psi)
    y_pred_exact.append(1 if s >= 0 else -1)

results["IQC_Exact_Backend"] = accuracy_score(y_test, y_pred_exact)

# =====
# IQC - Transition (circuit B')
# =====
transition_backend = TransitionBackend()

print("Evaluating IQC-Transition (Circuit-B')...")
y_pred_transition = []
for psi in tqdm(X_test, desc="IQC Transition"):
    s = transition_backend.score(chi=chi_single, psi=psi)
    y_pred_transition.append(1 if s >= 0 else -1)

results["IQC_Transition_Backend"] = accuracy_score(y_test,
y_pred_transition)

# =====
# ISDO - K-prototype interference ( Exact )
# =====
isdo = StaticISDOClassifier(PROTO_DIR, K_ISDO)
print(f"Evaluating ISDO-K (K={K_ISDO})...")
y_pred_isdo = isdo.predict(X_test)
results["ISDO_K"] = accuracy_score((y_test + 1) // 2, y_pred_isdo)

# =====
# Fidelity (SWAP test) - load cached result
# =====
results["Fidelity_SWAP"] = 0.8784 # from evaluate_swap_test_batch.py

# =====
# Classical baselines - load from logs
# =====
with open(os.path.join(LOG_DIR, "embedding_baseline_results.json")) as f:
    classical = json.load(f)

for k, v in classical.items():
    results[k] = v["accuracy"]

# =====
```

```

# QSVM (optional)
# -----
if INCLUDE_QSVM:
    print("Evaluating QSVM baseline...")
    try:
        K_train = np.load(os.path.join(QSVM_DIR, "qsvm_kernel_train.npy"))
        K_test = np.load(os.path.join(QSVM_DIR, "qsvm_kernel_test.npy"))
        y_train = np.load(os.path.join(QSVM_DIR, "y_train_sub.npy"))

        # Note: SVC expects kernel values, labels should correspond to
        # kernel indices
        qsvm = SVC(kernel="precomputed")
        qsvm.fit(K_train, y_train)

        y_test_sub = np.load(os.path.join(QSVM_DIR, "y_test_sub.npy"))
        y_pred_qsvm = qsvm.predict(K_test)
        results["QSVM"] = accuracy_score(y_test_sub, y_pred_qsvm)

    except Exception as e:
        print(f"QSVM evaluation skipped: {e}")
        results["QSVM"] = None

# -----
# Save
# -----
with open("final_comparison_results.json", "w") as f:
    json.dump(results, f, indent=2)

print("\n==== FINAL COMPARISON ====")
for k, v in results.items():
    if v is not None:
        print(f"{k:25s}: {v:.4f}")
    else:
        print(f"{k:25s}: N/A")

```

File: src/training/compare_best_iqc_vs_classical.py

```

import os
import json
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.IQL.training.adaptive_memory_trainer import AdaptiveMemoryTrainer

# -----
# Load paths
# -----
_, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]

```

```

LOG_DIR    = PATHS["logs"]

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy"))

train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx  = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_train, y_train = X[train_idx], y[train_idx]
X_test, y_test  = X[test_idx], y[test_idx]

X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
X_test  /= np.linalg.norm(X_test, axis=1, keepdims=True)

results = {}

# -----
# Best IQC
# -----
adaptive = AdaptiveMemoryTrainer()
adaptive.fit(X_train, y_train)
results["IQC_Adaptive"] = accuracy_score(
    y_test, adaptive.predict(X_test)
)

# -----
# Classical baselines (from logs)
# -----
with open(os.path.join(LOG_DIR, "embedding_baseline_results.json")) as f:
    classical = json.load(f)

for k, v in classical.items():
    results[k] = v["accuracy"]

print("\n== Best IQC vs Classical ==")
for k, v in results.items():
    print(f"{k:25s}: {v}")

```

File: src/training/test_fixed_memory_iqc.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.IQL.models.fixed_memory_iqc import FixedMemoryIQC

def main():
    # -----
    /

```

```

# Load paths
#
_, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]

#
# Load embeddings and labels (polar)
#
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy")) # ±1

train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_train, y_train = X[train_idx], y[train_idx]
X_test, y_test = X[test_idx], y[test_idx]

#
# Quantum-safe normalization (defensive)
#
X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
X_test /= np.linalg.norm(X_test, axis=1, keepdims=True)

#
# Train Fixed-Memory IQC
#
K = 5
model = FixedMemoryIQC(K=K, eta=0.1)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
acc = accuracy_score(y_test, y_pred)

print(f"✓ FixedMemoryIQC | K={K} | Test Accuracy: {acc:.4f}")

if __name__ == "__main__":
    main()

```

File: src/training/validate_backends.py

```

import numpy as np

from src.IQL.backends.exact import ExactBackend
from src.IQL.backends.hadamard import HadamardBackend
from src.IQL.backends.transition import TransitionBackend
from src.IQL.backends.prime_b import PrimeBBackend

def random_state(n_qubits, seed=None):
    /

```

```
if seed is not None:
    np.random.seed(seed)
dim = 2 ** n_qubits
v = np.random.randn(dim) + 1j * np.random.randn(dim)
return v / np.linalg.norm(v)

def run_backend_tests(n_qubits=3, n_tests=20):
    backends = {
        "Exact": ExactBackend(),
        "Hadamard": HadamardBackend(),
        "Transition": TransitionBackend(),
        "PrimeB": PrimeBBackend(),
    }

    print(f"\nRunning backend tests with {n_qubits} qubits\n")

    # Fix X
    chi = random_state(n_qubits, seed=42)

    scores = {name: [] for name in backends}

    for i in range(n_tests):
        psi = random_state(n_qubits, seed=100 + i)

        print(f"Test {i + 1}")
        for name, backend in backends.items():
            s = backend.score(chi, psi)
            scores[name].append(s)
            print(f"  {name:10s}: {s:+.6f}")
        print()

    # -----
    # Analysis
    # -----
    print("\n==== Backend Agreement Analysis ===\n")

    exact = np.array(scores["Exact"])

    for name in ["Hadamard", "Transition"]:
        diff = np.max(np.abs(exact - np.array(scores[name])))
        print(f"Max |Exact - {name}| = {diff:.2e}")

    # PrimeB: sign + ordering only
    primeb = np.array(scores["PrimeB"])

    sign_match = np.mean(np.sign(primeb) == np.sign(exact))
    print(f"\nPrimeB sign agreement with Exact: {sign_match * 100:.1f}%")

    # Rank correlation (ordering)
    exact_rank = np.argsort(exact)
    primeb_rank = np.argsort(primeb)
    rank_corr = np.corrcoef(exact_rank, primeb_rank)[0, 1]
    print(f"PrimeB rank correlation with Exact: {rank_corr:.3f}")
```

```
if __name__ == "__main__":
    run_backend_tests(n_qubits=3, n_tests=200)
```

```
"""
```

```
Test 194
```

```
    Exact      : -0.224492
    Hadamard   : -0.224492
    Transition: -0.224492
    PrimeB     : +0.095676
```

```
Test 195
```

```
    Exact      : -0.028519
    Hadamard   : -0.028519
    Transition: -0.028519
    PrimeB     : -0.423231
```

```
Test 196
```

```
    Exact      : +0.203938
    Hadamard   : +0.203938
    Transition: +0.203938
    PrimeB     : -0.201812
```

```
Test 197
```

```
    Exact      : +0.143895
    Hadamard   : +0.143895
    Transition: +0.143895
    PrimeB     : +0.035991
```

```
Test 198
```

```
    Exact      : -0.111603
    Hadamard   : -0.111603
    Transition: -0.111603
    PrimeB     : -0.143718
```

```
Test 199
```

```
    Exact      : +0.164120
    Hadamard   : +0.164120
    Transition: +0.164120
    PrimeB     : +0.107708
```

```
Test 200
```

```
    Exact      : +0.145881
    Hadamard   : +0.145881
    Transition: +0.145881
    PrimeB     : -0.250643
```

```
==== Backend Agreement Analysis ===
```

```
Max |Exact - Hadamard| = 3.22e-15
Max |Exact - Transition| = 4.97e-14
```

```
PrimeB sign agreement with Exact: 52.5%
PrimeB rank correlation with Exact: -0.004
"""
```

File: src/training/test_static_isdo_model.py

```
import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.IQL.models.static_isdo_model import StaticISDOModel

def main():
    # -----
    # Load paths
    # -----
    _, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]

    # -----
    # Load embeddings and labels
    # -----
    X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
    y = np.load(os.path.join(EMBED_DIR, "val_labels.npy")) # {0,1}

    train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
    test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

    X_train, y_train = X[train_idx], y[train_idx]
    X_test, y_test = X[test_idx], y[test_idx]

    # -----
    # Sanity: ensure quantum-safe normalization
    # -----
    X_train = X_train / np.linalg.norm(X_train, axis=1, keepdims=True)
    X_test = X_test / np.linalg.norm(X_test, axis=1, keepdims=True)

    # -----
    # Run Static ISDO Model
    # -----
    K = 5 # best K from sweep
    model = StaticISDOModel(K=K)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    print(f"✓ StaticISDOModel | K={K} | Test Accuracy: {acc:.4f}")
```

```
if __name__ == "__main__":
    main()
```

File: src/training/test_core_functionality.py

```
# tests/test_core_functionality.py
"""
Comprehensive test suite for quantum classifier core functionality.
Run with: python -m pytest tests/test_core_functionality.py -v
"""

import pytest
import numpy as np
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime2_online import OnlinePerceptron
from src.IQL.regimes.regime3a_wta import WinnerTakeAll
from src.utils.label_utils import binary_to_polar, polar_to_binary,
ensure_polar

class TestLabelConversions:
    """Test label conversion utilities."""

    def test_binary_to_polar(self):
        binary = np.array([0, 1, 0, 1])
        polar = binary_to_polar(binary)
        expected = np.array([-1, 1, -1, 1])
        np.testing.assert_array_equal(polar, expected)

    def test_polar_to_binary(self):
        polar = np.array([-1, 1, -1, 1])
        binary = polar_to_binary(polar)
        expected = np.array([0, 1, 0, 1])
        np.testing.assert_array_equal(binary, expected)

    def test_round_trip(self):
        binary = np.array([0, 1, 0, 1])
        polar = binary_to_polar(binary)
        back_to_binary = polar_to_binary(polar)
        np.testing.assert_array_equal(binary, back_to_binary)

    def test_ensure_polar_from_binary(self):
        binary = np.array([0, 1])
        polar = ensure_polar(binary)
        np.testing.assert_array_equal(polar, np.array([-1, 1]))

    def test_ensure_polar_from_polar(self):
        polar = np.array([-1, 1])
        result = ensure_polar(polar)
        np.testing.assert_array_equal(result, polar)
```

```
class TestClassState:  
    """Test ClassState functionality."""  
  
    def test_initialization(self):  
        backend = ExactBackend()  
        vec = np.array([1, 0, 0, 0], dtype=np.complex128)  
        cs = ClassState(vec, backend, label=1)  
  
        assert cs.label == 1  
        assert np.isclose(np.linalg.norm(cs.vector), 1.0)  
  
    def test_normalization(self):  
        backend = ExactBackend()  
        vec = np.array([3, 4, 0, 0], dtype=np.complex128)  
        cs = ClassState(vec, backend)  
  
        # Should be normalized  
        assert np.isclose(np.linalg.norm(cs.vector), 1.0)  
  
    def test_score_orthogonal(self):  
        backend = ExactBackend()  
        chi = np.array([1, 0, 0, 0], dtype=np.complex128)  
        psi = np.array([0, 1, 0, 0], dtype=np.complex128)  
  
        cs = ClassState(chi, backend)  
        score = cs.score(psi)  
  
        assert np.isclose(score, 0.0, atol=1e-10)  
  
    def test_score_parallel(self):  
        backend = ExactBackend()  
        vec = np.array([1, 0, 0, 0], dtype=np.complex128)  
  
        cs = ClassState(vec, backend)  
        score = cs.score(vec)  
  
        assert np.isclose(score, 1.0, atol=1e-10)  
  
class TestMemoryBank:  
    """Test MemoryBank functionality."""  
  
    def test_initialization(self):  
        backend = ExactBackend()  
        cs1 = ClassState(np.array([1, 0, 0, 0], dtype=np.complex128),  
        backend, label=0)  
        cs2 = ClassState(np.array([0, 1, 0, 0], dtype=np.complex128),  
        backend, label=1)  
  
        mb = MemoryBank([cs1, cs2])  
        assert len(mb.class_states) == 2  
  
    def test_add_memory(self):  
        backend = ExactBackend()  
        /
```

```
    cs = ClassState(np.array([1, 0, 0, 0], dtype=np.complex128),
backend)
    mb = MemoryBank([cs])

    new_vec = np.array([0, 1, 0, 0], dtype=np.complex128)
    mb.add_memory(new_vec, backend, label=1)

    assert len(mb.class_states) == 2
    assert mb.class_states[1].label == 1

def test_remove_memory(self):
    backend = ExactBackend()
    cs1 = ClassState(np.array([1, 0, 0, 0], dtype=np.complex128),
backend)
    cs2 = ClassState(np.array([0, 1, 0, 0], dtype=np.complex128),
backend)
    cs3 = ClassState(np.array([0, 0, 1, 0], dtype=np.complex128),
backend)

    mb = MemoryBank([cs1, cs2, cs3])
    mb.remove(1)

    assert len(mb.class_states) == 2

def test_winner(self):
    backend = ExactBackend()
    cs1 = ClassState(np.array([1, 0, 0, 0], dtype=np.complex128),
backend)
    cs2 = ClassState(np.array([0, 1, 0, 0], dtype=np.complex128),
backend)

    mb = MemoryBank([cs1, cs2])

    # Test with state close to cs1
    psi = np.array([0.9, 0.1, 0, 0], dtype=np.complex128)
    psi /= np.linalg.norm(psi)

    idx, score = mb.winner(psi)
    assert idx == 0 # Should select cs1

class TestOnlinePerceptron:
    """Test OnlinePerceptron regime."""

    def test_training_convergence(self):
        backend = ExactBackend()

        # Simple linearly separable data
        X = np.array([
            [1, 0, 0, 0],
            [0.9, 0.1, 0, 0],
            [0, 1, 0, 0],
            [0, 0.9, 0.1, 0],
        ], dtype=np.complex128)
```

```
y = np.array([1, 1, -1, -1])

# Initialize with random state
chi0 = np.array([0.5, 0.5, 0, 0], dtype=np.complex128)
chi0 /= np.linalg.norm(chi0)

cs = ClassState(chi0, backend)
trainer = OnlinePerceptron(cs, eta=0.1)

acc = trainer.fit(X, y)

# Should achieve reasonable accuracy on this simple problem
assert acc >= 0.5

def test_save_load(self):
    import tempfile
    import os

    backend = ExactBackend()
    chi = np.array([1, 0, 0, 0], dtype=np.complex128)
    cs = ClassState(chi, backend)
    trainer = OnlinePerceptron(cs, eta=0.1)

    # Train a bit
    X = np.array([[1, 0, 0, 0], [0, 1, 0, 0]], dtype=np.complex128)
    y = np.array([1, -1])
    trainer.fit(X, y)

    # Save
    with tempfile.NamedTemporaryFile(delete=False, suffix='.pkl') as f:
        temp_path = f.name

    try:
        trainer.save(temp_path)

        # Load
        loaded_trainer = OnlinePerceptron.load(temp_path)

        # Check attributes
        assert loaded_trainer.eta == trainer.eta
        assert loaded_trainer.num_updates == trainer.num_updates
        assert len(loaded_trainer.history["scores"]) ==
len(trainer.history["scores"])
    finally:
        os.unlink(temp_path)

def test_integration():
    """Integration test for full pipeline."""
    from src.IQL.regimes.regime3a_wta import WinnerTakeAll

    backend = ExactBackend()

    # Create simple data
    X = np.array([
        /
```

```

[1, 0, 0, 0],
[0.9, 0.1, 0, 0],
[0, 1, 0, 0],
[0.1, 0.9, 0, 0],
], dtype=np.complex128)
y = np.array([1, 1, -1, -1])

# Initialize memory bank
cs1 = ClassState(np.array([1, 0, 0, 0], dtype=np.complex128), backend,
label=1)
cs2 = ClassState(np.array([0, 1, 0, 0], dtype=np.complex128), backend,
label=-1)
mb = MemoryBank([cs1, cs2])

# Train with WTA
wta = WinnerTakeAll(mb, eta=0.1, backend=backend)
acc = wta.fit(X, y)

# Predict
predictions = wta.predict(X)

assert len(predictions) == len(X)
assert acc >= 0.5 # Should do reasonably well

if __name__ == "__main__":
    pytest.main([__file__, "-v"])

```

File: src/training/compare_iqc_algorithms.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier
from src.IQL.training.online_perceptron_trainer import
OnlinePerceptronTrainer
from src.IQL.training.adaptive_memory_trainer import AdaptiveMemoryTrainer
from src.IQL.states.class_state import ClassState
from src.IQL.memory.memory_bank import MemoryBank
import pickle

# -----
# Load data
# -----
_, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]
PROTO_DIR = PATHS["class_prototypes"]

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy"))

```

```
train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_train, y_train = X[train_idx], y[train_idx]
X_test, y_test = X[test_idx], y[test_idx]

X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
X_test /= np.linalg.norm(X_test, axis=1, keepdims=True)

results = {}

# -----
# Static ISDO
# -----
isdo = StaticISDOClassifier(PROTO_DIR, K=3)
results["Static_ISDO"] = accuracy_score((y_test + 1)//2,
                                         isdo.predict(X_test))

# -----
# IQC-Online (Regime-2)
# -----

# bootstrap initialization (important!)
chi0 = np.zeros_like(X_train[0])
for psi, label in zip(X_train[:10], y_train[:10]):
    chi0 += label * psi
chi0 = chi0 / np.linalg.norm(chi0)

class_state = ClassState(chi0)
online = OnlinePerceptronTrainer(class_state, eta=0.1)
online.fit(X_train, y_train)
results["IQC_Online"] = accuracy_score(y_test, online.predict(X_test))

# -----
# IQC-Adaptive Memory (Regime-3C)
# -----

MEMORY_PATH = os.path.join(PATHS["artifacts"], "regime3c_memory.pkl")

with open(MEMORY_PATH, "rb") as f:
    memory_bank = pickle.load(f)

adaptive = AdaptiveMemoryTrainer(
    memory_bank=memory_bank,
    eta=0.1,
    percentile=5,          # τ = 5th percentile of margins
    tau_abs = -0.121,
    margin_window=500
)
adaptive.fit(X_train, y_train)

results["IQC_Adaptive"] = accuracy_score(
    y_test, adaptive.predict(X_test))
```

```
)  
results["Adaptive_Memory_Size"] = adaptive.memory_size()  
  
print("\n==== IQC Algorithm Comparison ===")  
for k, v in results.items():  
    print(f"{k:25s}: {v}")  
  
## output  
"""  
==== IQC Algorithm Comparison ===  
Static_ISDO : 0.8806666666666667  
IQC_Online : 0.904  
IQC_Adaptive : 0.56  
Adaptive_Memory_Size : 45  
"""
```

File: src/training/protocol_online/train_perceptron.py

```
import numpy as np  
import os  
  
from src.IQL.learning.class_state import ClassState  
from src.IQL.encoding.embedding_to_state import embedding_to_state  
from src.IQL.regimes.regime2_online import OnlinePerceptron  
from src.IQL.learning.metrics import summarize_training  
from src.IQL.backends.exact import ExactBackend  
from src.utils.paths import load_paths  
from src.utils.seed import set_seed  
  
# -----  
# Reproducibility  
# -----  
set_seed(42)  
  
# -----  
# Load paths  
# -----  
_, PATHS = load_paths()  
EMBED_DIR = PATHS["embeddings"]  
  
os.makedirs(EMBED_DIR, exist_ok=True)  
  
# -----  
# Load embeddings (TRAIN ONLY)  
# -----  
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))  
y = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy"))  
train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))  
  
X_train = X[train_idx]  
y_train = y[train_idx]
```

```

print("Loaded train embeddings:", X_train.shape)

def main():

    chi0 = np.zeros_like(X_train[0])
    for psi, label in zip(X_train[:10], y_train[:10]):
        chi0 += label * psi
    chi0 = chi0 / np.linalg.norm(chi0)

    class_state = ClassState(chi0, backend=ExactBackend())
    trainer = OnlinePerceptron(class_state, eta=0.1)

    acc = trainer.fit(X_train, y_train)
    stats = summarize_training(trainer.history)

    print("Final accuracy:", acc)
    print("Training stats:", stats)

if __name__ == "__main__":
    main()

### output
"""
🌱 Global seed set to 42
Loaded train embeddings: (3500, 32)
Final accuracy: 0.8562857142857143
Training stats: {'mean_margin': 0.14930659062683652, 'min_margin': -0.7069261085786833, 'num_updates': 503, 'update_rate': 0.1437142857142857}
"""

```

File: src/training/protocol_static/evaluate_isdo_k_sweep.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier
from src.utils.paths import load_paths
import matplotlib.pyplot as plt

BASE_ROOT, PATHS = load_paths()

EMBED_DIR = PATHS["embeddings"]
PROTO_BASE = PATHS["class_prototypes"]

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

```

```

test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_test = X[test_idx]
y_test = y[test_idx]

accuracy = []
for K in PATHS["class_count"]["K_values"]:

    clf = StaticISDOClassifier(PROTO_BASE, K)

    y_pred = clf.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    accuracy.append(acc)
    print(f"ISDO | K={K} | Accuracy: {acc:.4f}")

"""
ISDO | K=1 | Accuracy: 0.8827
ISDO | K=2 | Accuracy: 0.8800
ISDO | K=3 | Accuracy: 0.8960 ## best
ISDO | K=5 | Accuracy: 0.8840
ISDO | K=7 | Accuracy: 0.8840
ISDO | K=11 | Accuracy: 0.8820
ISDO | K=13 | Accuracy: 0.8800
ISDO | K=17 | Accuracy: 0.8740
ISDO | K=19 | Accuracy: 0.8780
ISDO | K=23 | Accuracy: 0.8747
"""

plt.plot(PATHS["class_count"]["K_values"], accuracy, marker="o")
plt.xlabel("Number of prototypes per class (K)")
plt.ylabel("Test Accuracy")
plt.title("ISDO Accuracy vs Interference Capacity")
plt.grid(True)
plt.savefig(os.path.join(PATHS["figures"], "isdo_k_sweep.png"))

```

File: src/training/protocol_static/evaluate_static_isdo.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier
from src.utils.paths import load_paths

BASE_ROOT, PATHS = load_paths()

EMBED_DIR = PATHS["embeddings"]
PROTO_DIR = PATHS["class_prototypes"]
K = int(PATHS["class_count"]["K"])

```

```
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_test = X[test_idx]
y_test = y[test_idx]

clf = StaticISDOClassifier(PROTO_DIR, K)
y_pred = clf.predict(X_test)

acc = accuracy_score(y_test, y_pred)
print(f"ISDO Accuracy (test): {acc:.4f}")

"""
ISDO Accuracy (test): 0.8840
"""
```

File: src/training/classical/make_embedding_split.py

```
import os
import numpy as np
from sklearn.model_selection import train_test_split

from src.utils.paths import load_paths
from src.utils.seed import set_seed
set_seed(42)

BASE_ROOT, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

indices = np.arange(len(y))

train_idx, test_idx = train_test_split(
    indices,
    test_size=0.3,
    random_state=42,
    stratify=y
)

np.save(os.path.join(EMBED_DIR, "split_train_idx.npy"), train_idx)
np.save(os.path.join(EMBED_DIR, "split_test_idx.npy"), test_idx)

print("Saved split:")
print("Train:", len(train_idx))
print("Test :", len(test_idx))
```

File: src/training/classical/train_embedding_models.py

```
import os
import json
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, normalize
from sklearn.metrics import accuracy_score, roc_auc_score

from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.neighbors import KNeighborsClassifier

from src.utils.paths import load_paths
from src.utils.seed import set_seed
set_seed(42)

# -----
# Load paths
# -----
BASE_ROOT, PATHS = load_paths()

EMBED_DIR = PATHS["embeddings"]
LOG_DIR = PATHS["logs"]
os.makedirs(LOG_DIR, exist_ok=True)

# -----
# Load embeddings
# -----
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

print("Loaded embeddings:", X.shape)

# -----
# Preprocessing (DEPRECATED: Now handled in extract_embeddings.py)
# -----
# # 1) Standardize (important for linear models)
# scaler = StandardScaler()
# X_std = scaler.fit_transform(X)
#
# # 2) L2-normalize (important for similarity & quantum)
# X_l2 = normalize(X_std, norm="l2")

# -----
# Train / test split
# -----
```

```
# Using raw pre-normalized float64 embeddings for all models
Xtr = X[train_idx]
Xte = X[test_idx]
ytr = y[train_idx]
yte = y[test_idx]

results = {}

# =====
# [1] Logistic Regression (Linear separability)
# =====
print("\nTraining Logistic Regression...")
logreg = LogisticRegression(
    max_iter=1000,
    n_jobs=-1
)
logreg.fit(Xtr, ytr)

pred_lr = logreg.predict(Xte)
proba_lr = logreg.predict_proba(Xte)[:, 1]

results["LogisticRegression"] = {
    "accuracy": accuracy_score(yte, pred_lr),
    "auc": roc_auc_score(yte, proba_lr)
}

# =====
# [2] Linear SVM (Max-margin)
# =====
print("Training Linear SVM...")
svm = LinearSVC()
svm.fit(Xtr, ytr)

pred_svm = svm.predict(Xte)

results["LinearSVM"] = {
    "accuracy": accuracy_score(yte, pred_svm),
    "auc": None    # LinearSVC has no probability estimates
}

# =====
# [3] k-NN (Distance-based similarity)
# =====
print("Training k-NN...")
knn = KNeighborsClassifier(
    n_neighbors=5,
    metric="euclidean"
)
knn.fit(Xtr, ytr)
print("Knn neighbors:", knn.n_neighbors)
pred_knn = knn.predict(Xte)
proba_knn = knn.predict_proba(Xte)[:, 1]

results["kNN"] = {
```

```

    "accuracy": accuracy_score(yte, pred_knn),
    "auc": roc_auc_score(yte, proba_knn)
}

# -----
# Save results
# -----
with open(os.path.join(LOG_DIR, "embedding_baseline_results.json"), "w") as f:
    json.dump(results, f, indent=2)

# -----
# Print summary
# -----
print("\n==== Embedding Baseline Results ===")
for model, metrics in results.items():
    print(
        f"{model:>18} | "
        f"Acc: {metrics['accuracy']:.4f} | "
        f"AUC: {metrics['auc']}"
    )

## output
"""
🌱 Global seed set to 42
Loaded embeddings: (5000, 32)

Training Logistic Regression...
Training Linear SVM...
Training k-NN...
Knn neighbors: 5

==== Embedding Baseline Results ===
LogisticRegression | Acc: 0.9047 | AUC: 0.9664224751066857
    LinearSVM | Acc: 0.9053 | AUC: None
        kNN | Acc: 0.9260 | AUC: 0.9711219772403983
"""

```

File: src/training/classical/extract_embeddings.py

```

import os
import torch
import numpy as np
from torch.utils.data import DataLoader, Subset
from tqdm import tqdm

from src.classical.cnn import PCamCNN
from src.data.pcam_loader import get_pcam_dataset
from src.data.transforms import get_eval_transforms
from src.utils.paths import load_paths
from src.utils.seed import set_seed

```

```

set_seed(42)

BASE_ROOT, PATHS = load_paths()
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

CHECKPOINT = os.path.join(PATHS["checkpoints"], "pcam_cnn_best.pt")
os.makedirs(PATHS["embeddings"], exist_ok=True)

model = PCamCNN(embedding_dim=32).to(DEVICE)
model.load_state_dict(torch.load(CHECKPOINT, map_location=DEVICE))
model.eval()

dataset = get_pcam_dataset(PATHS["dataset"], "val", get_eval_transforms())
subset = Subset(dataset, range(5000))
loader = DataLoader(subset, batch_size=128, num_workers=6, pin_memory=True)

embeds, labels, lable_polar = [], [], []

with torch.no_grad():
    for x, y in tqdm(loader):
        z = model(x.to(DEVICE), return_embedding=True)
        # Convert to float64 FIRST, then normalize for maximum precision
        z = z.to(torch.float64)
        z = torch.nn.functional.normalize(z, p=2, dim=1)

        embeds.append(z.cpu().numpy())
        labels.append(y.numpy().astype(np.float64))
        lable_polar.append((y.numpy())*2 - 1).astype(np.float64)

np.save(os.path.join(PATHS["embeddings"], "val_embeddings.npy"),
np.vstack(embeds).astype(np.float64))
np.save(os.path.join(PATHS["embeddings"], "val_labels.npy"),
np.concatenate(labels).astype(np.float64))
np.save(os.path.join(PATHS["embeddings"], "val_labels_polar.npy"),
np.concatenate(lable_polar).astype(np.float64))

```

File: src/training/classical/visualize_embeddings.py

```

import os
import numpy as np
import matplotlib.pyplot as plt
from sklearn.manifold import TSNE
from src.utils.paths import load_paths
from src.utils.seed import set_seed

set_seed(42)

_, PATHS = load_paths()

X = np.load(os.path.join(PATHS["embeddings"], "val_embeddings.npy"))

```

```
y = np.load(os.path.join(PATHS["embeddings"], "val_labels.npy"))

tsne = TSNE(n_components=2, perplexity=30, max_iter=1000, random_state=42)
X2 = tsne.fit_transform(X)

plt.figure(figsize=(7, 6))
plt.scatter(X2[y == 0, 0], X2[y == 0, 1], s=8, label="Benign")
plt.scatter(X2[y == 1, 0], X2[y == 1, 1], s=8, label="Malignant")
plt.legend()
plt.savefig(os.path.join(PATHS["figures"], "embedding_tsne.png"), dpi=300)
plt.show()
```

File: src/training/classical/train_cnn.py

```
import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader
from tqdm import tqdm
import json
import matplotlib.pyplot as plt

from src.classical.cnn import PCamCNN
from src.data.pcam_loader import get_pcam_dataset
from src.data.transforms import get_train_transforms, get_eval_transforms
from src.utils.paths import load_paths
from src.utils.seed import set_seed

set_seed(42)
#torch.backends.cudnn.benchmark = True

# -----
# Load paths
# -----
BASE_ROOT, PATHS = load_paths()
DATA_ROOT = PATHS["dataset"]

# -----
# Config
# -----
BATCH_SIZE = 64
EPOCHS = 30
LR = 1e-3
EMBEDDING_DIM = 32
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

os.makedirs(PATHS["checkpoints"], exist_ok=True)
os.makedirs(PATHS["logs"], exist_ok=True)
os.makedirs(PATHS["figures"], exist_ok=True)
```

```
# -----
# Training / Evaluation loops
# -----
def train_one_epoch(model, loader, criterion, optimizer):
    model.train()
    running_loss, correct, total = 0.0, 0, 0

    for images, labels in tqdm(loader, desc="Training", leave=False):
        images, labels = images.to(DEVICE), labels.to(DEVICE)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * images.size(0)
        correct += outputs.argmax(1).eq(labels).sum().item()
        total += labels.size(0)

    return running_loss / total, correct / total

@torch.no_grad()
def evaluate(model, loader, criterion):
    model.eval()
    running_loss, correct, total = 0.0, 0, 0

    for images, labels in tqdm(loader, desc="Validation", leave=False):
        images, labels = images.to(DEVICE), labels.to(DEVICE)
        outputs = model(images)
        loss = criterion(outputs, labels)

        running_loss += loss.item() * images.size(0)
        correct += outputs.argmax(1).eq(labels).sum().item()
        total += labels.size(0)

    return running_loss / total, correct / total

def main():
    print(f"🚀 Training on device: {DEVICE}")

    train_set = get_pciam_dataset(DATA_ROOT, "train",
get_train_transforms())
    val_set = get_pciam_dataset(DATA_ROOT, "val", get_eval_transforms())

    train_loader = DataLoader(train_set, BATCH_SIZE, shuffle=True,
num_workers=6, pin_memory=True)
    val_loader = DataLoader(val_set, BATCH_SIZE, shuffle=False,
num_workers=6, pin_memory=True)

    model = PCamCNN(embedding_dim=EMBEDDING_DIM).to(DEVICE)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=LR,
```

```
weight_decay=1e-4)
    scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
        optimizer, mode="max", factor=0.5, patience=2
    )

    best_val_acc, patience, wait = 0.0, 10, 0
    history = {k: [] for k in ["train_loss", "train_acc", "val_loss",
"val_acc"]}

    for epoch in range(1, EPOCHS + 1):
        print(f"\nEpoch {epoch}/{EPOCHS}")

        tr_loss, tr_acc = train_one_epoch(model, train_loader, criterion,
optimizer)
        val_loss, val_acc = evaluate(model, val_loader, criterion)
        scheduler.step(val_acc)

        history["train_loss"].append(tr_loss)
        history["train_acc"].append(tr_acc)
        history["val_loss"].append(val_loss)
        history["val_acc"].append(val_acc)

        print(f"Train Acc {tr_acc:.4f} | Val Acc {val_acc:.4f}")

        if val_acc > best_val_acc:
            best_val_acc = val_acc
            torch.save(model.state_dict(),
os.path.join(PATHS["checkpoints"], "pcam_cnn_best.pt"))
            print("✓ Best validation accuracy reached : Saved checkpoint")
            wait = 0
        else:
            wait += 1

        if wait >= patience:
            print("■ Early stopping")
            break

    torch.save(model.state_dict(), os.path.join(PATHS["checkpoints"],
"pcam_cnn_final.pt"))
    print("✓ Final checkpoint saved")
    # Save logs
    with open(os.path.join(PATHS["logs"], "train_history.json"), "w") as f:
        json.dump(history, f, indent=2)

    # Plots
    epochs = range(1, len(history["train_loss"])) + 1
    plt.figure()
    plt.plot(epochs, history["train_acc"], label="Train")
    plt.plot(epochs, history["val_acc"], label="Val")
    plt.legend()
    plt.savefig(os.path.join(PATHS["figures"], "cnn_accuracy.png"))
    plt.close()

    plt.figure()
```

```
plt.plot(epochs, history["train_loss"], label="Train")
plt.plot(epochs, history["val_loss"], label="Val")
plt.legend()
plt.savefig(os.path.join(PATHS["figures"], "cnn_loss.png"))
plt.close()

if __name__ == "__main__":
    main()
```

File: src/training/classical/verify_embeddings.py

```
import os
import numpy as np
from src.utils.paths import load_paths

def verify_embeddings():
    BASE_ROOT, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]

    file_path = os.path.join(EMBED_DIR, "val_embeddings.npy")
    if not os.path.exists(file_path):
        print(f"File not found: {file_path}")
        return

    print(f"Verifying: {file_path}")
    X = np.load(file_path)
    print(f"Shape: {X.shape}, Dtype: {X.dtype}")

    # Calculate norm-squared for each sample
    norms_sq = np.sum(X**2, axis=1)

    max_val = np.max(norms_sq)
    min_val = np.min(norms_sq)
    mean_val = np.mean(norms_sq)

    print(f"Max norm squared: {max_val:.15f}")
    print(f"Min norm squared: {min_val:.15f}")
    print(f"Mean norm squared: {mean_val:.15f}")

    # Qiskit usually has a tolerance around 1e-8 or 1e-10
    tolerance = 1e-8
    violations = np.sum(np.abs(norms_sq - 1.0) > tolerance)

    print(f"Violations (> {tolerance} absolute diff from 1.0): {violations}")

    if violations > 0:
        idx = np.argmax(np.abs(norms_sq - 1.0))
        print(f"Worst violation at index {idx}: {norms_sq[idx]:.15f}")
```

```
if __name__ == "__main__":
    verify_embeddings()
```

File: src/training/classical/visualize_pcam.py

```
import matplotlib.pyplot as plt
from src.data.pcam_loader import get_pcam_dataset
from src.utils.paths import load_paths
from src.utils.seed import set_seed

set_seed(42)

_, PATHS = load_paths()

dataset = get_pcam_dataset(PATHS["dataset"], "test")

plt.figure(figsize=(10, 5))
for i in range(2):
    img, label = dataset[i]
    plt.subplot(1, 2, i + 1)
    plt.imshow(img.permute(1, 2, 0))
    plt.title("Malignant" if label else "Benign")
    plt.axis("off")

plt.show()
```

File: src/training/protocol_adaptive/consolidate_memory.py

```
import os
import numpy as np

from src.utils.paths import load_paths
from src.utils.seed import set_seed

from src.IQL.encoding.embedding_to_state import embedding_to_state
from src.IQL.regimes.regime3a_wta import WinnerTakeAll
from src.IQL.inference.weighted_vote_classifier import
WeightedVoteClassifier
from src.IQL.backends.exact import ExactBackend
from src.IQL.learning.memory_bank import MemoryBank
import pickle

# -----
# Reproducibility
# -----
set_seed(42)
```

```
# -----
# Load paths
# -----
_, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]
os.makedirs(EMBED_DIR, exist_ok=True)

# -----
# Load embeddings (TRAIN SPLIT)
# -----
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy"))
train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))

X_train = X[train_idx]
y_train = y[train_idx]

print("Loaded train embeddings:", X_train.shape)

# -----
# ❾ LOAD MEMORY BANK FROM REGIME 3-C
# -----
# IMPORTANT:
# This must be the SAME memory_bank produced by Regime 3-C

MEMORY_PATH = os.path.join(PATHS["artifacts"], "regime3c_memory.pkl")

with open(MEMORY_PATH, "rb") as f:
    memory_bank = pickle.load(f)

print("Loaded memory bank with",
      len(memory_bank.class_states),
      "memories")

# -----
# 🔄 CONSOLIDATION PHASE (NO GROWTH)
# -----
# Use Regime 3-A trainer:
# - updates memories
# - NO spawning logic
trainer = WinnerTakeAll(
    memory_bank=memory_bank,
    eta=0.05,          # slightly smaller eta for stabilization
    backend=ExactBackend()
)

acc_train = trainer.fit(X_train, y_train)
print("Consolidation pass accuracy:", acc_train)
print("Updates during consolidation:", trainer.num_updates)
```

```
# -----
# \|l FINAL EVALUATION (Regime 3-B inference)
# -----
classifier = WeightedVoteClassifier(memory_bank)

correct = 0
for x, y in zip(X_train, y_train):
    if classifier.predict(x) == y:
        correct += 1

final_acc = correct / len(X_train)
print("FINAL Regime 3-C accuracy:", final_acc)

### output
"""
🌱 Global seed set to 42
Loaded train embeddings: (3500, 32)
Loaded memory bank with 22 memories
Consolidation pass accuracy: 0.8048571428571428
Updates during consolidation: 683
FINAL Regime 3-C accuracy: 0.884
"""
```

File: src/training/protocol_adaptive/train_adaptive_memory.py

```
import os
import numpy as np
from collections import Counter

from src.utils.paths import load_paths
from src.utils.seed import set_seed

from src.IQL.learning.class_state import ClassState
from src.IQL.encoding.embedding_to_state import embedding_to_state
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend

from src.IQL.regimes.regime3c_adaptive import AdaptiveMemory
from src.IQL.inference.weighted_vote_classifier import
WeightedVoteClassifier
import pickle

# -----
# Reproducibility
# -----
set_seed(42)
```

```
# -----
# Load paths
# -----
_, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]
MEMORY_PATH = os.path.join(PATHS["artifacts"], "regime3c_memory.pkl")

os.makedirs(EMBED_DIR, exist_ok=True)
os.makedirs(PATHS["artifacts"], exist_ok=True)

# -----
# Load embeddings (TRAIN SPLIT)
# -----
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy"))
train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))

X_train = X[train_idx]
y_train = y[train_idx]

print("Loaded train embeddings:", X_train.shape)

# -----
# Initialize memory bank (M = 3)
# -----
d = X_train[0].shape[0]

backend = ExactBackend()

class_states = []
for _ in range(3):
    v = np.random.randn(d)
    v /= np.linalg.norm(v)
    class_states.append(ClassState(v, backend=backend))

memory_bank = MemoryBank(
    class_states=class_states
)

print("Initial number of memories:", len(memory_bank.class_states))

# -----
# Train Regime 3-C (percentile-based τ)
# -----
trainer = AdaptiveMemory(
    memory_bank=memory_bank,
    eta=0.1,
    percentile=5,          # τ = 5th percentile of margins
    tau_abs = -0.121,
```

```
margin_window=500,    # sliding window for stability
backend=backend,
)

trainer.fit(X_train, y_train)

print("Training finished.")
print("Number of memories after training:", len(memory_bank.class_states))
print("Number of spawned memories:", trainer.num_spawns)
print("Number of updates:", trainer.num_updates)

# -----
# Evaluate using Regime 3-B inference
# -----
classifier = WeightedVoteClassifier(memory_bank)

correct = 0
for psi, y in zip(X_train, y_train):
    if classifier.predict(psi) == y:
        correct += 1

acc_3c = correct / len(X_train)
print("Regime 3-C accuracy (3-B inference):", acc_3c)

# -----
# Optional diagnostics
# -----
print("Final memory count:", len(memory_bank.class_states))

with open(MEMORY_PATH, "wb") as f:
    pickle.dump(memory_bank, f)

print("Saved Regime 3-C memory bank.")

### output
"""
 Global seed set to 42
Loaded train embeddings: (3500, 32)
Initial number of memories: 3
Training finished.
Number of memories after training: 22
Number of spawned memories: 19
Number of updates: 429
Regime 3-C accuracy (3-B inference): 0.788
Final memory count: 22
Saved Regime 3-C memory bank.
"""
```

File: src/classical/cnn.py

```
import torch
import torch.nn as nn
import torch.nn.functional as F

class PCamCNN(nn.Module):
    """
    Lightweight CNN for PCam feature extraction.
    Produces low-dimensional embeddings suitable for quantum encoding.
    """

    def __init__(self, embedding_dim: int = 32, num_classes: int = 2):
        super().__init__()

        # ----- Convolutional backbone -----
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2), # 48x48

            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2), # 24x24

            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),

            nn.AdaptiveAvgPool2d((1, 1)) # 128 x 1 x 1
        )

        # ----- Embedding head -----
        self.embedding = nn.Linear(128, embedding_dim)

        # ----- Temporary classifier (used ONLY for CNN training) -----
        --
        self.classifier = nn.Linear(embedding_dim, num_classes)

    def forward(self, x, return_embedding: bool = False):
        x = self.features(x)
        x = x.view(x.size(0), -1) # flatten

        embedding = self.embedding(x)
        embedding = F.relu(embedding)

        if return_embedding:
            return embedding

        logits = self.classifier(embedding)
        return logits
```

File: src/classical/**init**.py

```
{  
    "embedding_dim": 32,  
    "classes": [  
        0,  
        1  
    ],  
    "normalization": "l2",  
    "source": "mean_of_class_embeddings"  
}
```

File: results/logs/train_history.json

```
{  
    "train_loss": [  
        0.323274524165754,  
        0.23782655238210282,  
        0.1993992631250876,  
        0.17781282487430872,  
        0.16371910747557195,  
        0.1537160865741498,  
        0.14631224803679288,  
        0.1406550945730487,  
        0.13497550318106732,  
        0.11610426991182976,  
        0.11148261162816198,  
        0.10862913947039488,  
        0.09541817462331892,  
        0.09339981933680974,  
        0.0910517916313438,  
        0.08256717906601807  
    "train_acc": [  
        0.8622550964355469,  
        0.9047431945800781,  
        0.9226036071777344,  
        0.9327926635742188,  
        0.9391098022460938,  
        0.9431877136230469,  
        0.9461746215820312,  
    ]  
}
```

```
    0.9482002258300781,  
    0.9505386352539062,  
    0.9581375122070312,  
    0.9599342346191406,  
    0.9608840942382812,  
    0.9665145874023438,  
    0.9672470092773438,  
    0.9680290222167969,  
    0.9713249206542969  
],  
"val_loss": [  
    0.7608505549724214,  
    0.3770719189342344,  
    0.33603281057730783,  

```

File: results/logs/embedding_baseline_results.json

```
{  
  "LogisticRegression": {  
    "accuracy": 0.9046666666666666,  
    "auc": 0.9664224751066857  
  },  
  "LinearSVM": {  
    "accuracy": 0.9053333333333333,  
    "auc": null  
  },  
  "kNN": {  
    "accuracy": 0.926,  
    "auc": 0.9711219772403983  
  }  
}
```