

Project Summary

Directory Structure

```
measurement-free-quantum-classifier/  
  configs/  
    paths.yaml  
  src/  
    run_with_shorts_IQL_QSVM_VQC.py  
    evaluate_capacity_sweep_quantum_vs_knn.py  
    evaluate_all_qmls.py  
    evaluate_iqc_vs_classical.py  
    __init__.py  
    IQL/  
      __init__.py  
      models/  
        fixed_memory_iqc.py  
        static_isdo_model.py  
        adaptive_memory_model.py  
        __init__.py  
      learning/  
        class_state.py  
        update.py  
        metrics.py  
        prototype.py  
        memory_bank.py  
        __init__.py  
      backends/  
        base.py  
        hadamard.py  
        transition.py  
        exact.py  
        hardware_native.py  
        __init__.py  
      regimes/  
        regime3b_responsible.py  
        regime4a_spawn.py  
        regime4b_pruning.py  
        regime3a_wta.py  
        regime2_online.py  
        __init__.py  
      inference/  
        weighted_vote_classifier.py  
        __init__.py  
      encoding/  
        embedding_to_state.py  
        __init__.py  
      baselines/  
        static_isdo_classifier.py  
  scripts/
```

```
test_adaptive_memory_trainer_with_frames.py
frames_to_video.py
iqc/
utils/
    common_backup.py
    common.py
    paths.py
    seed.py
    load_data.py
    label_utils.py
    __init__.py
data/
    pcap_loader.py
    transforms.py
    __init__.py
quantum/
    train_test_vqc_amp_encode.py
    train_test_qsvm_amp_encode.py
    __init__.py
training/
    sweep_k_fixed_memory_iqc_copy.py
    test_fixed_memory_iqc.py
    validate_backends.py
    test_static_isdo_model.py
    sweep_k_fixed_memory_iqc.py
    test_adaptive_memory_trainer.py
    protocol_online/
        train_perceptron.py
    protocol_adaptive_pruning_regime4b/
        test_regime4b_pruning.py
    protocol_adaptive_regime4A/
        test_regime4a.py
    protocol_static/
        evaluate_isdo_k_sweep.py
        plot_isdo_hilbert_geometry.py
        evaluate_static_isdo.py
    classical/
        make_embedding_split.py
        train_embedding_models.py
        extract_embeddings.py
        visualize_embeddings.py
        train_cnn.py
        verify_embeddings.py
        visualize_pcam.py
    protocol_fixed_regime3b_responsible/
        test_regime3b_responsible.py
classical/
    cnn.py
    __init__.py
```

File: configs/paths.yaml

```

base_root: "/home/tarakesh/Work/Repo/measurement-free-quantum-classifier"

paths:
  dataset: "dataset"
  checkpoints: "results/checkpoints"
  embeddings: "results/embeddings"
  figures: "results/figures"
  logs: "results/logs"
  class_prototypes: "results/embeddings/class_prototypes"
  artifacts: "results/artifacts"
  frames_adaptive: "results/frames/adaptive"
  video_adaptive: "results/video/adaptive.mp4"
  frames_fixed: "results/frames/fixed"
  video_fixed: "results/video/fixed.mp4"

class_count:
  K: 3
  K_values: [1, 2, 3, 5, 7, 11, 13, 17, 19, 23]

```

File: src/run_with_shorts_IQL_QSVM_VQC.py

```

import os
import json
import numpy as np
import time
import pickle
import joblib
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import Normalizer

# Qiskit components
from qiskit_machine_learning.algorithms import VQC, QSVC
from qiskit_machine_learning.kernels import FidelityQuantumKernel
from qiskit_algorithms.state_fidelities import ComputeUncompute
from qiskit.primitives import BackendSamplerV2
from qiskit_aer import AerSimulator

# Project components
from src.utils.paths import load_paths
from src.utils.seed import set_seed
from src.IQL.models.fixed_memory_iqc import FixedMemoryIQC
from src.IQL.backends.hardware_native import HardwareNativeBackend

def main():
    set_seed()
    BASE_ROOT, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]

    # Shot counts to compare

```

```

SHOT_LIST = [25, 50, 128, 512, 1024, 2048]

# Subsample for faster evaluation (optional, but recommended for shots
> 512)
EVAL_SAMPLES = 20

print("Loading embeddings...")
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_test = X[test_idx][:EVAL_SAMPLES]
y_test = y[test_idx][:EVAL_SAMPLES]

# y_test_pol for IQL (uses -1, +1)
y_test_pol = np.where(y_test == 0, -1, 1) if np.min(y_test) == 0 else
y_test

# L2 Normalization for Amplitude Encoding
normalizer = Normalizer(norm='l2')
X_test_norm = normalizer.fit_transform(X_test)

# Model Paths
IQL_PATH = os.path.join(BASE_ROOT, "results", "fixed_memory_iqc",
"fixed_memory_iqc.pkl")
QSVM_PATH = os.path.join(BASE_ROOT, "results", "qsvm",
"qsvm_amp_model.dill")
VQC_PATH = os.path.join(BASE_ROOT, "results", "vqc_amp_simple",
"vqc_amp_simple.dill")

print(f"Loading Models...")
# Load IQL
iql_model = FixedMemoryIQC.load(IQL_PATH)

# Load QSVM
qsvm_model = QSVC.load(QSVM_PATH)

# Load VQC
vqc_model = VQC.load(VQC_PATH)

results = {
    "shots": SHOT_LIST,
    "iql_acc": [],
    "qsvm_acc": [],
    "vqc_acc": []
}

for shots in SHOT_LIST:
    print(f"\nEvaluating with {shots} shots...")

    # --- 1. IQL Evaluation ---
    print(" Evaluating IQL...")
    # Inject shots into HardwareNativeBackend
    new_backend = HardwareNativeBackend(shots=shots)

```

```

# Update each ClassState in the memory bank (deep update)
for cs in iql_model.memory_bank.class_states:
    cs.backend = new_backend

y_pred = iql_model.predict(X_test_norm)
acc = accuracy_score(y_test_pol, y_pred)
results["iql_acc"].append(acc)
print(f" IQL Accuracy: {acc:.4f}")

# --- 2. VQC Evaluation ---
print(" Evaluating VQC...")
vqc_model.sampler =
BackendSamplerV2(backend=AerSimulator(shots=shots))
y_pred = vqc_model.predict(X_test_norm)
acc = accuracy_score(y_test, y_pred)
results["vqc_acc"].append(acc)
print(f" VQC Accuracy: {acc:.4f}")

print(" Evaluating QSVM...")
# Since FidelityQuantumKernel with RawFeatureVector fails on
inverse()
# when parameters are unbound, and AerSimulator crashes on large
batches,
# we compute the probabilities using Statevector and then sample.

from qiskit.quantum_info import Statevector

support_vectors = qsvm_model._BaseLibSVM__Xfit[qsvm_model.support_]
n_test = len(X_test_norm)
n_support = len(support_vectors)

# Pre-compute support statevectors
print(f" Pre-computing {n_support} support statevectors...")
sv_support = [Statevector(sj) for sj in support_vectors]

# Build K matrix
K = np.zeros((n_test, n_support))
print(f" Computing {n_test * n_support} overlaps and sampling
{shots} shots...")

for i in range(n_test):
    sv_test = Statevector(X_test_norm[i])
    for j in range(n_support):
        prob = np.abs(sv_test.data @ sv_support[j].data.conj())**2
        # Sample from binomial distribution to simulate shots
        n0 = np.random.binomial(shots, prob)
        K[i, j] = n0 / shots

# Decision function:  $f(x) = \sum \alpha_i y_i K(x, s_i) + b$ 
decision = K @ qsvm_model.dual_coef_.T + qsvm_model.intercept_
y_pred = np.where(decision >= 0, 1, 0).flatten()

acc = accuracy_score(y_test, y_pred)
results["qsvm_acc"].append(acc)

```

```

        print(f"  QSVM Accuracy: {acc:.4f}")

    # Save results
    RESULTS_FILE = os.path.join(BASE_ROOT, "results",
    "shots_comparison_results.json")
    with open(RESULTS_FILE, "w") as f:
        json.dump(results, f, indent=2)
    print(f"\nResults saved to {RESULTS_FILE}")

    # Plotting
    plt.figure(figsize=(10, 6))
    plt.plot(SHOT_LIST, results["iq_l_acc"], 'o-', label='IQL
(FixedMemoryIQC)', linewidth=2)
    plt.plot(SHOT_LIST, results["qsvm_acc"], 's-', label='QSVM
(Amplitude)', linewidth=2)
    plt.plot(SHOT_LIST, results["vqc_acc"], '^-', label='VQC (Amplitude)',
linewidth=2)

    plt.xscale('log')
    plt.xticks(SHOT_LIST, [str(s) for s in SHOT_LIST])
    plt.xlabel('Shot Count (log scale)')
    plt.ylabel('Test Accuracy')
    plt.title(f'Quantum Model Performance vs. Shot Count (N=
{EVAL_SAMPLES})')
    plt.legend()
    plt.grid(True, linestyle='--', alpha=0.7)

    PLOT_FILE = os.path.join(BASE_ROOT, "results", "shots_comparison.png")
    plt.savefig(PLOT_FILE)
    print(f"Plot saved to {PLOT_FILE}")

    # plt.show() # Commented out for non-interactive run

if __name__ == "__main__":
    main()

```

File: src/evaluate_capacity_sweep_quantum_vs_knn.py

```

import os
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier

from src.utils.paths import load_paths
from src.utils.label_utils import ensure_polar, ensure_binary
from src.utils.load_data import load_data

# Quantum models
from src.IQL.models.fixed_memory_iqc import FixedMemoryIQC

```

```

from src.IQL.models.adaptive_memory_model import AdaptiveMemoryModel
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime4a_spawn import Regime4ASpawn
from src.IQL.regimes.regime4b_pruning import Regime4BPruning

def eval_fixed_iqc_k_sweep(Xtr, Xte, ytr_pol, yte_pol, K_values):
    accs = []
    for K in K_values:
        model = FixedMemoryIQC(K=K, eta=0.1)
        model.fit(Xtr, ytr_pol)
        y_pred = model.predict(Xte)
        acc = accuracy_score(yte_pol, y_pred)
        accs.append(acc)
        print(f"Fixed IQC | K={K:<2} | Acc={acc:.4f}")
    return accs

def eval_adaptive_initial_k_sweep(Xtr, Xte, ytr_pol, yte_pol, K_values):
    accs, final_sizes = [], []

    for K in K_values:
        backend = ExactBackend()
        class_states = []

        for cls in [-1, +1]:
            idxs = np.where(ytr_pol == cls)[0][:K]
            for idx in idxs:
                chi = Xtr[idx].astype(np.complex128)
                chi /= np.linalg.norm(chi)
                class_states.append(
                    ClassState(chi, label=cls, backend=backend)
                )

        memory_bank = MemoryBank(class_states)

        learner = Regime4ASpawn(
            memory_bank=memory_bank,
            eta=0.1,
            backend=backend,
            delta_cover=0.2,
            spawn_cooldown=100,
            min_polarized_per_class=2,
        )

        pruner = Regime4BPruning(
            memory_bank=memory_bank,
            tau_harm=-0.15,
            min_age=100,
            min_per_class=1,
            prune_interval=150,

```

```

    )

    model = AdaptiveMemoryModel(
        memory_bank=memory_bank,
        learner=learner,
        pruner=pruner,
        tau_responsible=0.1,
        beta=0.98,
    )

    model.fit(Xtr, ytr_pol)
    model.consolidate(Xtr, ytr_pol, epochs=5, eta_scale=0.4)

    y_pred = model.predict(Xte)
    acc = accuracy_score(yte_pol, y_pred)

    accs.append(acc)
    final_sizes.append(len(memory_bank.class_states))

    print(
        f"Adaptive IQC | init K={K:<2} | "
        f"final mem={final_sizes[-1]:<2} | Acc={acc:.4f}"
    )

    return accs, final_sizes

def eval_knn_k_sweep(Xtr, Xte, ytr_bin, yte_bin, k_values):
    accs = []
    for k in k_values:
        clf = KNeighborsClassifier(n_neighbors=k)
        clf.fit(Xtr, ytr_bin)
        y_pred = clf.predict(Xte)
        acc = accuracy_score(yte_bin, y_pred)
        accs.append(acc)
        print(f"k-NN | k={k:<2} | Acc={acc:.4f}")
    return accs

def main():
    Xtr, Xte, ytr_bin, yte_bin, ytr_pol, yte_pol = load_data("all")

    K = [i for i in range(1, 20)]

    print("\n=== FixedMemory IQC sweep ===")
    fixed_acc = eval_fixed_iqc_k_sweep(Xtr, Xte, ytr_pol, yte_pol, K)

    print("\n=== Adaptive IQC sweep ===")
    adapt_acc, adapt_sizes = eval_adaptive_initial_k_sweep(
        Xtr, Xte, ytr_pol, yte_pol, K
    )

    print("\n=== k-NN sweep ===")
    knn_acc = eval_knn_k_sweep(Xtr, Xte, ytr_bin, yte_bin, K)

```



```

# ----- Plot -----
plt.figure(figsize=(7, 5))
plt.plot(K, fixed_acc, marker="o", label="FixedMemory IQC")
plt.plot(K, adapt_acc, marker="s", label="Adaptive IQC")
plt.plot(K, knn_acc, marker="^", label="k-NN")

plt.xlabel("Capacity parameter (K or k)")
plt.ylabel("Test accuracy")
plt.title("Accuracy vs Capacity: Quantum IQC vs k-NN")
plt.legend()
plt.grid(True)

_, PATHS = load_paths()
out = os.path.join(PATHS["figures"],
"capacity_sweep_quantum_vs_knn.png")
plt.tight_layout()
plt.savefig(out)
plt.close()

print(f"\n👉 Plot saved to: {out}")

if __name__ == "__main__":
    main()

## output

"""

=== FixedMemory IQC sweep ===
Fixed IQC | K=1   | Acc=0.9000
Fixed IQC | K=2   | Acc=0.8927
Fixed IQC | K=3   | Acc=0.8827
Fixed IQC | K=4   | Acc=0.8947
Fixed IQC | K=5   | Acc=0.8927
Fixed IQC | K=6   | Acc=0.8913
Fixed IQC | K=7   | Acc=0.8900
Fixed IQC | K=8   | Acc=0.8847
Fixed IQC | K=9   | Acc=0.8873
Fixed IQC | K=10  | Acc=0.8893
Fixed IQC | K=11  | Acc=0.8913
Fixed IQC | K=12  | Acc=0.8807
Fixed IQC | K=13  | Acc=0.8867
Fixed IQC | K=14  | Acc=0.8893
Fixed IQC | K=15  | Acc=0.8920
Fixed IQC | K=16  | Acc=0.8860
Fixed IQC | K=17  | Acc=0.8887
Fixed IQC | K=18  | Acc=0.8840
Fixed IQC | K=19  | Acc=0.8853

=== Adaptive IQC sweep ===
No frames directory specified. Skipping frame saving.

```

```
Adaptive IQC | init K=1 | final mem=3 | Acc=0.9000  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=2 | final mem=3 | Acc=0.8820  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=3 | final mem=3 | Acc=0.8867  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=4 | final mem=8 | Acc=0.8580  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=5 | final mem=2 | Acc=0.8940  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=6 | final mem=3 | Acc=0.8913  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=7 | final mem=6 | Acc=0.8700  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=8 | final mem=4 | Acc=0.8927  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=9 | final mem=15 | Acc=0.8827  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=10 | final mem=20 | Acc=0.8840  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=11 | final mem=6 | Acc=0.8733  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=12 | final mem=20 | Acc=0.8760  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=13 | final mem=10 | Acc=0.8780  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=14 | final mem=16 | Acc=0.8867  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=15 | final mem=20 | Acc=0.8653  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=16 | final mem=19 | Acc=0.8633  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=17 | final mem=28 | Acc=0.8853  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=18 | final mem=28 | Acc=0.8740  
No frames directory specified. Skipping frame saving.  
Adaptive IQC | init K=19 | final mem=24 | Acc=0.8633
```

=== k-NN sweep ===

```
k-NN | k=1 | Acc=0.9140  
k-NN | k=2 | Acc=0.9187  
k-NN | k=3 | Acc=0.9233  
k-NN | k=4 | Acc=0.9340  
k-NN | k=5 | Acc=0.9260  
k-NN | k=6 | Acc=0.9300  
k-NN | k=7 | Acc=0.9267  
k-NN | k=8 | Acc=0.9307  
k-NN | k=9 | Acc=0.9260  
k-NN | k=10 | Acc=0.9267  
k-NN | k=11 | Acc=0.9267  
k-NN | k=12 | Acc=0.9253  
k-NN | k=13 | Acc=0.9247  
k-NN | k=14 | Acc=0.9253  
k-NN | k=15 | Acc=0.9247
```

```
k-NN | k=16 | Acc=0.9227
k-NN | k=17 | Acc=0.9233
k-NN | k=18 | Acc=0.9233
k-NN | k=19 | Acc=0.9220
```

```
Plot saved to: /home/tarakesh/Work/Repo/measurement-free-quantum-
classifier/results/figures/capacity_sweep_quantum_vs_knn.png
"""
```

File: src/evaluate_all_qmls.py

```
import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.utils.label_utils import ensure_polar, ensure_binary
from src.utils.load_data import load_data
# Models
from src.IQL.models.static_isdo_model import StaticISDOModel
from src.IQL.models.fixed_memory_iqc import FixedMemoryIQC
from src.IQL.models.adaptive_memory_model import AdaptiveMemoryModel

# Adaptive components
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.backends.hardware_native import HardwareNativeBackend
from src.IQL.regimes.regime4a_spawn import Regime4ASpawn
from src.IQL.regimes.regime4b_pruning import Regime4BPruning

def eval_static_isdo(X_train, X_test, y_train_bin, y_test_bin):
    model = StaticISDOModel(K=3)
    model.fit(X_train, y_train_bin)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test_bin, y_pred)
    return acc, "static", 6 # 2*K memories

def eval_fixed_iqc(X_train, X_test, y_train_pol, y_test_pol):
    model = FixedMemoryIQC(K=3, eta=0.1)
    model.fit(X_train, y_train_pol)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test_pol, y_pred)
    mem = len(model.memory_bank.class_states)
    return acc, "fixed", mem

def eval_adaptive_iqc(X_train, X_test, y_train_pol, y_test_pol):
    backend = HardwareNativeBackend()
```

```

# Bootstrap memory (1 per class)
class_states = []
for cls in [-1, +1]:
    idx = np.where(y_train_pol == cls)[0][0]
    chi = X_train[idx].astype(np.complex128)
    chi /= np.linalg.norm(chi)
    class_states.append(ClassState(chi, label=cls, backend=backend))

memory_bank = MemoryBank(class_states)

learner = Regime4ASpawn(
    memory_bank=memory_bank,
    eta=0.1,
    backend=backend,
    delta_cover=0.2,
    spawn_cooldown=100,
    min_polarized_per_class=2,
)

pruner = Regime4BPruning(
    memory_bank=memory_bank,
    tau_harm=-0.15,
    min_age=100,
    min_per_class=1,
    prune_interval=150,
)

model = AdaptiveMemoryModel(
    memory_bank=memory_bank,
    learner=learner,
    pruner=pruner,
    tau_responsible=0.1,
    beta=0.98,
)

# Adaptive phase
model.fit(X_train, y_train_pol)

# Consolidation phase
model.consolidate(X_train, y_train_pol, epochs=5, eta_scale=0.4)

y_pred = model.predict(X_test)
acc = accuracy_score(y_test_pol, y_pred)
mem = len(memory_bank.class_states)

return acc, "adaptive", mem

def main():
    print("\n\U0001F600 Unified Model Evaluation\n")

    Xtr, Xte, ytr_bin, yte_bin, ytr_pol, yte_pol = load_data("all")

```

```

results = []

acc, typ, mem = eval_static_isdo(Xtr, Xte, ytr_bin, yte_bin)
results.append(("Static ISDO", acc, typ, mem))

acc, typ, mem = eval_fixed_iqc(Xtr, Xte, ytr_pol, yte_pol)
results.append(("FixedMemory IQC (K=3)", acc, typ, mem))

acc, typ, mem = eval_adaptive_iqc(Xtr, Xte, ytr_pol, yte_pol)
results.append(("Adaptive IQC (with consolidation)", acc, typ, mem))

print("\n=== Final Comparison ===")
for name, acc, typ, mem in results:
    print(f"{name:35s} | Acc: {acc:.4f} | Type: {typ:8s} | Memory: {mem}")

if __name__ == "__main__":
    main()

"""

└─ Unified Model Evaluation

No frames directory specified. Skipping frame saving.

=== Final Comparison ===
Static ISDO                               | Acc: 0.8807 | Type: static   |
Memory: 6
FixedMemory IQC (K=3)                     | Acc: 0.8827 | Type: fixed    |
Memory: 6
Adaptive IQC (with consolidation)         | Acc: 0.9000 | Type: adaptive |
Memory: 3
"""

```

File: src/evaluate_iqc_vs_classical.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

# Classical models
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier

# Utilities
from src.utils.paths import load_paths
from src.utils.label_utils import ensure_polar, ensure_binary
from src.utils.load_data import load_data

```

```

# Adaptive IQC
from src.IQL.models.adaptive_memory_model import AdaptiveMemoryModel
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime4a_spawn import Regime4ASpawn
from src.IQL.regimes.regime4b_pruning import Regime4BPruning

def eval_classical_models(X_train, X_test, y_train_bin, y_test_bin):
    results = []

    models = {
        "Logistic Regression": LogisticRegression(
            max_iter=500,
            solver="lbfgs"
        ),
        "Linear SVM": SVC(
            kernel="linear"
        ),
        "RBF SVM": SVC(
            kernel="rbf",
            gamma="scale"
        ),
        "k-NN": KNeighborsClassifier(),
    }

    for name, model in models.items():
        model.fit(X_train, y_train_bin)
        y_pred = model.predict(X_test)
        acc = accuracy_score(y_test_bin, y_pred)
        results.append((name, acc))

    return results

def eval_adaptive_iqc(X_train, X_test, y_train_pol, y_test_pol):
    backend = ExactBackend()

    # Bootstrap 1 memory per class
    class_states = []
    for cls in [-1, +1]:
        idx = np.where(y_train_pol == cls)[0][0]
        chi = X_train[idx].astype(np.complex128)
        chi /= np.linalg.norm(chi)
        class_states.append(
            ClassState(chi, label=cls, backend=backend)
        )

    memory_bank = MemoryBank(class_states)

    learner = Regime4ASpawn(
        memory_bank=memory_bank,
        eta=0.1,

```

```

        backend=backend,
        delta_cover=0.2,
        spawn_cooldown=100,
        min_polarized_per_class=2,
    )

    pruner = Regime4BPruning(
        memory_bank=memory_bank,
        tau_harm=-0.15,
        min_age=100,
        min_per_class=1,
        prune_interval=150,
    )

    model = AdaptiveMemoryModel(
        memory_bank=memory_bank,
        learner=learner,
        pruner=pruner,
        tau_responsible=0.1,
        beta=0.98,
    )

    # Adaptive training
    model.fit(X_train, y_train_pol)

    # Consolidation
    model consolidate(
        X_train,
        y_train_pol,
        epochs=5,
        eta_scale=0.4,
    )

    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test_pol, y_pred)
    mem = len(memory_bank.class_states)

    return acc, mem

def main():
    print("\n\U0001F4D Adaptive IQC vs Classical Models\n")

    Xtr, Xte, ytr_bin, yte_bin, ytr_pol, yte_pol = load_data("all")

    # Classical models
    classical_results = eval_classical_models(
        Xtr, Xte, ytr_bin, yte_bin
    )

    # Adaptive IQC
    iqc_acc, iqc_mem = eval_adaptive_iqc(
        Xtr, Xte, ytr_pol, yte_pol
    )

```

```
print("\n=== Classical Models ===")
for name, acc in classical_results:
    print(f"{name:25s} | Acc: {acc:.4f}")

print("\n=== Adaptive Quantum Model ===")
print(
    f"Adaptive IQC (consolidated) | "
    f"Acc: {iqc_acc:.4f} | Memory: {iqc_mem}"
)

if __name__ == "__main__":
    main()
```

File: src/init.py

File: src/IQL/init.py

File: src/IQL/models/fixed_memory_iqc.py

```
# src/IQL/models/fixed_memory_iqc.py

import os
import pickle
import numpy as np

from src.utils.paths import load_paths
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.regimes.regime3a_wta import WinnerTakeAll
from src.IQL.inference.weighted_vote_classifier import
WeightedVoteClassifier
from src.IQL.backends.exact import ExactBackend
from src.IQL.learning.prototype import generate_prototypes, load_prototypes
from src.utils.label_utils import ensure_binary

class FixedMemoryIQC:
    """
    Fixed-Memory Interference Quantum Classifier (IQC)
    """
```


Training pipeline:

1. Generate K prototypes per class (if missing)
 2. Initialize K×2 quantum memory states
 3. Train with Winner-Take-All (Regime-3A)
 4. Freeze memory
- """

```
def __init__(self, K: int, eta: float = 0.1, backend=None, alpha: float
= 0, beta: float = 1):
    self.K = K
    self.eta = eta
    self.backend = backend or ExactBackend()
    self.alpha = alpha
    self.beta = beta

    self.memory_bank = None
    self.trainer = None
    self.classifier = None

def _ensure_prototypes(self, X, y):
    """
    Generate prototypes if they do not already exist.
    """
    _, PATHS = load_paths()
    proto_base = PATHS["class_prototypes"]
    proto_dir = os.path.join(proto_base, f"K{self.K}")

    os.makedirs(proto_dir, exist_ok=True)
    y_binary = ensure_binary(y)
    generate_prototypes(
        X=X,
        y=y_binary,
        K=self.K,
        output_dir=proto_dir
    )
    return load_prototypes(K=self.K, output_dir=proto_dir)

def fit(self, X, y):
    # -----
    # Step 1: ensure prototypes exist
    # -----
    proto = self._ensure_prototypes(X, y)

    # -----
    # Step 2: initialize memory bank
    # -----
    class_states = [
        ClassState(v["vector"], backend=self.backend, label=v["label"])
        for v in proto
    ]
    self.memory_bank = MemoryBank(class_states)

    # -----
```

```
# Step 3: Regime-3A training
# -----
self.trainer = WinnerTakeAll(
    memory_bank=self.memory_bank,
    eta=self.eta,
    backend=self.backend,
    alpha = self.alpha,
    beta = self.beta
)
self.trainer.fit(X, y)

# -----
# Step 4: freeze → inference
# -----
self.classifier = WeightedVoteClassifier(self.memory_bank)
return self

def predict(self, X):
    if self.classifier is None:
        raise RuntimeError("Model not trained. Call fit() first.")
    return [self.classifier.predict(x) for x in X]

def save(self, path):
    """
    Save the model state.
    """
    payload = {
        "K": self.K,
        "eta": self.eta,
        "alpha": self.alpha,
        "beta": self.beta,
        "memory_bank": self.memory_bank,
        "classifier": self.classifier,
        "backend": self.backend
    }
    with open(path, "wb") as f:
        pickle.dump(payload, f)

@classmethod
def load(cls, path):
    """
    Load the model state.
    """
    with open(path, "rb") as f:
        payload = pickle.load(f)

    obj = cls(
        K=payload["K"],
        eta=payload["eta"],
        alpha=payload["alpha"],
        beta=payload["beta"],
        backend=payload.get("backend")
    )
    obj.memory_bank = payload["memory_bank"]
```

```
obj.classifier = payload["classifier"]
return obj
```

File: src/IQL/models/static_isdo_model.py

```
# src/IQL/models/static_isdo_model.py

from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier
from src.utils.paths import load_paths
from src.IQL.learning.prototype import generate_prototypes
import os

class StaticISDOModel:
    """
    Static ISDO Model (Baseline)

    - K prototypes per class
    - No learning
    - Fixed interference reference state  $|\chi\rangle$ 
    """

    def __init__(self, K: int):
        _, PATHS = load_paths()
        self.proto_dir = PATHS["class_prototypes"]
        self.K = K
        self.classifier = None

    def _ensure_prototypes(self, X, y):
        """
        Generate prototypes if they do not already exist.
        """
        _, PATHS = load_paths()
        proto_base = PATHS["class_prototypes"]
        proto_dir = os.path.join(proto_base, f"K{self.K}")
        os.makedirs(proto_dir, exist_ok=True)
        generate_prototypes(
            X=X,
            y=y,
            K=self.K,
            output_dir=proto_dir,
            seed = 42
        )

    def fit(self, X, y):
        """
        Offline preparation only.
        Loads precomputed prototypes and builds classifier.
        """
        self._ensure_prototypes(X, y)
        self.classifier = StaticISDOClassifier(
```

```

        proto_dir=self.proto_dir,
        K=self.K
    )
    return self

def predict(self, X):
    if self.classifier is None:
        raise RuntimeError("Model not fitted. Call fit() first.")
    return self.classifier.predict(X)

```

File: src/IQL/models/adaptive_memory_model.py

```

# src/IQL/training/adaptive_memory_trainer.py
from src.IQL.regimes.regime4a_spawn import Regime4ASpawn
from src.IQL.regimes.regime4b_pruning import Regime4BPruning
from src.IQL.regimes.regime3a_wta import WinnerTakeAll
from src.utils.paths import load_paths
import os

class AdaptiveMemoryModel:
    """
    System-level adaptive controller for IQC.

    Responsibilities:
    - Enforce memory lifecycle invariants
    - Orchestrate learning (Regime-4A)
    - Handle aging, harm tracking, and pruning (Regime-4B)

    This class contains NO learning logic.
    """

    def __init__(
        self,
        memory_bank,
        learner : Regime4ASpawn,          # Regime4ASpawn
        pruner : Regime4BPruning,         # Regime4BPruning
        tau_responsible: float = 0.1,
        beta: float = 0.98,
        frames_dir: str = None,
        fps: int = 50,
    ):
        self.memory_bank = memory_bank
        self.learner = learner
        self.pruner = pruner
        self.tau_responsible = tau_responsible
        self.beta = beta

        self.step_count = 0
        self.history = {
            "action": [],          # spawned | updated | noop

```

```

        "memory_size": [],
        "num_pruned": [],
    }
    self.frames_dir = frames_dir
    self.FRAME_EVERY = fps
def consolidate(
    self,
    X,
    y,
    epochs: int = 5,
    eta_scale: float = 0.3,
):
    """
    Post-adaptive consolidation phase.

    - Freezes memory structure (no spawn, no prune)
    - Refines existing memories using WTA updates
    - Improves margins and accuracy

    Args:
        X: input states
        y: true labels (polar)
        epochs: number of consolidation passes
        eta_scale: scale factor for learning rate
    """

    #print(
    #    f"\n🌀 Consolidation phase started "
    #    f"(epochs={epochs}, eta_scale={eta_scale})"
    #)

    # Winner-Take-All learner (Regime-3A semantics)
    consolidator = WinnerTakeAll(
        memory_bank=self.memory_bank,
        eta=self.learner.eta * eta_scale,
        backend=self.learner.backend,
        alpha=0.0,    # update only on error
        beta=1.0,     # full update
    )

    # IMPORTANT: freeze adaptive structure
    original_spawn = self.learner.step
    original_prune = self.pruner.step

    try:
        # Disable spawn & prune
        self.learner.step = lambda psi, y: "noop"
        self.pruner.step = lambda: []

        for ep in range(epochs):
            consolidator.fit(X, y)
            #print(f" ✓ Consolidation epoch {ep+1}/{epochs}")

    finally:

```

```

        # Restore adaptive behavior
        self.learner.step = original_spawn
        self.pruner.step = original_prune

    #print("👉 Consolidation phase completed\n")
    return self

def step(self, psi, y, frames=False):
    """
    Execute ONE adaptive training step.

    Ordering is STRICT and MUST NOT be changed.
    """

    # -----
    # STEP 1: learning + possible memory spawn
    # -----
    action = self.learner.step(psi, y)

    # -----
    # STEP 2: age update (MANDATORY)
    # -----
    self.memory_bank.increment_age()

    # -----
    # STEP 3: harm EMA update (MANDATORY)
    # -----
    self.memory_bank.update_harm_ema(
        psi,
        y_true=y,
        tau_responsible=self.tau_responsible,
        beta=self.beta,
    )

    # -----
    # STEP 4: pruning (PERIODIC)
    # -----
    pruned = self.pruner.step()
    num_pruned = len(pruned) if pruned else 0

    # -----
    # STEP 5: bookkeeping
    # -----
    self.step_count += 1

    if frames:
        if self.step_count % self.FRAME_EVERY == 0:
            frame_path = f"
{self.frames_dir}/frame_{self.step_count:05d}.png"
            self.memory_bank.visualize(
                qubit=0,
                title="Adaptive IQC – Memory States (Final Snapshot)",
                save_path=frame_path,

```

```
        show=False,
    )

    self.history["action"].append(action)
    self.history["memory_size"].append(
        len(self.memory_bank.class_states)
    )
    self.history["num_pruned"].append(num_pruned)

    return action, pruned

def fit(self, X, y):
    """
    Online adaptive training loop.
    """
    if self.frames_dir is None:
        print("No frames directory specified. Skipping frame saving.")
        frames = False
    else:
        frames = True
        os.makedirs(self.frames_dir, exist_ok=True)
    for psi, label in zip(X, y):
        self.step(psi, label, frames)
    return self

def predict(self, X):
    """
    Inference using winner-take-all interference.
    """
    preds = []
    for psi in X:
        _, score = self.memory_bank.winner(psi)
        preds.append(1 if score >= 0 else -1)
    return preds

def summary(self):
    return {
        "steps": self.step_count,
        "final_memory_size": len(self.memory_bank.class_states),
        "num_spawns": self.history["action"].count("spawned"),
        "num_pruned": sum(self.history["num_pruned"]),
    }
```

File: src/IQL/models/init.py

File: src/IQL/learning/class_state.py

```

import numpy as np
from src.IQL.backends.base import InterferenceBackend

def normalize(v: np.ndarray) -> np.ndarray:
    norm = np.linalg.norm(v)
    if norm == 0:
        raise ValueError("Zero-norm vector cannot be normalized")
    return v / norm

class ClassState:
    """
    Represents the quantum class memory |chi>.
    Invariant: ||chi|| = 1 always.
    """

    def __init__(self, vector: np.ndarray, label: int, backend:
InterferenceBackend = None):
        self.vector = normalize(vector.astype(np.complex128))
        self.backend = backend
        self.label = label
        self.age = 0
        self.harm_ema = 0.0

    def score(self, psi: np.ndarray) -> float:
        """
        ISDO score: Re <chi | psi>
        """
        return self.backend.score(self.vector, psi)

    def update(self, delta: np.ndarray):
        """
        Update |chi> <- normalize(|chi> + delta)
        """
        self.vector = normalize(self.vector + delta)

```

File: src/IQL/learning/update.py

```

import numpy as np
from src.IQL.backends.base import InterferenceBackend

def update(
    chi: np.ndarray,
    psi: np.ndarray,
    y: int,
    eta: float,
    backend: InterferenceBackend,

```



```

):
    """
    Regime-2 update rule (quantum perceptron):

    If  $y * \text{Re}\langle \chi | \psi \rangle \geq 0$ :
        no update
    else:
         $\chi \leftarrow \text{normalize}(\chi + \eta * y * \psi)$ 
    """
    s = backend.score(chi, psi)

    if y * s >= 0:
        return chi, False # correct classification

    delta = eta * y * psi
    chi_new = chi + delta
    chi_new = chi_new / np.linalg.norm(chi_new)

    return chi_new, True

```

File: src/IQL/learning/metrics.py

```

import numpy as np

def summarize_training(history: dict):
    margins = np.array(history["margins"])
    updates = np.array(history["updates"])

    return {
        "mean_margin": float(margins.mean()),
        "min_margin": float(margins.min()),
        "num_updates": int(updates.sum()),
        "update_rate": float(updates.mean()),
    }

```

File: src/IQL/learning/prototype.py

```

import os
import numpy as np
from sklearn.cluster import KMeans

from src.utils.seed import set_seed

# -----
# Helper: quantum-safe normalization
# -----

```

```

def to_quantum_state(x):
    x = np.asarray(x, dtype=np.float64).reshape(-1)
    x = x / np.sqrt(np.sum(x ** 2))
    assert np.isclose(np.sum(x ** 2), 1.0, atol=1e-12)
    return x

# -----
# Core function (IMPORTABLE)
# -----
def generate_prototypes(X, y, K, output_dir, seed=42):
    """
    Generate K prototypes per class using KMeans clustering.
    Prototypes are saved WITH labels.
    """
    set_seed(seed)
    os.makedirs(output_dir, exist_ok=True)

    for cls in [0, 1]:
        X_cls = X[y == cls].astype(np.float64)

        if len(X_cls) < K:
            raise ValueError(
                f"Not enough samples for class {cls}: "
                f"{len(X_cls)} < K={K}"
            )

        kmeans = KMeans(
            n_clusters=K,
            random_state=seed,
            n_init=10
        )
        kmeans.fit(X_cls)

        centers = kmeans.cluster_centers_

        for i in range(K):
            proto = to_quantum_state(centers[i])

            path = os.path.join(
                output_dir, f"class{cls}_proto{i}.npz"
            )

            # ---- SAVE VECTOR + LABEL ----
            np.savez(
                path,
                vector=proto,
                label=cls
            )

def load_prototypes(K, output_dir):
    """
    Load prototypes generated by generate_prototypes.

```

```

Returns:
    List[dict]: each dict has keys { "vector", "label" }
    """
    prototypes = []

    for cls in [0, 1]:
        for i in range(K):
            # New format (.npz)
            npz_path = os.path.join(
                output_dir, f"class{cls}_proto{i}.npz"
            )

            if os.path.exists(npz_path):
                data = np.load(npz_path)
                prototypes.append({
                    "vector": data["vector"],
                    "label": int(data["label"]),
                })
            else:
                # ---- BACKWARD COMPATIBILITY (.npy) ----
                npy_path = os.path.join(
                    output_dir, f"class{cls}_proto{i}.npy"
                )
                vec = np.load(npy_path)
                prototypes.append({
                    "vector": vec,
                    "label": None,
                })

    return prototypes

# -----
# Script mode (EXPERIMENTS ONLY)
# -----
if __name__ == "__main__":
    from src.utils.paths import load_paths

    # Reproducibility
    set_seed(42)

    # Load paths
    _, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]
    PROTO_BASE = PATHS["class_prototypes"]

    os.makedirs(EMBED_DIR, exist_ok=True)
    os.makedirs(PROTO_BASE, exist_ok=True)

    # Load embeddings (TRAIN ONLY)
    X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
    y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))
    train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))

    X_train = X[train_idx]
    y_train = y[train_idx]

```

```

print("Loaded train embeddings:", X_train.shape)

K_VALUES = PATHS["class_count"]["K_values"]

for K in K_VALUES:
    print(f"\n=== Computing prototypes for K={K} ===")
    CLASS_DIR = os.path.join(PROTO_BASE, f"K{K}")
    generate_prototypes(
        X=X_train,
        y=y_train,
        K=K,
        output_dir=CLASS_DIR,
        seed=42
    )
    print(f"Saved prototypes to {CLASS_DIR}")

```

File: src/IQL/learning/memory_bank.py

```

from src.IQL.learning.class_state import ClassState
from qiskit.visualization.bloch import Bloch
import numpy as np
import matplotlib.pyplot as plt

class MemoryBank:
    def __init__(self, class_states):
        self.class_states = class_states

    def scores(self, psi):
        return [
            cs.score(psi)
            for cs in self.class_states
        ]

    def increment_age(self):
        """
        Increment age of all memories by 1.
        Call once per training step.
        """
        for cs in self.class_states:
            cs.age += 1

    def update_harm_ema(self, psi, y_true, tau_responsible, beta):
        """
        Update harm EMA for responsible memories.

        Args:
            psi: input state
            tau_responsible: responsibility threshold
            beta: EMA decay factor
        """

```

```

    """
    scores = self.scores(psi)

    for cs, s in zip(self.class_states, scores):
        if abs(s) > tau_responsible and cs.label is not None:
            harm = -y_true * s
            cs.harm_ema = beta * cs.harm_ema + (1 - beta) * harm

def winner(self, psi):
    scores = self.scores(psi)
    idx = int(max(range(len(scores)), key=lambda i: abs(scores[i])))
    #idx = int(max(range(len(scores)), key=lambda i: scores[i])) ##
causes lower score ??
    return idx, scores[idx]

def add_memory(self, chi_vector, backend, label: int):
    """
    Add a new memory to the bank.

    Args:
        chi_vector: quantum state vector
        backend: interference backend
        label: class label (mandatory)
    """
    self.class_states.append(ClassState(chi_vector, backend=backend,
label=label))

def remove(self, idx):
    """Remove memory at index idx."""
    if 0 <= idx < len(self.class_states):
        del self.class_states[idx]

def prune(self, prune_states):
    """
    Remove given ClassState objects from the memory bank.
    """
    self.class_states = [
        cs for cs in self.class_states
        if cs not in prune_states
    ]

def visualize(
    self,
    qubit: int = 0,
    title: str | None = None,
    save_path: str | None = None,
    show: bool = True,
):
    """
    Visualize the MEMORY-BANK-LEVEL geometry on a single Bloch sphere.

    - Points : individual memory states (projected)
    - Arrows : class centroids
    - Colors : red = class -1 / 0, blue = class +1 / 1

```

```

- STATIC snapshot (no learning, no dynamics)
"""

if not self.class_states:
    raise RuntimeError("MemoryBank is empty")

bloch = Bloch()
bloch.vector_color = []

red_pts, blue_pts = [], []

# --- Pauli matrices ---
X = np.array([[0, 1], [1, 0]], dtype=complex)
Y = np.array([[0, -1j], [1j, 0]], dtype=complex)
Z = np.array([[1, 0], [0, -1]], dtype=complex)
I = np.eye(2, dtype=complex)

def pauli_on_qubit(P, q, n):
    ops = [I] * n
    ops[q] = P
    out = ops[0]
    for op in ops[1:]:
        out = np.kron(out, op)
    return out

# --- Project each memory ---
for cs in self.class_states:
    chi = cs.vector
    n = int(np.log2(len(chi)))
    if 2**n != len(chi):
        raise ValueError("State dimension must be 2^n")

    Xq = pauli_on_qubit(X, qubit, n)
    Yq = pauli_on_qubit(Y, qubit, n)
    Zq = pauli_on_qubit(Z, qubit, n)

    v = np.array([
        float(np.real(np.vdot(chi, Xq @ chi))),
        float(np.real(np.vdot(chi, Yq @ chi))),
        float(np.real(np.vdot(chi, Zq @ chi))),
    ])

    bloch.add_vectors(v)
    bloch.vector_color.append(
        "red" if cs.label in [-1, 0] else "blue"
    )

    if cs.label in [-1, 0]:
        red_pts.append(v)
    else:
        blue_pts.append(v)

# --- Add centroid arrows ---
def add_centroid(vectors, color):

```

```

"""
Add a class centroid as an ARROW on the Bloch sphere.

- vectors : list of Bloch vectors (Nx3)
- color    : color for the centroid arrow
"""

if len(vectors) == 0:
    return

mu = np.mean(vectors, axis=0)
norm = np.linalg.norm(mu)

if norm < 1e-9:
    return

# Keep centroid inside Bloch ball
mu = mu / max(1.0, norm)

# --- Temporarily force arrow rendering ---
previous_style = bloch.vector_style
bloch.vector_style = "arrow"

bloch.add_vectors(mu)
bloch.vector_color.append(color)

# --- Restore previous style (points) ---
bloch.vector_style = previous_style

add_centroid(red_pts, "darkred")
add_centroid(blue_pts, "darkblue")

# --- Title / save / show ---
if title:
    bloch.title = title

if save_path:
    bloch.save(save_path)

if show:
    plt.show()
else:
    plt.close(bloch.fig)

```

File: src/IQL/learning/**init**.py

File: src/IQL/backends/base.py

```

from abc import ABC, abstractmethod

class InterferenceBackend(ABC):
    """
    Abstract interface for computing interference scores.
    """

    @abstractmethod
    def score(self, chi, psi) -> float:
        """
        Return  $\text{Re}\langle \chi | \psi \rangle$  as a real scalar.
        """
        pass

```

File: src/IQL/backends/hadamard.py

```

import numpy as np
from qiskit import QuantumCircuit
from qiskit.quantum_info import Statevector, Pauli
from qiskit.circuit.library import StatePreparation # ✓ Correct import
from .base import InterferenceBackend

# If you also want the conceptual/oracle version:
class HadamardBackend(InterferenceBackend):
    """
    CONCEPTUAL Hadamard-test using oracle state preparation.

    WARNING: This uses non-unitary StatePreparation and is NOT
    physically realizable. Use only for conceptual understanding.
    For actual implementation, use TransitionInterferenceBackend.

    Computes  $\text{Re}\langle \chi | \psi \rangle$  in oracle model.
    """

    def score(self, chi, psi) -> float:
        chi = np.asarray(chi, dtype=np.complex128)
        psi = np.asarray(psi, dtype=np.complex128)

        # Normalize
        chi = chi / np.linalg.norm(chi)
        psi = psi / np.linalg.norm(psi)

        assert chi.shape == psi.shape
        n = int(np.log2(len(psi)))
        assert 2**n == len(psi)

        qc = QuantumCircuit(1 + n)

```



```

anc = 0
data = list(range(1, 1 + n))

# Hadamard on ancilla
qc.h(anc)

# Controlled state preparation (ORACLE ASSUMPTION)
# When anc=0: prepare |psi>
state_prep_psi = StatePreparation(psi)
qc.append(state_prep_psi.control(1), [anc] + data)

# Flip ancilla
qc.x(anc)

# When anc=1 (after flip, so anc=0): prepare |chi>
state_prep_chi = StatePreparation(chi)
qc.append(state_prep_chi.control(1), [anc] + data)

# Flip back
qc.x(anc)

# Final Hadamard
qc.h(anc)

# Get statevector and measure Z on ancilla
sv = Statevector.from_instruction(qc)
z_exp = sv.expectation_value(Pauli('Z'), [anc]).real

return float(z_exp)

```

File: src/IQL/backends/transition.py

```

import numpy as np
from qiskit import QuantumCircuit
from qiskit.quantum_info import Statevector, Pauli
from qiskit.circuit.library import UnitaryGate, StatePreparation # ✓
Correct import
from .base import InterferenceBackend

class TransitionBackend(InterferenceBackend):
    """
    CORRECT physical Hadamard-test using transition unitary.

    This is the physically realizable ISDO implementation.
    Computes  $\text{Re}\langle \chi | \psi \rangle$  using  $U_{\chi\psi} = U_{\chi} @ U_{\psi}^{\dagger}$ 

    This should be used for all hardware experiments and claims.
    """

    @staticmethod

```

```

def _statevector_to_unitary(vec):
    """Build unitary that prepares vec from |0...0>"""
    vec = np.asarray(vec, dtype=np.complex128)
    vec = vec / np.linalg.norm(vec)
    dim = len(vec)

    U = np.zeros((dim, dim), dtype=complex)
    U[:, 0] = vec

    # Gram-Schmidt to complete the unitary
    for i in range(1, dim):
        v = np.zeros(dim, dtype=complex)
        v[i] = 1.0

        for j in range(i):
            v -= np.vdot(U[:, j], v) * U[:, j]

        v_norm = np.linalg.norm(v)
        if v_norm > 1e-10:
            U[:, i] = v / v_norm
        else:
            v = np.random.randn(dim) + 1j * np.random.randn(dim)
            for j in range(i):
                v -= np.vdot(U[:, j], v) * U[:, j]
            U[:, i] = v / np.linalg.norm(v)

    return U

@staticmethod
def _build_transition_unitary(psi, chi):
    """Build  $U_{\chi\psi} = U_{\chi} @ U_{\psi}^{\dagger}$ """
    U_psi = TransitionBackend._statevector_to_unitary(psi)
    U_chi = TransitionBackend._statevector_to_unitary(chi)

    # Transition unitary
    U_chi_psi = U_chi @ U_psi.conj().T

    return UnitaryGate(U_chi_psi)

def score(self, chi, psi) -> float:
    chi = np.asarray(chi, dtype=np.complex128)
    psi = np.asarray(psi, dtype=np.complex128)

    # Normalize
    chi = chi / np.linalg.norm(chi)
    psi = psi / np.linalg.norm(psi)

    assert chi.shape == psi.shape
    n = int(np.log2(len(psi)))
    assert 2**n == len(psi)

    qc = QuantumCircuit(1 + n)
    anc = 0
    data = list(range(1, 1 + n))

```

```

# Prepare  $|\psi\rangle$  on data qubits
qc.append(StatePreparation(psi), data)

# Hadamard on ancilla
qc.h(anc)

# Controlled transition unitary
U_chi_psi = self._build_transition_unitary(psi, chi)
qc.append(U_chi_psi.control(1), [anc] + data)

# Final Hadamard
qc.h(anc)

# Get statevector and measure Z on ancilla
sv = Statevector.from_instruction(qc)
z_exp = sv.expectation_value(Pauli('Z'), [anc]).real

return float(z_exp)

```

File: src/IQL/backends/exact.py

```

import numpy as np
from .base import InterferenceBackend

class ExactBackend(InterferenceBackend):
    """
    Numpy-based interference backend.
    This reproduces existing behavior exactly.
    """

    def score(self, chi, psi) -> float:
        return float(np.real(np.vdot(chi, psi)))

```

File: src/IQL/backends/hardware_native.py

```

import numpy as np
from qiskit import QuantumCircuit, transpile
from qiskit.circuit.library import StatePreparation
from qiskit_aer import AerSimulator
import time

class HardwareNativeBackend:
    """
    Hardware-native Hadamard test implementation.
    Computes  $\text{Re}\langle \chi | \psi \rangle$  using controlled state-preparation circuits.
    """

```

```

def __init__(self, backend=None, shots=25):
    self.backend = backend or AerSimulator()
    self.shots = shots

def score(self, chi, psi) -> float:
    chi = np.asarray(chi, dtype=np.complex128)
    psi = np.asarray(psi, dtype=np.complex128)

    chi = chi / np.linalg.norm(chi)
    psi = psi / np.linalg.norm(psi)

    assert chi.shape == psi.shape
    n = int(np.log2(len(psi)))
    assert 2**n == len(psi)

    qc = QuantumCircuit(1 + n, 1)
    anc = 0
    data = list(range(1, 1 + n))

    psi_state = StatePreparation(psi)
    chi_state = StatePreparation(chi)

    # Prepare |psi>
    qc.append(psi_state, data)

    # Hadamard on ancilla
    qc.h(anc)

    # Controlled Uψ†
    qc.append(psi_state.inverse().control(1), [anc] + data)

    # Controlled Uχ
    qc.append(chi_state.control(1), [anc] + data)

    # Final Hadamard
    qc.h(anc)

    # Measure ancilla
    qc.measure(anc, 0)

    # Transpile for backend
    tqc = transpile(qc, self.backend)

    # Execute
    job = self.backend.run(tqc, shots=self.shots)
    counts = job.result().get_counts()

    # Compute expectation value
    n0 = counts.get('0', 0)
    n1 = counts.get('1', 0)

    z_exp = (n0 - n1) / self.shots
    return float(z_exp)

```

File: src/IQL/backends/init.py

File: src/IQL/regimes/regime3b_responsible.py

```
import numpy as np
from src.IQL.learning.update import update
from src.IQL.backends.exact import ExactBackend

class Regime3BResponsible:
    """
    Regime-3B: Responsible-Set Corrective Learning

    - Same as Regime-3A, but:
      instead of updating only the winner,
      update all RESPONSIBLE memories.

    - Direction still comes from y_true
    - Uses existing update() primitive
    - Guard-A: update energy normalized by |responsible set|
    """

    def __init__(
        self,
        memory_bank,
        eta,
        backend=ExactBackend(),
        alpha_correct: float = 0.0,
        alpha_wrong: float = 1.0,
        tau: float = 0.1,  # responsibility threshold
    ):
        self.memory_bank = memory_bank
        self.eta = eta
        self.backend = backend

        self.alpha_correct = alpha_correct
        self.alpha_wrong = alpha_wrong
        self.tau = tau

        self.num_updates = 0

    # -----
    # Core step
    # -----
    def step(self, psi, y_true):
        # Compute all scores
        scores = self.memory_bank.scores(psi)
```

```

# Winner (for prediction only)
idx_star = int(np.argmax(np.abs(scores)))
score_star = scores[idx_star]

# Correctness
misclassified = (y_true * score_star) < 0
alpha = self.alpha_wrong if misclassified else self.alpha_correct

# Prediction
y_hat = 1 if score_star >= 0 else -1

# -----
# Responsible set
# -----
responsible = [
    cs for cs, s in zip(self.memory_bank.class_states, scores)
    if abs(s) > self.tau
]

# Fallback: always update winner at least
if not responsible:
    responsible = [self.memory_bank.class_states[idx_star]]

# Guard-A normalization
scale = self.eta * alpha / len(responsible)

# -----
# Update ALL responsible memories
# -----
for cs in responsible:
    chi_new, updated = update(
        cs.vector,
        psi,
        y_true,      # ← direction = truth (unchanged)
        scale,
        self.backend,
    )
    if updated:
        cs.vector = chi_new
        self.num_updates += 1

return y_hat

# -----
# Training loop
# -----
def fit(self, X, y):
    correct = 0
    for psi, label in zip(X, y):
        y_hat = self.step(psi, label)
        if y_hat == label:
            correct += 1
    return correct / len(X)

```

```

# -----
# Prediction
# -----
def predict_one(self, psi):
    _, score = self.memory_bank.winner(psi)
    return 1 if score >= 0 else -1

def predict(self, X):
    return [self.predict_one(x) for x in X]

# -----
# Summary
# -----
def summary(self):
    return {
        "memory_size": len(self.memory_bank.class_states),
        "num_updates": self.num_updates,
        "tau": self.tau,
        "alpha_correct": self.alpha_correct,
        "alpha_wrong": self.alpha_wrong,
    }

```

File: src/IQL/regimes/regime4a_spawn.py

```

import numpy as np
from collections import defaultdict

from src.IQL.learning.update import update
from src.IQL.backends.exact import ExactBackend

class Regime4ASpawn:
    """
    Regime-4A: Interference-Coverage Adaptive Memory

    Properties:
    - New implementation (does NOT inherit from Regime-3C)
    - Uses EXACT Regime-3A semantics:
      * Winner-Take-All selection
      * Same misclassification rule
      * Same Regime-2 update
    - Adds memory ONLY when:
      * Winner interference is weak (poor coverage)
      * Winner misclassifies the sample
      * Spawn cooldown allows it
    - Early phase: class-polarized spawning
    - Later phase: class-agnostic spawning
    """

```

```

def __init__(
    self,
    memory_bank,
    eta: float = 0.1,
    backend=None,
    delta_cover: float = 0.2,
    spawn_cooldown: int = 100,
    min_polarized_per_class: int = 1,
):
    """
    Args:
        memory_bank (MemoryBank): existing memory bank
        eta (float): learning rate (Regime-2 update)
        backend (InterferenceBackend): scoring backend
        delta_cover (float): minimum |interference| required to avoid
    spawning
        spawn_cooldown (int): minimum steps between spawns
        min_polarized_per_class (int): bootstrap polarized memories per
    class
    """
    self.memory_bank = memory_bank
    self.eta = eta
    self.backend = backend or ExactBackend()

    # Regime-4A parameters
    self.delta_cover = delta_cover
    self.spawn_cooldown = spawn_cooldown
    self.min_polarized_per_class = min_polarized_per_class

    # Internal state
    self.steps_since_spawn = spawn_cooldown
    self.polarized_count = defaultdict(int)

    # Logging / diagnostics
    self.num_spawns = 0
    self.num_updates = 0
    self.history = {
        "action": [], # "spawned" | "updated" | "noop"
        "winner_score": [],
        "memory_size": [],
    }

    # -----
    # Core step
    # -----
    def step(self, psi: np.ndarray, y: int):
        """
        Process a single training example.

        Args:
            psi (np.ndarray): input quantum state |psi>
            y (int): label in {-1, +1}

        Returns:

```



```

        action (str): "spawned", "updated", or "noop"
    """

    # -----
    # 1. Compute interference scores
    # -----
    scores = self.memory_bank.scores(psi)

    if len(scores) == 0:
        raise RuntimeError("MemoryBank is empty – cannot run Regime-
4A.")

    # -----
    # 2. Winner-Take-All (EXACT Regime-3A semantics)
    # -----
    winner_idx = max(range(len(scores)), key=lambda i: abs(scores[i]))
    s_star = scores[winner_idx]

    # -----
    # 3. Coverage + misclassification checks
    # -----
    poor_coverage = abs(s_star) < self.delta_cover
    misclassified = (y * s_star) < 0
    spawn_allowed = self.steps_since_spawn >= self.spawn_cooldown

    # -----
    # 4. Regime-4A: Spawn new memory if needed
    # -----
    if poor_coverage and misclassified and spawn_allowed:
        residual = psi.astype(np.complex128, copy=True)

        # Orthogonalize against existing memory
        for cs in self.memory_bank.class_states:
            proj = np.vdot(cs.vector, psi)
            residual -= proj * cs.vector

        norm = np.linalg.norm(residual)

        if norm > 1e-8:
            residual /= norm

            # -----
            # Polarized → agnostic transition
            # -----
            # Polarized → agnostic transition
            if self.polarized_count[y] < self.min_polarized_per_class:
                chi_new = y * residual
                self.polarized_count[y] += 1
                label = y # ✓ SET LABEL for polarized memories
            else:
                chi_new = residual
                label = y # Agnostic memories now also carry a label
for consistency

```

```

        self.memory_bank.add_memory(chi_new, self.backend,
label=label) # ✓ PASS LABEL
        self.steps_since_spawn = 0
        self.num_spawns += 1

        self.history["action"].append("spawned")
        self.history["winner_score"].append(float(s_star))
        self.history["memory_size"].append(
            len(self.memory_bank.class_states)
        )

        return "spawned"

# -----
# 5. Otherwise: standard Regime-3A update
# -----
cs = self.memory_bank.class_states[winner_idx]

chi_new, updated = update(
    cs.vector, psi, y, self.eta, self.backend
)

if updated:
    cs.vector = chi_new
    self.num_updates += 1

self.steps_since_spawn += 1

self.history["action"].append("updated" if updated else "noop")
self.history["winner_score"].append(float(s_star))
self.history["memory_size"].append(
    len(self.memory_bank.class_states)
)

return "updated" if updated else "noop"

# -----
# Training loop
# -----
def fit(self, X, y):
    """
    Online training over dataset.

    Args:
        X (Iterable[np.ndarray]): input states
        y (Iterable[int]): labels in {-1, +1}
    """
    for psi, label in zip(X, y):
        self.step(psi, label)
    return self

# -----
# Convenience helpers
# -----

```

```

def memory_size(self) -> int:
    return len(self.memory_bank.class_states)

def summary(self) -> dict:
    return {
        "memory_size": self.memory_size(),
        "num_spawns": self.num_spawns,
        "num_updates": self.num_updates,
    }

```

File: src/IQL/regimes/regime4b_pruning.py

```

class Regime4BPruning:
    """
    Regime-4B: Responsible EMA-Based Memory Pruning

    Removes memories that:
    - are old enough
    - are responsible for interference
    - consistently interfere destructively with their own class
    """

    def __init__(
        self,
        memory_bank,
        tau_harm=-0.2,
        min_age=200,
        min_per_class=1,
        prune_interval=200,
    ):
        self.memory_bank = memory_bank
        self.tau_harm = tau_harm
        self.min_age = min_age
        self.min_per_class = min_per_class
        self.prune_interval = prune_interval

        self.step_count = 0
        self.num_pruned = 0

    # -----
    # Called once per training step
    # -----
    def step(self):
        self.step_count += 1

        if self.step_count % self.prune_interval != 0:
            return []

        return self.prune()

```

```

# -----
# Core pruning logic
# -----
def prune(self):
    to_prune = []

    # Count memories per class
    class_counts = {}
    for cs in self.memory_bank.class_states:
        class_counts.setdefault(cs.label, 0)
        class_counts[cs.label] += 1

    # Identify prune candidates
    for cs in self.memory_bank.class_states:
        if cs.age < self.min_age:
            continue

        if cs.harm_ema < self.tau_harm:
            # enforce class floor
            if class_counts.get(cs.label, 0) > self.min_per_class:
                to_prune.append(cs)
                class_counts[cs.label] -= 1

    if to_prune:
        self.memory_bank.prune(to_prune)
        self.num_pruned += len(to_prune)

    return to_prune

# -----
# Diagnostics
# -----
def summary(self):
    return {
        "num_pruned": self.num_pruned,
        "current_memory_size": len(self.memory_bank.class_states),
        "steps": self.step_count,
    }

```

File: src/IQL/regimes/regime3a_wta.py

```

from src.IQL.learning.update import update
from src.IQL.backends.exact import ExactBackend
import pickle
from tqdm import tqdm

class WinnerTakeAll:
    """
    Regime 3-A: Winner-Takes-All IQC ( $\alpha$ -scaled formulation)

```

```

Special case:
    alpha_correct = 0
    alpha_wrong   = 1
reproduces the original Regime-3A exactly.
"""

def __init__(
    self,
    memory_bank,
    eta,
    backend=ExactBackend(),
    alpha: float = 0.0,
    beta: float = 1.0,
):
    self.memory_bank = memory_bank
    self.eta = eta
    self.backend = backend

    # Scaling factors
    self.alpha_correct = alpha
    self.alpha_wrong = beta

    self.num_updates = 0

    self.history = {
        "winner_idx": [],
        "scores": [],
        "updates": [],
        "alpha": [],
    }

def step(self, psi, y):
    # -----
    # Winner selection (unchanged)
    # -----
    idx, score = self.memory_bank.winner(psi)
    cs = self.memory_bank.class_states[idx]

    # -----
    # Correctness check
    # -----
    misclassified = (y * score) < 0

    #  $\alpha$ -scaling (THIS is the only real change)
    alpha = self.alpha_wrong if misclassified else self.alpha_correct

    # -----
    # Scaled update (winner only)
    # -----
    chi_new, updated = update(
        cs.vector,
        psi,
        y,

```

```

        self.eta * alpha,
        self.backend,
    )

    if updated:
        cs.vector = chi_new
        self.num_updates += 1

    # Prediction (unchanged)
    y_hat = 1 if score >= 0 else -1

    # -----
    # Logging
    # -----
    self.history["winner_idx"].append(idx)
    self.history["scores"].append(score)
    self.history["updates"].append(updated)
    self.history["alpha"].append(alpha)

    return y_hat, idx, updated

def fit(self, X, y):
    correct = 0
    pbar = tqdm(zip(X, y), total=len(X), desc="Training")
    for x, label in pbar:
        y_hat, _, _ = self.step(x, label)
        if y_hat == label:
            correct += 1
        pbar.set_postfix(acc=f"{correct/(pbar.n+1):.4f}")
    return correct / len(X)

def predict_one(self, X):
    _, score = self.memory_bank.winner(X)
    return 1 if score >= 0 else -1

def predict(self, X):
    return [self.predict_one(x) for x in X]

def save(self, path):
    """
    Save trained memory bank and history.
    """
    payload = {
        "memory_bank": self.memory_bank,
        "eta": self.eta,
        "num_updates": self.num_updates,
        "history": self.history,
        "backend": self.backend,
        "alpha_correct": self.alpha_correct,
        "alpha_wrong": self.alpha_wrong,
    }

    with open(path, "wb") as f:
        pickle.dump(payload, f)

```

```

@classmethod
def load(cls, path):
    """
    Load a trained Winner-Take-All model.
    """
    with open(path, "rb") as f:
        payload = pickle.load(f)

    obj = cls(
        memory_bank=payload["memory_bank"],
        eta=payload["eta"],
        backend=payload["backend"],
        alpha_correct=payload.get("alpha_correct", 0.0),
        alpha_wrong=payload.get("alpha_wrong", 1.0),
    )

    obj.num_updates = payload["num_updates"]
    obj.history = payload["history"]

    return obj

```

File: src/IQL/regimes/regime2_online.py

```

import numpy as np
from src.IQL.learning.update import update
import pickle

class OnlinePerceptron:
    """
    Online Interference Quantum Classifier (Regime 2)

    Fixed circuit.
    Trainable object: |chi>
    """

    def __init__(self, class_state, eta: float):
        self.class_state = class_state
        self.eta = eta
        # logs
        self.num_updates = 0
        self.history = {
            "scores": [],
            "margins": [],
            "updates": [],
        }

    def step(self, psi: np.ndarray, y: int):
        """
        Process a single training example.

```

```
        """
        s = self.class_state.score(psi)
        margin = y * s
        y_hat = 1 if s >= 0 else -1

        chi_new, updated = update(
            self.class_state.vector, psi, y, self.eta,
            self.class_state.backend
        )

        if updated:
            self.class_state.vector = chi_new
            self.num_updates += 1

        # logging
        self.history["scores"].append(s)
        self.history["margins"].append(margin)
        self.history["updates"].append(updated)

        return y_hat, s, updated

def fit(self, X, y):
    """
    Single-pass online training.
    dataset: iterable of (psi, y)
    """
    correct = 0

    for i in range(len(X)):
        y_hat, _, _ = self.step(X[i], y[i])
        if y_hat == y[i]:
            correct += 1

    accuracy = correct / len(X)
    return accuracy

def predict_one(self, X):
    s = self.class_state.score(X)
    return 1 if s >= 0 else -1

def predict(self, X):
    return [self.predict_one(x) for x in X]

def save(self, path):
    """
    Save trained perceptron state and history.
    """
    payload = {
        "class_state": self.class_state,
        "eta": self.eta,
        "num_updates": self.num_updates,
        "history": self.history,
        "backend": self.class_state.backend,
    }
```



```
        with open(path, "wb") as f:
            pickle.dump(payload, f)

    @classmethod
    def load(cls, path):
        """
        Load a trained perceptron model.
        """
        with open(path, "rb") as f:
            payload = pickle.load(f)

        obj = cls(
            class_state=payload["class_state"],
            eta=payload["eta"],
        )

        # restore training statistics
        obj.num_updates = payload["num_updates"]
        obj.history = payload["history"]

        return obj
```

File: src/IQL/regimes/init.py

File: src/IQL/inference/weighted_vote_classifier.py

```
class WeightedVoteClassifier:
    def __init__(self, memory_bank, weights=None):
        self.memory_bank = memory_bank
        self.M = len(memory_bank.class_states)

        if weights is None:
            self.weights = [1.0 / self.M] * self.M
        else:
            s = sum(weights)
            self.weights = [w / s for w in weights]

    def score(self, psi):
        scores = self.memory_bank.scores(psi)
        return sum(w * s for w, s in zip(self.weights, scores))

    def predict(self, psi):
        return 1 if self.score(psi) >= 0 else -1

    def save(self, path):
        import pickle
```

```

        payload = {
            "memory_bank": self.memory_bank,
            "weights": self.weights,
        }
        with open(path, "wb") as f:
            pickle.dump(payload, f)

    @classmethod
    def load(cls, path):
        import pickle
        with open(path, "rb") as f:
            payload = pickle.load(f)
        obj = cls(payload["memory_bank"], payload["weights"])
        return obj

```

File: src/IQL/inference/**init**.py

File: src/IQL/encoding/embedding_to_state.py

```

import numpy as np

def embedding_to_state(x: np.ndarray) -> np.ndarray:
    """
    Maps a real embedding  $x \in \mathbb{R}^d$  to a quantum state  $|\psi\rangle$ .
    This is a purely geometric normalization.
    """
    x = x.astype(np.complex128)
    norm = np.linalg.norm(x)
    if norm == 0:
        raise ValueError("Zero embedding encountered")
    return x / norm

```

File: src/IQL/encoding/**init**.py

File: src/IQL/baselines/static_isdo_classifier.py

```

import os
import numpy as np

```

```

from tqdm import tqdm
from src.IQL.backends.exact import ExactBackend
from src.IQL.learning.prototype import load_prototypes

class StaticISDOClassifier:
    def __init__(self, proto_dir, K):
        self.proto_dir = proto_dir
        self.K = K
        self.exact = ExactBackend()
        protos = load_prototypes(
            K=K,
            output_dir=os.path.join(proto_dir, f"K{K}")
        )
        # Binary split (ignore labels even if present)
        self.prototypes = {0: [], 1: []}
        for p in protos:
            # class index is encoded in filename order,
            # OR we can rely on p["label"] if present
            cls = p["label"] if p["label"] is not None else None
            self.prototypes[cls].append(p["vector"])
        self.chi = sum(self.prototypes[0]) - sum(self.prototypes[1])
        self.chi /= np.linalg.norm(self.chi)
    def predict_one(self, psi):
        return 1 if self.exact.score(self.chi, psi) < 0 else 0

    def predict(self, X):
        return np.array([self.predict_one(x) for x in tqdm(X, desc="ISDO
Prediction", leave=False)])

```

File: src/scripts/test_adaptive_memory_trainer_with_frames.py

```

# src/training/protocol_adaptive/test_adaptive_memory_trainer.py

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.utils.load_data import load_data

from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank

from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime4a_spawn import Regime4ASpawn
from src.IQL.regimes.regime4b_pruning import Regime4BPruning
from src.IQL.models.adaptive_memory_model import AdaptiveMemoryModel

import matplotlib.pyplot as plt
from collections import Counter

```

```

def main():
    print("\n🚀 Testing AdaptiveMemoryTrainer (V0)\n")
    _ , paths = load_paths()
    # -----
    # Load data
    # -----
    X_train, X_test, y_train, y_test = load_data("polar")

    print(f"Train samples: {len(X_train)}")
    print(f"Test samples : {len(X_test)}")

    # -----
    # Bootstrap initial memory (1 per class)
    # -----
    backend = ExactBackend()
    class_states = []

    for cls in [-1, +1]:
        idx = np.where(y_train == cls)[0][0]
        chi = X_train[idx].astype(np.complex128)
        chi /= np.linalg.norm(chi)
        class_states.append(
            ClassState(chi, label=cls, backend=backend)
        )

    memory_bank = MemoryBank(class_states)
    print("Initial memory size:", len(memory_bank.class_states))

    # -----
    # Regime-4A (spawn)
    # -----
    learner = Regime4ASpawn(
        memory_bank=memory_bank,
        eta=0.1,
        backend=backend,
        delta_cover=0.2,
        spawn_cooldown=100,
        min_polarized_per_class=1,
    )

    # -----
    # Regime-4B (pruning)
    # -----
    pruner = Regime4BPruning(
        memory_bank=memory_bank,
        tau_harm=-0.15,
        min_age=200,
        min_per_class=1,
        prune_interval=200,
    )

    # -----

```

```

# Adaptive trainer
# -----
trainer = AdaptiveMemoryModel(
    memory_bank=memory_bank,
    learner=learner,
    pruner=pruner,
    tau_responsible=0.1,
    beta=0.98,
    frames_dir=paths["frames_adaptive"],
)

# -----
# Train
# -----
trainer.fit(X_train, y_train)

# -----
# Consolidation phase
# -----
trainer consolidate(
    X_train,
    y_train,
    epochs=5,
    eta_scale=0.3,
)

# -----
# Evaluate
# -----
y_pred = trainer.predict(X_test)
acc = accuracy_score(y_test, y_pred)

print("\n=== Adaptive Trainer Summary ===")
print(trainer.summary())

print("\n=== Evaluation ===")
print(f"Test Accuracy      : {acc:.4f}")
print(f"Final Memory Size   : {len(memory_bank.class_states)}")

print("\n✓ AdaptiveMemoryTrainer test completed.\n")

# -----
# Save adaptive diagnostics
# -----
RESULTS_DIR = paths["frames_adaptive"][:-9]
#save_adaptive_plots(trainer, memory_bank, RESULTS_DIR)
memory_bank.visualize(qubit=0, title="Adaptive IQC - Memory States
(Final Snapshot)", save_path=os.path.join(RESULTS_DIR,
"memory_states.png"), show=True,)

def save_adaptive_plots(trainer, memory_bank, out_dir):
    os.makedirs(out_dir, exist_ok=True)

# -----

```

```
# 1. Memory size over time
# -----
plt.figure(figsize=(6, 4))
plt.plot(trainer.history["memory_size"])
plt.xlabel("Training step")
plt.ylabel("Memory size")
plt.title("Adaptive Memory Size Over Time")
plt.grid(True)
plt.tight_layout()
plt.savefig(os.path.join(out_dir, "memory_size_over_time.png"))
plt.close()

# -----
# 2. Action distribution
# -----
action_counts = Counter(trainer.history["action"])

plt.figure(figsize=(5, 4))
plt.bar(action_counts.keys(), action_counts.values())
plt.xlabel("Action type")
plt.ylabel("Count")
plt.title("Adaptive Actions Distribution")
plt.tight_layout()
plt.savefig(os.path.join(out_dir, "action_distribution.png"))
plt.close()

# -----
# 3. Harm EMA distribution
# -----
harm = [cs.harm_ema for cs in memory_bank.class_states]

plt.figure(figsize=(6, 4))
plt.hist(harm, bins=20)
plt.axvline(x=0.0, linestyle="--")
plt.xlabel("Harm EMA")
plt.ylabel("Count")
plt.title("Harm EMA Distribution (Final)")
plt.tight_layout()
plt.savefig(os.path.join(out_dir, "harm_ema_distribution.png"))
plt.close()

# -----
# 4. Memory age distribution
# -----
ages = [cs.age for cs in memory_bank.class_states]

plt.figure(figsize=(6, 4))
plt.hist(ages, bins=15)
plt.xlabel("Memory age")
plt.ylabel("Count")
plt.title("Memory Age Distribution (Final)")
plt.tight_layout()
plt.savefig(os.path.join(out_dir, "memory_age_distribution.png"))
plt.close()
```

```

        print(f"\n\n Adaptive plots saved to: {out_dir}")

if __name__ == "__main__":
    main()

```

File: src/scripts/frames_to_video.py

```

# scripts/frames_to_video.py
import cv2
import glob
import os
from src.utils.paths import load_paths

def frames_to_video(
    frames_dir: str,
    output_path: str,
    fps: int = 15,
    pattern: str = "frame*.png",
):
    """
    Convert saved Bloch-sphere frames into a video.

    Parameters
    -----
    frames_dir : str
        Directory containing frame images.
    output_path : str
        Path to output video (e.g. .mp4).
    fps : int
        Frames per second.
    pattern : str
        Glob pattern for frame files.
    """
    frame_paths = sorted(glob.glob(os.path.join(frames_dir, pattern)))

    if not frame_paths:
        raise RuntimeError("No frames found to convert into video.")

    # Read first frame to get dimensions
    first = cv2.imread(frame_paths[0])
    if first is None:
        raise RuntimeError("Failed to read first frame.")

    height, width, _ = first.shape

    fourcc = cv2.VideoWriter_fourcc(*"mp4v")
    video = cv2.VideoWriter(
        output_path,

```

```

        cv2.CAP_FFMPEG,
        fourcc,
        fps,
        (width, height),
    )

    for path in frame_paths:
        img = cv2.imread(path)
        if img is None:
            raise RuntimeError(f"Failed to read frame: {path}")
        video.write(img)

    video.release()
    print(f"✓ Video written to: {output_path}")

if __name__ == "__main__":
    _ , path = load_paths()
    frames_dir = path["frames_adaptive"]
    output_path = path["video_adaptive"]
    os.makedirs(os.path.dirname(output_path), exist_ok=True)
    frames_to_video(
        frames_dir=frames_dir,
        output_path=output_path,
        fps=15,
    )

```

File: src/utils/common_backup.py

```

import numpy as np
from qiskit import QuantumCircuit
from qiskit.circuit.library import StatePreparation, UnitaryGate

def load_statevector(vec):
    """
    Create a Qiskit StatePreparation gate from a normalized vector.

    NOTE: This is for CONCEPTUAL/ORACLE model only (Circuit A)
    For physical implementation, use build_transition_unitary instead
    """
    vec = np.asarray(vec, dtype=np.complex128)
    norm = np.linalg.norm(vec)
    if not np.isclose(norm, 1.0, atol=1e-12):
        raise ValueError("Statevector must be normalized")
    return StatePreparation(vec)

def statevector_to_unitary(psi):
    """

```



```

    Convert a statevector to a unitary operator that creates it from
     $|0\dots 0\rangle$ 
    Uses Gram-Schmidt to complete the unitary matrix.

    This creates  $U_{\text{psi}}$  such that  $U_{\text{psi}} |0\dots 0\rangle = |\text{psi}\rangle$ 

    Used for building transition unitaries in Circuit B'.
    """
    psi = np.asarray(psi, dtype=np.complex128)
    dim = len(psi)

    # Normalize
    psi = psi / np.linalg.norm(psi)

    # Create unitary matrix where first column is psi
    U = np.zeros((dim, dim), dtype=complex)
    U[:, 0] = psi

    # Complete to full unitary using Gram-Schmidt orthogonalization
    for i in range(1, dim):
        # Start with standard basis vector
        v = np.zeros(dim, dtype=complex)
        v[i] = 1.0

        # Orthogonalize against all previous columns
        for j in range(i):
            v -= np.vdot(U[:, j], v) * U[:, j]

        # Normalize and store
        v_norm = np.linalg.norm(v)
        if v_norm > 1e-10:
            U[:, i] = v / v_norm
        else:
            # Use random vector if degenerate
            v = np.random.randn(dim) + 1j * np.random.randn(dim)
            for j in range(i):
                v -= np.vdot(U[:, j], v) * U[:, j]
            U[:, i] = v / np.linalg.norm(v)

    return U

def build_transition_unitary(psi, chi):
    """
    Build the transition unitary  $U_{\text{chi\_psi}} = U_{\text{chi}} @ U_{\text{psi}}^{\dagger}$ 

    This is the KEY OPERATION for physically realizable ISDO (Circuit B').

    This unitary satisfies:  $U_{\text{chi\_psi}} |\text{psi}\rangle = |\text{chi}\rangle$ 

    Args:
        psi: Source statevector
        chi: Target statevector

```

```

Returns:
    UnitaryGate that implements the transition
    """
    # Build unitaries that prepare each state from |0...0>
    U_psi = statevector_to_unitary(psi)
    U_chi = statevector_to_unitary(chi)

    # Transition unitary: U_chi @ U_psi^dagger
    U_chi_psi = U_chi @ U_psi.conj().T

    # Verify it works
    psi_normalized = np.asarray(psi, dtype=np.complex128)
    psi_normalized = psi_normalized / np.linalg.norm(psi_normalized)
    chi_normalized = np.asarray(chi, dtype=np.complex128)
    chi_normalized = chi_normalized / np.linalg.norm(chi_normalized)

    result = U_chi_psi @ psi_normalized
    if not np.allclose(result, chi_normalized, atol=1e-10):
        raise ValueError("Transition unitary does not correctly map |psi>
to |chi>")

    return UnitaryGate(U_chi_psi)

def build_chi_state(class0_protos, class1_protos):
    """
    Build |chi> = sum_k |phi_k^0> - sum_k |phi_k^1>, normalized

    This constructs the reference state for ISDO classification.
    """
    chi = np.zeros_like(class0_protos[0], dtype=np.float64)

    for p in class0_protos:
        chi += p
    for p in class1_protos:
        chi -= p

    chi /= np.linalg.norm(chi)
    return chi

```

File: src/utls/common.py

```

import numpy as np
from qiskit import QuantumCircuit
from qiskit.circuit.library import StatePreparation, UnitaryGate

def load_statevector(vec):
    """
    Create a Qiskit StatePreparation gate from a normalized vector.

```

```

NOTE: This is for CONCEPTUAL/ORACLE model only (Circuit A)
For physical implementation, use build_transition_unitary instead
"""

vec = np.asarray(vec, dtype=np.complex128)
norm = np.linalg.norm(vec)
if not np.isclose(norm, 1.0, atol=1e-12):
    raise ValueError("Statevector must be normalized")
return StatePreparation(vec)

def statevector_to_unitary(psi):
    """
    Convert a statevector to a unitary operator using Householder
    efficiency.
    Construct a Householder reflection U such that  $U |e_1\rangle = |\psi\rangle$ 
    where  $e_1 = [1, 0, \dots, 0]^T$ .

    This is  $O(D^2)$  to build the matrix, compared to  $O(D^3)$  for Gram-
    Schmidt.
    """
    psi = np.asarray(psi, dtype=np.complex128)
    norm = np.linalg.norm(psi)
    if norm > 1e-15:
        psi = psi / norm

    dim = len(psi)
    e1 = np.zeros(dim, dtype=np.complex128)
    e1[0] = 1.0

    # Adjust phase to avoid numerical instability (choose phase to make w
    large)
    # We want to map phase * e1 to psi where phase has same angle as psi[0]
    # This ensures  $w = \text{phase} * e_1 - \psi$  is stable.
    angle = np.angle(psi[0]) if np.abs(psi[0]) > 1e-10 else 0.0
    phase = np.exp(1j * angle)

    target = phase * e1
    w = target - psi
    w_norm = np.linalg.norm(w)

    if w_norm < 1e-12:
        # psi is already phase * e1, so just return identity * phase
        return np.eye(dim, dtype=np.complex128) * phase

    v = w / w_norm
    #  $R = I - 2vv^*$  maps target (phase * e1) to psi
    #  $R * \text{phase} * e_1 = \psi \Rightarrow R * e_1 = \psi * \text{phase}^*$ 
    # To get  $U * e_1 = \psi$ , we need  $U = R * \text{phase}$ 
    H = (np.eye(dim, dtype=np.complex128) - 2.0 * np.outer(v, v.conj())) *
    phase
    return H

```

```

def build_transition_unitary(psi, chi):
    """
    Build the transition unitary  $U_{\chi\psi} = U_{\chi} @ U_{\psi}^{\dagger}$ 

    This is the KEY OPERATION for physically realizable ISDO (Circuit B').

    This unitary satisfies:  $U_{\chi\psi} |\psi\rangle = |\chi\rangle$ 

    Args:
        psi: Source statevector
        chi: Target statevector

    Returns:
        UnitaryGate that implements the transition
    """
    # Build unitaries that prepare each state from  $|0\dots0\rangle$ 
    U_psi = statevector_to_unitary(psi)
    U_chi = statevector_to_unitary(chi)

    # Transition unitary:  $U_{\chi} @ U_{\psi}^{\dagger}$ 
    U_chi_psi = U_chi @ U_psi.conj().T

    # Verify it works
    psi_normalized = np.asarray(psi, dtype=np.complex128)
    psi_normalized = psi_normalized / np.linalg.norm(psi_normalized)
    chi_normalized = np.asarray(chi, dtype=np.complex128)
    chi_normalized = chi_normalized / np.linalg.norm(chi_normalized)

    result = U_chi_psi @ psi_normalized
    if not np.allclose(result, chi_normalized, atol=1e-10):
        raise ValueError("Transition unitary does not correctly map  $|\psi\rangle$  to  $|\chi\rangle$ ")

    return UnitaryGate(U_chi_psi)

def build_chi_state(class0_protos, class1_protos):
    """
    Build  $|\chi\rangle = \sum_k |\phi_k^0\rangle - \sum_k |\phi_k^1\rangle$ , normalized

    This constructs the reference state for ISDO classification.
    """
    chi = np.zeros_like(class0_protos[0], dtype=np.float64)

    for p in class0_protos:
        chi += p
    for p in class1_protos:
        chi -= p

    chi /= np.linalg.norm(chi)
    return chi

```

File: src/utils/paths.py

```
import yaml
import os

def load_paths(config_path="configs/paths.yaml"):
    with open(config_path, "r") as f:
        cfg = yaml.safe_load(f)

    base_root = cfg["base_root"]
    paths = {
        k: os.path.join(base_root, v)
        for k, v in cfg["paths"].items()
    }
    paths["class_count"] = cfg["class_count"]
    return base_root, paths
```

File: src/utils/seed.py

```
import random
import numpy as np
import torch
import os

def set_seed(seed: int = 42):
    # Python
    random.seed(seed)

    # NumPy
    np.random.seed(seed)

    # PyTorch
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)

    # cuDNN (important)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

    # Extra safety (hash-based ops)
    os.environ["PYTHONHASHSEED"] = str(seed)
```

File: src/utils/load_data.py

```
import os
import numpy as np
from src.utils.paths import load_paths
from src.utils.seed import set_seed
from src.utils.label_utils import ensure_polar
from src.utils.label_utils import ensure_binary

def load_data(y="all", limit=None):
    set_seed(42)

    _, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]

    X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
    y_bin = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))
    y_pol = np.load(os.path.join(EMBED_DIR, "val_labels_polar.npy"))

    y_bin = ensure_binary(y_bin)
    y_pol = ensure_polar(y_pol)

    train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
    test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

    if limit :
        np.random.shuffle(train_idx)
        np.random.shuffle(test_idx)
        train_idx = train_idx[:limit]
        test_idx = test_idx[:limit]

    X_train, X_test = X[train_idx], X[test_idx]
    y_train_bin, y_test_bin = y_bin[train_idx], y_bin[test_idx]
    y_train_pol, y_test_pol = y_pol[train_idx], y_pol[test_idx]

    # Defensive normalization
    X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
    X_test /= np.linalg.norm(X_test, axis=1, keepdims=True)

    if y == "binary":
        return X_train, X_test, y_train_bin, y_test_bin
    elif y == "polar":
        return X_train, X_test, y_train_pol, y_test_pol
    elif y == "all":
        return X_train, X_test, y_train_bin, y_test_bin, y_train_pol,
        y_test_pol
    else:
        raise ValueError("Invalid value for y")
```

File: src/utils/label_utils.py

```
# src/utils/label_utils.py
"""
Unified label conversion utilities for quantum classifier.

Standard convention:
- Binary: {0, 1} for storage and classical models
- Polar: {-1, +1} for quantum interference calculations
"""
import numpy as np

def binary_to_polar(labels):
    """
    Convert binary labels {0, 1} to polar {-1, +1}.

    Args:
        labels: array-like with values in {0, 1}

    Returns:
        numpy array with values in {-1, +1}
    """
    labels = np.asarray(labels)
    return 2 * labels - 1

def polar_to_binary(labels):
    """
    Convert polar labels {-1, +1} to binary {0, 1}.

    Args:
        labels: array-like with values in {-1, +1}

    Returns:
        numpy array with values in {0, 1}
    """
    labels = np.asarray(labels)
    return (labels + 1) // 2

def ensure_polar(labels):
    """
    Ensure labels are in polar format {-1, +1}.
    Automatically detects format and converts if needed.
    """
    labels = np.asarray(labels)
    unique_vals = np.unique(labels)

    if set(unique_vals).issubset({0, 1}):
        return binary_to_polar(labels)
    elif set(unique_vals).issubset({-1, 1}):
        return labels
    else:
        raise ValueError(f"Labels must be binary {{0,1}} or polar {{-1,+1}}. Got: {unique_vals}")

def ensure_binary(labels):
```

```

"""
Ensure labels are in binary format {0, 1}.
Automatically detects format and converts if needed.
"""
labels = np.asarray(labels)
unique_vals = np.unique(labels)

if set(unique_vals).issubset({0, 1}):
    return labels
elif set(unique_vals).issubset({-1, 1}):
    return polar_to_binary(labels)
else:
    raise ValueError(f"Labels must be binary {{0,1}} or polar
{{-1,+1}}. Got: {unique_vals}")

```

File: src/utils/init.py

File: src/data/pcam_loader.py

```

from torchvision import datasets, transforms
from torch.utils.data import DataLoader

def get_pcam_dataset(data_dir='/home/tarakesh/Work/Repo/measurement-free-
quantum-classifier/dataset', split='train', download=True, transform=None):
    """
    Wrapper for torchvision's built-in PCAM dataset.
    Automatically handles downloading and formatting.
    """
    if transform is None:
        # Default transformation for the hybrid model
        transform = transforms.Compose([
            transforms.ToTensor(), # Scales [0, 255] to [0.0, 1.0] and HWC
to CHW
            ])

    dataset = datasets.PCAM(
        root=data_dir,
        split=split,
        download=download,
        transform=transform
    )
    return dataset

if __name__ == "__main__":
    print("PCAM Loader (using torchvision) initialized.")

```


File: src/data/transforms.py

```
from torchvision import transforms

def get_train_transforms():
    """
    Minimal, label-preserving augmentations for CNN training only.
    """
    return transforms.Compose([
        transforms.RandomHorizontalFlip(),
        transforms.RandomVerticalFlip(),
        transforms.ColorJitter(
            brightness=0.1,
            contrast=0.1,
            saturation=0.05,
        ),
        transforms.ToTensor(),
        transforms.Normalize(
            mean=[0.5, 0.5, 0.5],
            std=[0.5, 0.5, 0.5],
        ),
    ])

def get_eval_transforms():
    """
    Deterministic transforms for validation, testing, and embedding
    extraction.
    """
    return transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize(
            mean=[0.5, 0.5, 0.5],
            std=[0.5, 0.5, 0.5],
        ),
    ])
```

File: src/data/init.py

File: src/quantum/train_test_vqc_amp_encode.py

```

import os
import json
import numpy as np
import time
import joblib
import dill
from sklearn.preprocessing import Normalizer
from sklearn.metrics import accuracy_score

from qiskit import QuantumCircuit
from qiskit.circuit import ParameterVector
from qiskit_machine_learning.algorithms import VQC
from qiskit.circuit.library import StatePreparation
from qiskit_machine_learning.utils.loss_functions import CrossEntropyLoss
from qiskit_algorithms.optimizers import COBYLA
from qiskit.primitives import Sampler # Using reference Sampler for better
compatibility with parameterized initialize

from src.utils.paths import load_paths
from src.utils.seed import set_seed

# =====
# 0. Reproducibility
# =====
set_seed()

# =====
# 1. Load Data
# =====
BASE_ROOT, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]

OUT_DIR = os.path.join(BASE_ROOT, "results", "vqc_amp_simple")
os.makedirs(OUT_DIR, exist_ok=True)

MODEL_PATH = os.path.join(OUT_DIR, "vqc_amp_simple.dill")
RESULT_PATH = os.path.join(OUT_DIR, "vqc_amp_simple_results.json")

print("Loading embeddings...")
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_train = X[train_idx]
y_train = y[train_idx]
X_test = X[test_idx]
y_test = y[test_idx]

# =====

```

```

# 2. Amplitude Normalization
# =====
normalizer = Normalizer(norm='l2')
X_train = normalizer.fit_transform(X_train)
X_test = normalizer.transform(X_test)

dim = X_train.shape[1]
num_qubits = int(np.log2(dim))
assert 2**num_qubits == dim, "Feature dimension must be power of 2"

print(f"Using {num_qubits} qubits (Amplitude Encoding)")

# =====
# 3. Define Feature Map
# =====
# Feature map (Amplitude Encoding)
# We use a ParameterVector to represent the features.
# RawFeatureVector is the standard way, but if it fails, we can try to wrap
it.
from qiskit_machine_learning.circuit.library import RawFeatureVector
feature_map = RawFeatureVector(feature_dimension=dim)

# =====
# 4. Define SIMPLE Ansatz
# =====
theta = ParameterVector("θ", length=num_qubits)
ansatz = QuantumCircuit(num_qubits)

# Single rotation layer
for i in range(num_qubits):
    ansatz.ry(theta[i], i)

# Simple entanglement ring
for i in range(num_qubits - 1):
    ansatz.cx(i, i + 1)
ansatz.cx(num_qubits - 1, 0)

# =====
# 5. Optimizer & Sampler
# =====
optimizer = COBYLA(maxiter=100)
sampler = Sampler()

# =====
# 6. Train or Load
# =====
if os.path.exists(MODEL_PATH) and 1:
    print(f"Loading saved model from {MODEL_PATH}...")
    # NOTE: Qiskit's VQC.load often requires the same environment
    vqc = VQC.load(MODEL_PATH)

```

```

    vqc.sampler = sampler
    train_time = 0.0
else:
    print("Training simple amplitude-VQC...")

    vqc = VQC(
        sampler=sampler,
        feature_map=feature_map,
        ansatz=ansatz,
        optimizer=optimizer,
        loss=CrossEntropyLoss()
    )

    start = time.time()
    vqc.fit(X_train, y_train)
    end = time.time()

    train_time = end - start
    print(f"Training time: {train_time:.2f}s")

    vqc.save(MODEL_PATH)
    print(f"Model saved to {MODEL_PATH}")

# =====
# 7. Evaluate
# =====
print("Evaluating...")
start = time.time()
y_pred = vqc.predict(X_test)
end = time.time()

test_time = end - start
accuracy = accuracy_score(y_test, y_pred)

print("=" * 60)
print(f"Simple Amplitude-VQC Accuracy: {accuracy:.4f}")
print("=" * 60)

# =====
# 8. Save Results
# =====
results = {
    "accuracy": float(accuracy),
    "num_train": len(X_train),
    "num_test": len(X_test),
    "num_features": dim,
    "num_qubits": num_qubits,
    "encoding": "Amplitude Encoding",
    "ansatz": "Single RY layer + CNOT ring",
    "num_parameters": num_qubits,
    "training_time": train_time,
    "test_time": test_time
}

```

```

}

with open(RESULT_PATH, "w") as f:
    json.dump(results, f, indent=2)

print(f"Results saved to {RESULT_PATH}")

"""
Loading embeddings...
Using 5 qubits (Amplitude Encoding)
/home/tarakesh/Work/Repo/measurement-free-quantum-
classifier/src/quantum/train_test_vqc_amp_encode.py:96: DeprecationWarning:
The class ``qiskit.primitives.sampler.Sampler`` is deprecated as of qiskit
1.2. It will be removed no earlier than 3 months after the release date.
All implementations of the `BaseSamplerV1` interface have been deprecated
in favor of their V2 counterparts. The V2 alternative for the `Sampler`
class is `StatevectorSampler`.
    sampler = Sampler()
Training simple amplitude-VQC...
Training time: 650.60s
Model saved to /home/tarakesh/Work/Repo/measurement-free-quantum-
classifier/results/vqc_amp_simple/vqc_amp_simple.dill
Evaluating...
=====
Simple Amplitude-VQC Accuracy: 0.8187
=====
Results saved to /home/tarakesh/Work/Repo/measurement-free-quantum-
classifier/results/vqc_amp_simple/vqc_amp_simple_results.json
"""

```

File: src/quantum/train_test_qsvm_amp_encode.py

```

import os
import json
import numpy as np
import time

from sklearn.metrics import accuracy_score
from sklearn.preprocessing import Normalizer

from qiskit_machine_learning.algorithms import QSVC

from src.utils.paths import load_paths
from src.utils.seed import set_seed

# =====
# 0. Reproducibility
# =====
set_seed()

```

```

# =====
# 1. Load paths and embeddings
# =====
BASE_ROOT, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]
OUT_DIR = os.path.join(BASE_ROOT, "results", "qsvm")
os.makedirs(OUT_DIR, exist_ok=True)
MODEL_PATH = os.path.join(OUT_DIR, "qsvm_amp_model.dill")

print("Loading embeddings...")
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

# Subsample for feasibility
TRAIN_SIZE = 100
TEST_SIZE = 50

X_train = X[train_idx][:TRAIN_SIZE]
y_train = y[train_idx][:TRAIN_SIZE]
X_test = X[test_idx][:TEST_SIZE]
y_test = y[test_idx][:TEST_SIZE]

print(f"Original Shape: {X_train.shape}")

# =====
# 2. Preprocessing for Amplitude Encoding
# =====
# Amplitude encoding requires the input vector to be normalized (L2 norm = 1)
print("Normalizing features (L2) for Amplitude Encoding...")
# Using sklearn Normalizer to ensure L2 norm is exactly 1
normalizer = Normalizer(norm='l2')
X_train_norm = normalizer.fit_transform(X_train)
X_test_norm = normalizer.transform(X_test)

dim = X_train.shape[1]
num_qubits = int(np.log2(dim))
assert 2**num_qubits == dim, f"Dimension {dim} must be a power of 2 for amplitude encoding (2^n)"

print(f"Using {num_qubits} qubits to encode {dim} features.")

# =====
# 3. Define Amplitude Encoding Feature Map using RawFeatureVector
# =====
from qiskit_machine_learning.circuit.library import RawFeatureVector

# RawFeatureVector implements amplitude encoding and handles parameter binding correctly
feature_map = RawFeatureVector(feature_dimension=dim)

```

```

# =====
# 4. Quantum Kernel Setup
# =====
print("Setting up FidelityStatevectorKernel for Amplitude Encoding...")
from qiskit_machine_learning.kernels import FidelityStatevectorKernel

# FidelityStatevectorKernel calculates  $|\langle \psi(x) | \psi(y) \rangle|^2$  directly using
statevectors.
# It does NOT require circuit inversion, so it works with
RawFeatureVector/Amplitude Encoding.
qkernel = FidelityStatevectorKernel(feature_map=feature_map)

if os.path.exists(MODEL_PATH) and 1:
    print(f"Loading saved model from {MODEL_PATH} (skipping training)...")
    qsvm = joblib.load(MODEL_PATH)
    train_time = 0.0
else:
    print("Training QSVC (Amplitude Encoding)...")
    start_time = time.time()

    qsvm = QSVC(quantum_kernel=qkernel)
    qsvm.fit(X_train_norm, y_train)

    end_time = time.time()
    train_time = end_time - start_time
    print(f"Training time: {train_time:.4f}s")
    print(f"Time per sample: {train_time / len(X_train_norm):.4f}s")

    qsvm.save(MODEL_PATH)
    print(f"Model saved to {MODEL_PATH}")

# =====
# 5. Evaluate and Save
# =====
print("Predicting on test set...")
start_time = time.time()
y_pred = qsvm.predict(X_test_norm)
end_time = time.time()
test_time = end_time - start_time
print(f"Test time: {test_time:.4f}s")
print(f"Time per sample: {test_time / len(X_test_norm):.4f}s")

accuracy = accuracy_score(y_test, y_pred)

print("=" * 60)
print(f"QSVC (Amplitude Encoding) Test Accuracy: {accuracy:.4f}")
print("=" * 60)

# Save Results
results = {
    "accuracy": float(accuracy),
    "num_train": len(X_train),
    "num_test": len(X_test),

```

```

    "num_features": dim,
    "num_qubits": num_qubits,
    "encoding": "Amplitude Encoding",
    "training_time": train_time
}

out_path = os.path.join(OUT_DIR, "qsvm_amp_results.json")
with open(out_path, "w") as f:
    json.dump(results, f, indent=2)

print(f"Saved results to {out_path}")

"""
Loading embeddings...
Original Shape: (3500, 32)
Normalizing features (L2) for Amplitude Encoding...
Using 5 qubits to encode 32 features.
Setting up FidelityStatevectorKernel for Amplitude Encoding...
Training QSVC (Amplitude Encoding)...
Training time: 79.7180s
Time per sample: 0.0228s
Predicting on test set...
Test time: 37.7612s
Time per sample: 0.0252s
=====
QSVC (Amplitude Encoding) Test Accuracy: 0.9093
=====
Saved results to /home/tarakesh/Work/Repo/measurement-free-quantum-
classifier/results/qsvm_final/qsvm_amp_results.json
"""

```

File: src/quantum/init.py

File: src/training/sweep_k_fixed_memory_iqc copy.py

```

import os
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
import time

from src.utils.load_data import load_data
from src.IQL.models.fixed_memory_iqc import FixedMemoryIQC
from src.IQL.backends.hardware_native import HardwareNativeBackend
from src.utils.paths import load_paths
from src.utils.seed import set_seed

```



```

def main():
    # -----
    # Load paths and Set Seed
    # -----
    set_seed()
    BASE_ROOT, PATHS = load_paths()

    OUT_DIR = os.path.join(BASE_ROOT, "results", "fixed_memory_iqc_sweep")
    os.makedirs(OUT_DIR, exist_ok=True)

    MODELS_DIR = os.path.join(OUT_DIR, "models")
    os.makedirs(MODELS_DIR, exist_ok=True)

    # -----
    # Load data
    # -----
    X_train, X_test, y_train, y_test = load_data("polar")

    # Quantum-safe normalization (defensive)
    X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
    X_test /= np.linalg.norm(X_test, axis=1, keepdims=True)

    # -----
    # Sweep Setup
    # -----
    # List of K values to sweep over
    k_values = [i for i in range(1, 20)]
    accuracies = []

    print(f"🚀 Starting K sweep for FixedMemoryIQC: {k_values}")

    for k in k_values:
        print(f"\n--- Training K={k} ---")
        start_time = time.time()

        # Initialize and train model
        model = FixedMemoryIQC(K=k, eta=0.1,
backend=HardwareNativeBackend())
        model.fit(X_train, y_train)

        # Evaluate model
        y_pred = model.predict(X_test)
        acc = accuracy_score(y_test, y_pred)
        accuracies.append(acc)

        train_time = time.time() - start_time
        print(f"✅ K={k} | Test Accuracy: {acc:.4f} | Time:
{train_time:.2f}s")

        # Save model
        model_name = f"fixed_memory_iqc_k{k}.pkl"
        model_path = os.path.join(MODELS_DIR, model_name)
        model.save(model_path)
        print(f"💾 Saved model to {model_path}")

```

```

# -----
# Plotting
# -----
plot_path = os.path.join(OUT_DIR, "k_vs_accuracy.png")

plt.figure(figsize=(10, 6))
plt.plot(k_values, accuracies, marker='o', linestyle='--', color='b',
label='Test Accuracy')
plt.title("FixedMemoryIQC: K vs Accuracy")
plt.xlabel("K (Memory Size)")
plt.ylabel("Test Accuracy")
plt.xticks(k_values)
plt.grid(True, linestyle='--', alpha=0.7)
plt.legend()

plt.savefig(plot_path, dpi=300, bbox_inches='tight')
plt.close()

print(f"\n📊 K vs Accuracy plot saved to: {plot_path}")

# Optional: Save results summary to a text file
summary_path = os.path.join(OUT_DIR, "sweep_results.txt")
with open(summary_path, "w") as f:
    f.write("K\tAccuracy\n")
    for k, acc in zip(k_values, accuracies):
        f.write(f"{k}\t{acc:.4f}\n")
print(f"📝 Sweep results summary saved to: {summary_path}")

if __name__ == "__main__":
    main()

```

File: src/training/test_fixed_memory_iqc.py

```

import numpy as np
import os
from sklearn.metrics import accuracy_score

from src.utils.load_data import load_data
from src.IQL.models.fixed_memory_iqc import FixedMemoryIQC
from src.IQL.backends.hardware_native import HardwareNativeBackend
from src.utils.paths import load_paths
from src.utils.seed import set_seed

def main():
    # -----
    # Load paths
    # -----
    set_seed()
    BASE_ROOT, PATHS = load_paths()

```

```

OUT_DIR = os.path.join(BASE_ROOT, "results", "fixed_memory_iqc")
os.makedirs(OUT_DIR, exist_ok=True)

MODEL_PATH = os.path.join(OUT_DIR, "fixed_memory_iqc.pkl")

# -----
# Load data
# -----
X_train, X_test, y_train, y_test = load_data("polar")

# -----
# Quantum-safe normalization (defensive)
# -----
X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
X_test /= np.linalg.norm(X_test, axis=1, keepdims=True)

# -----
# Train Fixed-Memory IQC
# -----
K = 1
model = FixedMemoryIQC(K=K, eta=0.1, backend=HardwareNativeBackend())#,
alpha=0.3, beta=1.5)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
acc = accuracy_score(y_test, y_pred)

print(f"✔ FixedMemoryIQC | K={K} | Test Accuracy: {acc:.4f}")

# -----
# Save/Load Demonstration
# -----
print(f"\n💾 Saving model to {MODEL_PATH}...")
model.save(MODEL_PATH)

#print(f"📂 Loading model from {MODEL_PATH}...")
#loaded_model = FixedMemoryIQC.load(MODEL_PATH)

#y_pred_loaded = loaded_model.predict(X_test)
#acc_loaded = accuracy_score(y_test, y_pred_loaded)
#print(f"✔ Loaded Model | Test Accuracy: {acc_loaded:.4f}")

#if np.all(y_pred == y_pred_loaded):
#    print("🚀 Success: Predictions match exactly!")
#else:
#    print("✖ Error: Predictions do not match!")

if __name__ == "__main__":
    main()

```

File: src/training/validate_backends.py

```

import numpy as np

from src.IQL.backends.exact import ExactBackend
from src.IQL.backends.hadamard import HadamardBackend
from src.IQL.backends.transition import TransitionBackend
from src.IQL.backends.hardware_native import HardwareNativeBackend

def random_state(n_qubits, seed=None):
    if seed is not None:
        np.random.seed(seed)
    dim = 2 ** n_qubits
    v = np.random.randn(dim) + 1j * np.random.randn(dim)
    return v / np.linalg.norm(v)

def run_backend_tests(n_qubits=3, n_tests=20):
    backends = {
        "Exact": ExactBackend(),
        "Hadamard": HadamardBackend(),
        "Transition": TransitionBackend(),
        "HardwareNative": HardwareNativeBackend(),
    }

    print(f"\nRunning backend tests with {n_qubits} qubits\n")

    # Fix  $\chi$ 
    chi = random_state(n_qubits, seed=42)

    scores = {name: [] for name in backends}

    for i in range(n_tests):
        psi = random_state(n_qubits, seed=100 + i)

        print(f"Test {i + 1}")
        for name, backend in backends.items():
            s = backend.score(chi, psi)
            scores[name].append(s)
            print(f"  {name:10s}: {s:+.6f}")
        print()

    # -----
    # Analysis
    # -----
    print("\n=== Backend Agreement Analysis ===\n")

    exact = np.array(scores["Exact"])

    for name in ["Hadamard", "Transition", "HardwareNative"]:
        diff = np.max(np.abs(exact - np.array(scores[name])))

```

```

    print(f"Max |Exact - {name}| = {diff:.2e}")

    # PrimeB: sign + ordering only
    primeb = np.array(scores["HardwareNative"])

    sign_match = np.mean(np.sign(primeb) == np.sign(exact))
    print(f"\nHardwareNative sign agreement with Exact: {sign_match *
100:.1f}%")

    # Rank correlation (ordering)
    exact_rank = np.argsort(exact)
    primeb_rank = np.argsort(primeb)
    rank_corr = np.corrcoef(exact_rank, primeb_rank)[0, 1]
    print(f"HardwareNative rank correlation with Exact: {rank_corr:.3f}")

if __name__ == "__main__":
    run_backend_tests(n_qubits=3, n_tests=20)

"""
Test 194
Exact      : -0.224492
Hadamard   : -0.224492
Transition: -0.224492
PrimeB     : +0.095676

Test 195
Exact      : -0.028519
Hadamard   : -0.028519
Transition: -0.028519
PrimeB     : -0.423231

Test 196
Exact      : +0.203938
Hadamard   : +0.203938
Transition: +0.203938
PrimeB     : -0.201812

Test 197
Exact      : +0.143895
Hadamard   : +0.143895
Transition: +0.143895
PrimeB     : +0.035991

Test 198
Exact      : -0.111603
Hadamard   : -0.111603
Transition: -0.111603
PrimeB     : -0.143718

Test 199
Exact      : +0.164120
Hadamard   : +0.164120

```

```

    Transition: +0.164120
    PrimeB      : +0.107708

Test 200
    Exact       : +0.145881
    Hadamard    : +0.145881
    Transition: +0.145881
    PrimeB      : -0.250643

=== Backend Agreement Analysis ===

Max |Exact - Hadamard| = 3.22e-15
Max |Exact - Transition| = 4.97e-14

PrimeB sign agreement with Exact: 52.5%
PrimeB rank correlation with Exact: -0.004
"""

```

File: src/training/test_static_isdo_model.py

```

import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.load_data import load_data
from src.IQL.models.static_isdo_model import StaticISDOModel

def main():
    # -----
    # Load data
    # -----
    X_train, X_test, y_train, y_test = load_data("polar")

    # -----
    # Sanity: ensure quantum-safe normalization
    # -----
    X_train = X_train / np.linalg.norm(X_train, axis=1, keepdims=True)
    X_test = X_test / np.linalg.norm(X_test, axis=1, keepdims=True)

    # -----
    # Run Static ISDO Model
    # -----
    K = 4 # best K from sweep
    model = StaticISDOModel(K=K)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    print(f"✓ StaticISDOModel | K={K} | Test Accuracy: {acc:.4f}")

```

```
if __name__ == "__main__":  
    main()
```

File: src/training/sweep_k_fixed_memory_iqc.py

```
import os  
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.metrics import accuracy_score  
import time  
from multiprocessing import Pool, cpu_count  
  
from src.utils.load_data import load_data  
from src.IQL.models.fixed_memory_iqc import FixedMemoryIQC  
from src.IQL.backends.hardware_native import HardwareNativeBackend  
from src.utils.paths import load_paths  
from src.utils.seed import set_seed  
  
def train_and_eval_k(k, X_train, y_train, X_test, y_test, models_dir):  
    """Worker function to train and evaluate a single K value."""  
    print(f"\n--- Starting Training K={k} ---")  
    start_time = time.time()  
  
    # Initialize and train model  
    # Note: HardwareNativeBackend and models usually need to be initialized  
    # inside the process  
    # to avoid pickling issues, especially with quantum simulators/hardware  
    # backends.  
    model = FixedMemoryIQC(K=k, eta=0.1, backend=HardwareNativeBackend())  
    model.fit(X_train, y_train)  
  
    # Evaluate model  
    y_pred = model.predict(X_test)  
    acc = accuracy_score(y_test, y_pred)  
  
    train_time = time.time() - start_time  
    print(f"✓ Finished K={k} | Test Accuracy: {acc:.4f} | Time:  
{train_time:.2f}s")  
  
    # Save model  
    model_name = f"fixed_memory_iqc_k{k}.pkl"  
    model_path = os.path.join(models_dir, model_name)  
    model.save(model_path)  
    print(f"💾 Saved K={k} model to {model_path}")  
  
    return k, acc  
  
def main():  
    # -----  
    # Load paths and Set Seed
```

```

# -----
set_seed()
BASE_ROOT, PATHS = load_paths()

OUT_DIR = os.path.join(BASE_ROOT, "results", "fixed_memory_iqc_sweep")
os.makedirs(OUT_DIR, exist_ok=True)

MODELS_DIR = os.path.join(OUT_DIR, "models")
os.makedirs(MODELS_DIR, exist_ok=True)

# -----
# Load data
# -----
X_train, X_test, y_train, y_test = load_data("polar")

# Quantum-safe normalization (defensive)
X_train /= np.linalg.norm(X_train, axis=1, keepdims=True)
X_test /= np.linalg.norm(X_test, axis=1, keepdims=True)

# -----
# Sweep Setup
# -----
# List of K values to sweep over
k_values = [10]#[i for i in range(1,20)]

# Use 3 workers or max CPU count, whichever is smaller
num_workers = min(3, cpu_count())

print(f"🚀 Starting K sweep for FixedMemoryIQC: {k_values}")
print(f"🏃 Running with {num_workers} parallel workers...")

overall_start_time = time.time()

results = []
# Create arguments for the worker pool
pool_args = [(k, X_train, y_train, X_test, y_test, MODELS_DIR) for k in
k_values]

# Run multiprocessing pool
with Pool(processes=num_workers) as pool:
    # Using starmap to pass multiple arguments to the worker function
    results = pool.starmap(train_and_eval_k, pool_args)

# Sort results by K to ensure consistent plotting
results.sort(key=lambda x: x[0])

# Unpack sorted results
sorted_k_values = [res[0] for res in results]
accuracies = [res[1] for res in results]

total_time = time.time() - overall_start_time
print(f"\n🎉 All K sweeps completed in {total_time:.2f}s
({total_time/60:.2f} mins)")

```



```

# -----
# Plotting
# -----
plot_path = os.path.join(OUT_DIR, "k_vs_accuracy.png")

plt.figure(figsize=(10, 6))
plt.plot(sorted_k_values, accuracies, marker='o', linestyle='--',
color='b', label='Test Accuracy')
plt.title("FixedMemoryIQC: K vs Accuracy")
plt.xlabel("K (Memory Size)")
plt.ylabel("Test Accuracy")
plt.xticks(sorted_k_values)
plt.grid(True, linestyle='--', alpha=0.7)
plt.legend()

plt.savefig(plot_path, dpi=300, bbox_inches='tight')
plt.close()

print(f"\n📊 K vs Accuracy plot saved to: {plot_path}")

# Optional: Save results summary to a text file
summary_path = os.path.join(OUT_DIR, "sweep_results.txt")
with open(summary_path, "w") as f:
    f.write("K\tAccuracy\n")
    for k, acc in zip(sorted_k_values, accuracies):
        f.write(f"{k}\t{acc:.4f}\n")
print(f"📝 Sweep results summary saved to: {summary_path}")

if __name__ == "__main__":
    main()

```

File: src/training/test_adaptive_memory_trainer.py

```

# src/training/protocol_adaptive/test_adaptive_memory_trainer.py

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.utils.load_data import load_data

from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank

from src.IQL.backends.exact import ExactBackend
from src.IQL.backends.hardware_native import HardwareNativeBackend
from src.IQL.regimes.regime4a_spawn import Regime4ASpawn
from src.IQL.regimes.regime4b_pruning import Regime4BPruning
from src.IQL.models.adaptive_memory_model import AdaptiveMemoryModel

```

```

import matplotlib.pyplot as plt
from collections import Counter

def main():
    print("\n🚀 Testing AdaptiveMemoryTrainer (V0)\n")

    # -----
    # Load data
    # -----
    X_train, X_test, y_train, y_test = load_data("polar")

    print(f"Train samples: {len(X_train)}")
    print(f"Test samples : {len(X_test)}")

    # -----
    # Bootstrap initial memory (1 per class)
    # -----
    backend = HardwareNativeBackend()
    class_states = []

    for cls in [-1, +1]:
        idx = np.where(y_train == cls)[0][0]
        chi = X_train[idx].astype(np.complex128)
        chi /= np.linalg.norm(chi)
        class_states.append(
            ClassState(chi, label=cls, backend=backend)
        )

    memory_bank = MemoryBank(class_states)
    print("Initial memory size:", len(memory_bank.class_states))

    # -----
    # Regime-4A (spawn)
    # -----
    learner = Regime4ASpawn(
        memory_bank=memory_bank,
        eta=0.1,
        backend=backend,
        delta_cover=0.2,
        spawn_cooldown=100,
        min_polarized_per_class=1,
    )

    # -----
    # Regime-4B (pruning)
    # -----
    pruner = Regime4BPruning(
        memory_bank=memory_bank,
        tau_harm=-0.15,
        min_age=200,
        min_per_class=1,
        prune_interval=200,

```

```

)

# -----
# Adaptive trainer
# -----
trainer = AdaptiveMemoryModel(
    memory_bank=memory_bank,
    learner=learner,
    pruner=pruner,
    tau_responsible=0.1,
    beta=0.98,
)

# -----
# Train
# -----
trainer.fit(X_train, y_train)

# -----
# Consolidation phase
# -----
trainer consolidate(
    X_train,
    y_train,
    epochs=5,
    eta_scale=0.3,
)

# -----
# Evaluate
# -----
y_pred = trainer.predict(X_test)
acc = accuracy_score(y_test, y_pred)

print("\n=== Adaptive Trainer Summary ===")
print(trainer.summary())

print("\n=== Evaluation ===")
print(f"Test Accuracy      : {acc:.4f}")
print(f"Final Memory Size   : {len(memory_bank.class_states)}")

print("\n✓ AdaptiveMemoryTrainer test completed.\n")

# -----
# Save adaptive diagnostics
# -----
RESULTS_DIR = "results/figures/adaptive"
save_adaptive_plots(trainer, memory_bank, RESULTS_DIR)
#memory_bank.visualize(
#    qubit=0,
#    title="Adaptive IQC - Memory States (Final Snapshot)",
#    save_path=os.path.join(RESULTS_DIR, "memory_states.png"),
#    show=True,
#)

```

```

def save_adaptive_plots(trainer, memory_bank, out_dir):
    os.makedirs(out_dir, exist_ok=True)

    # -----
    # 1. Memory size over time
    # -----
    plt.figure(figsize=(6, 4))
    plt.plot(trainer.history["memory_size"])
    plt.xlabel("Training step")
    plt.ylabel("Memory size")
    plt.title("Adaptive Memory Size Over Time")
    plt.grid(True)
    plt.tight_layout()
    plt.savefig(os.path.join(out_dir, "memory_size_over_time.png"))
    plt.close()

    # -----
    # 2. Action distribution
    # -----
    action_counts = Counter(trainer.history["action"])

    plt.figure(figsize=(5, 4))
    plt.bar(action_counts.keys(), action_counts.values())
    plt.xlabel("Action type")
    plt.ylabel("Count")
    plt.title("Adaptive Actions Distribution")
    plt.tight_layout()
    plt.savefig(os.path.join(out_dir, "action_distribution.png"))
    plt.close()

    # -----
    # 3. Harm EMA distribution
    # -----
    harm = [cs.harm_ema for cs in memory_bank.class_states]

    plt.figure(figsize=(6, 4))
    plt.hist(harm, bins=20)
    plt.axvline(x=0.0, linestyle="--")
    plt.xlabel("Harm EMA")
    plt.ylabel("Count")
    plt.title("Harm EMA Distribution (Final)")
    plt.tight_layout()
    plt.savefig(os.path.join(out_dir, "harm_ema_distribution.png"))
    plt.close()

    # -----
    # 4. Memory age distribution
    # -----
    ages = [cs.age for cs in memory_bank.class_states]

    plt.figure(figsize=(6, 4))
    plt.hist(ages, bins=15)
    plt.xlabel("Memory age")

```

```

plt.ylabel("Count")
plt.title("Memory Age Distribution (Final)")
plt.tight_layout()
plt.savefig(os.path.join(out_dir, "memory_age_distribution.png"))
plt.close()

print(f"\n\Adaptive plots saved to: {out_dir}")

if __name__ == "__main__":
    main()

```

File: src/training/protocol_online/train_perceptron.py

```

import numpy as np
import os

from src.IQL.learning.class_state import ClassState
from src.IQL.encoding.embedding_to_state import embedding_to_state
from src.IQL.regimes.regime2_online import OnlinePerceptron
from src.IQL.learning.metrics import summarize_training
from src.IQL.backends.exact import ExactBackend
from src.utils.paths import load_paths
from src.utils.seed import set_seed
from src.utils.load_data import load_data

# -----
# Reproducibility
# -----
set_seed(42)

def main():
    X_train, X_test, y_train, y_test = load_data("polar")

    chi0 = np.zeros_like(X_train[0])
    for psi, label in zip(X_train[:10], y_train[:10]):
        chi0 += label * psi
    chi0 = chi0 / np.linalg.norm(chi0)

    class_state = ClassState(chi0, backend=ExactBackend(), label=+1)
    trainer = OnlinePerceptron(class_state, eta=0.1)

    acc = trainer.fit(X_train, y_train)
    stats = summarize_training(trainer.history)

    print("Final accuracy:", acc)
    print("Training stats:", stats)

if __name__ == "__main__":

```

```

main()

### output
"""
🌱 Global seed set to 42
Loaded train embeddings: (3500, 32)
Final accuracy: 0.8562857142857143
Training stats: {'mean_margin': 0.14930659062683652, 'min_margin':
-0.7069261085786833, 'num_updates': 503, 'update_rate': 0.1437142857142857}
"""

```

File:
src/training/protocol_adaptive_pruning_regime4b/test_regime4b_pruning.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.utils.load_data import load_data
from src.utils.label_utils import ensure_polar
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime3b_responsible import Regime3BResponsible
from src.IQL.regimes.regime4b_pruning import Regime4BPruning

def main():
    print("\n🚀 Testing Regime-4B (EMA-Based Pruning)\n")

    X_train, X_test, y_train, y_test = load_data("polar")

    # -----
    # Initialize memory with extra capacity
    # -----
    backend = ExactBackend()
    class_states = []

    for cls in [-1, +1]:
        idxs = np.where(y_train == cls)[0][:4] # 4 per class
        for idx in idxs:
            chi = X_train[idx].astype(np.complex128)
            chi /= np.linalg.norm(chi)
            class_states.append(
                ClassState(chi, backend=backend, label=cls)
            )

    memory_bank = MemoryBank(class_states)

```

```

print("Initial memory size:", len(memory_bank.class_states))

# -----
# Regime-3B (learning)
# -----
learner = Regime3BResponsible(
    memory_bank=memory_bank,
    eta=0.1,
    alpha_correct=0.0,
    alpha_wrong=1.0,
    tau=0.1,
)

# -----
# Regime-4B (pruning)
# -----
pruner = Regime4BPruning(
    memory_bank=memory_bank,
    tau_harm=-0.15,
    min_age=200,
    min_per_class=1,
    prune_interval=200,
)

# -----
# Training loop
# -----
for step, (psi, label) in enumerate(zip(X_train, y_train)):
    learner.step(psi, label)

    # update metadata
    memory_bank.increment_age()
    memory_bank.update_harm_ema(
        psi,
        y_true=label,
        tau_responsible=0.1,
        beta=0.98,
    )

    pruned = pruner.step()

    if pruned:
        print(
            f"Step {step}: pruned {len(pruned)} memories "
            f"(current size = {len(memory_bank.class_states)})"
        )

# -----
# Evaluation
# -----
y_pred = [learner.predict_one(x) for x in X_test]
test_acc = accuracy_score(y_test, y_pred)

print("\n=== Evaluation ===")

```

```

print(f"Test Accuracy      : {test_acc:.4f}")
print(f"Final Memory Size : {len(memory_bank.class_states)}")

print("\n=== Pruning Summary ===")
print(pruner.summary())

print("\n✓ Regime-4B pruning test completed.\n")

if __name__ == "__main__":
    main()

```

File: src/training/protocol_adaptive_regime4A/test_regime4a.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.utils.load_data import load_data
from src.utils.label_utils import ensure_polar
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime4a_spawn import Regime4ASpawn
from src.IQL.inference.weighted_vote_classifier import
WeightedVoteClassifier

def main():
    print("\n🚀 Testing Regime-4A (Coverage-Based Adaptive Memory)\n")

    # -----
    # Load data
    # -----
    X_train, X_test, y_train, y_test = load_data("polar")

    print(f"Train samples: {len(X_train)}")
    print(f"Test samples : {len(X_test)}")

    # -----
    # Initialize memory bank (bootstrap like Regime-3A)
    # -----
    backend = ExactBackend()

    # Simple bootstrap: one memory per class
    class_states = []

    for cls in [-1, +1]:
        idx = np.where(y_train == cls)[0][0]

```



```

        chi0 = X_train[idx].copy()
        chi0 /= np.linalg.norm(chi0)
        class_states.append(ClassState(chi0, label=cls, backend=backend))

memory_bank = MemoryBank(class_states)

print("Initial memory size:", len(memory_bank.class_states))

# -----
# Train Regime-4A
# -----
model = Regime4ASpawn(
    memory_bank=memory_bank,
    eta=0.1,
    backend=backend,
    delta_cover=0.2,
    spawn_cooldown=100,
    min_polarized_per_class=1,
)

model.fit(X_train, y_train)

print("\n=== Regime-4A Training Summary ===")
print(model.summary())

# -----
# Inference (Regime-3B style)
# -----
classifier = WeightedVoteClassifier(memory_bank)

y_pred = [classifier.predict(x) for x in X_test]
acc = accuracy_score(y_test, y_pred)

print("\n=== Regime-4A Evaluation ===")
print(f"Test Accuracy      : {acc:.4f}")
print(f"Final Memory Size : {len(memory_bank.class_states)}")

# -----
# Sanity checks
# -----
print("\n=== Sanity Checks ===")

actions = model.history["action"]
num_spawned = actions.count("spawned")
num_updated = actions.count("updated")

print(f"Spawn events  : {num_spawned}")
print(f"Update events : {num_updated}")

if num_spawned == 0:
    print("⚠ No memories spawned – try lowering delta_cover")
elif len(memory_bank.class_states) > 50:
    print("⚠ Memory may be growing too fast")
else:

```

```

        print("✓ Memory growth appears controlled")

    print("\n✓ Regime-4A test completed successfully.\n")

if __name__ == "__main__":
    main()

```

File: src/training/protocol_static/evaluate_isdo_k_sweep.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier
from src.utils.paths import load_paths
import matplotlib.pyplot as plt

BASE_ROOT, PATHS = load_paths()

EMBED_DIR = PATHS["embeddings"]
PROTO_BASE = PATHS["class_prototypes"]

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_test = X[test_idx]
y_test = y[test_idx]

accuracy = []
for K in PATHS["class_count"]["K_values"]:

    clf = StaticISDOClassifier(PROTO_BASE, K)

    y_pred = clf.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    accuracy.append(acc)
    print(f"ISDO | K={K:<2} | Accuracy: {acc:.4f}")

"""
ISDO | K=1 | Accuracy: 0.8827
ISDO | K=2 | Accuracy: 0.8800
ISDO | K=3 | Accuracy: 0.8960 ## best
ISDO | K=5 | Accuracy: 0.8840
ISDO | K=7 | Accuracy: 0.8840
ISDO | K=11 | Accuracy: 0.8820
ISDO | K=13 | Accuracy: 0.8800
ISDO | K=17 | Accuracy: 0.8740

```

```
ISDO | K=19 | Accuracy: 0.8780
ISDO | K=23 | Accuracy: 0.8747
"""
```

```
plt.plot(PATHS["class_count"]["K_values"], accuracy, marker="o")
plt.xlabel("Number of prototypes per class (K)")
plt.ylabel("Test Accuracy")
plt.title("ISDO Accuracy vs Interference Capacity")
plt.grid(True)
plt.savefig(os.path.join(PATHS["figures"], "isdo_k_sweep.png"))
```

File: src/training/protocol_static/plot_isdo_hilbert_geometry.py

```
"""
Standalone visualization for Static ISDO classifier (split diagnostics).

PLOT A:
- Prototypes (class 0 / class 1)
-  $\chi$  (interference state)

PLOT B:
-  $\chi$  (interference state)
- Test input

Both plots:
- One figure
- One Bloch sphere per qubit
- Uses ONLY public members of StaticISDOClassifier
"""

import os
import numpy as np
import matplotlib.pyplot as plt
from qiskit.visualization.bloch import Bloch

from src.utils.paths import load_paths
from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier

# =====
# Pauli utilities (identical to MemoryBank)
# =====
X = np.array([[0, 1], [1, 0]], dtype=complex)
Y = np.array([[0, -1j], [1j, 0]], dtype=complex)
Z = np.array([[1, 0], [0, -1]], dtype=complex)
I = np.eye(2, dtype=complex)

def pauli_on_qubit(P, q, n):
    ops = [I] * n
```

```

ops[q] = P
out = ops[0]
for op in ops[1:]:
    out = np.kron(out, op)
return out

def bloch_projection(state, qubit: int):
    """
    Safe projection of a single n-qubit pure state
    onto Bloch coordinates of one qubit.
    """
    # Define Pauli matrices locally
    X_local = np.array([[0, 1], [1, 0]], dtype=complex)
    Y_local = np.array([[0, -1j], [1j, 0]], dtype=complex)
    Z_local = np.array([[1, 0], [0, -1]], dtype=complex)
    I_local = np.eye(2, dtype=complex)

    def pauli_on_qubit_local(P, q, n):
        ops = [I_local] * n
        ops[q] = P
        out = ops[0]
        for op in ops[1:]:
            out = np.kron(out, op)
        return out

    state = np.asarray(state, dtype=np.complex128)

    if state.ndim != 1:
        raise ValueError(
            f"bloch_projection expects a single statevector, got shape {state.shape}"
        )

    state /= np.linalg.norm(state)

    n = int(np.log2(len(state)))
    if 2**n != len(state):
        raise ValueError("State dimension must be 2^n")

    Xq = pauli_on_qubit_local(X_local, qubit, n)
    Yq = pauli_on_qubit_local(Y_local, qubit, n)
    Zq = pauli_on_qubit_local(Z_local, qubit, n)

    return np.array([
        float(np.real(np.vdot(state, Xq @ state))),
        float(np.real(np.vdot(state, Yq @ state))),
        float(np.real(np.vdot(state, Zq @ state))),
    ])

# =====
# PLOT A: Prototypes +  $\chi$ 
# =====
def plot_prototypes_and_chi(

```

```

    model,
    title: str,
    save_path: str | None = None,
    show: bool = True,
):
    """
    Visualize:
    - class 0 prototypes
    - class 1 prototypes
    -  $\chi$  (interference state)

    NO test input.
    """

    # infer dimension from  $\chi$ 
    chi = np.asarray(model.chi, dtype=np.complex128)
    chi /= np.linalg.norm(chi)
    n_qubits = int(np.log2(len(chi)))

    fig = plt.figure(figsize=(4 * n_qubits, 4))

    for q in range(n_qubits):
        ax = fig.add_subplot(1, n_qubits, q + 1, projection="3d")
        bloch = Bloch(fig=fig, axes=ax)
        bloch.vector_color = []

        # ---- prototypes ----
        for label, proto_list in model.prototypes.items():
            color = "red" if label in [0, -1] else "blue"

            for proto in proto_list:
                proto = np.asarray(proto, dtype=np.complex128)
                if proto.ndim > 1:
                    proto = proto.reshape(-1)

                if proto.shape != chi.shape:
                    continue

                proto /= np.linalg.norm(proto)
                v = bloch_projection(proto, q)
                bloch.add_vectors(v)
                bloch.vector_color.append(color)

        # ----  $\chi$  arrow ----
        chi_vec = bloch_projection(chi, q)
        prev = bloch.vector_style
        bloch.vector_style = "arrow"
        bloch.add_vectors(chi_vec)
        bloch.vector_color.append("green")
        bloch.vector_style = prev

        ax.set_title(f"Qubit {q}")
        bloch.render()
    fig.suptitle(title, fontsize=14)

```

```

    if save_path:
        plt.savefig(save_path)

    if show:
        plt.show()
    else:
        plt.close(fig)

# =====
# PLOT B:  $\chi$  + Test Input
# =====
def plot_chi_and_test(
    model,
    test_state,
    title: str,
    save_path: str | None = None,
    show: bool = True,
):
    """
    Visualize:
    -  $\chi$  (interference state)
    - test input

    NO prototypes.
    """

    chi = np.asarray(model.chi, dtype=np.complex128)
    chi /= np.linalg.norm(chi)

    test_state = np.asarray(test_state, dtype=np.complex128)
    test_state /= np.linalg.norm(test_state)

    n_qubits = int(np.log2(len(chi)))

    fig = plt.figure(figsize=(4 * n_qubits, 4))

    for q in range(n_qubits):
        ax = fig.add_subplot(1, n_qubits, q + 1, projection="3d")
        bloch = Bloch(fig=fig, axes=ax)
        bloch.vector_color = []

        # ----  $\chi$  arrow ----
        chi_vec = bloch_projection(chi, q)
        prev = bloch.vector_style
        bloch.vector_style = "arrow"
        bloch.add_vectors(chi_vec)
        bloch.vector_color.append("green")
        bloch.vector_style = prev

        # ---- test point ----
        test_vec = bloch_projection(test_state, q)
        bloch.add_vectors(test_vec)

```

```

        bloch.vector_color.append("black")

        ax.set_title(f"Qubit {q}")
        bloch.render()
    fig.suptitle(title, fontsize=14)

    if save_path:
        plt.savefig(save_path)

    if show:
        plt.show()
    else:
        plt.close(fig)

# =====
# MAIN
# =====
if __name__ == "__main__":

    _, PATHS = load_paths()

    EMBED_DIR = PATHS["embeddings"]
    PROTO_DIR = PATHS["class_prototypes"]
    K = int(PATHS["class_count"]["K"])

    X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
    y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))
    test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

    X_test = X[test_idx]
    y_test = y[test_idx]

    test_state_0 = X_test[np.where(y_test == 0)[0][0]]
    test_state_1 = X_test[np.where(y_test == 1)[0][0]]

    # ---- Static ISDO ----
    model = StaticISDOClassifier(PROTO_DIR, K)
    model.predict_one(test_state_0) # populate  $\chi$  + prototypes

    print(model.chi.shape)
    print(test_state_0.shape)
    print(model.prototypes[0][0].shape)
    print(model.prototypes[1][0].shape)

    # ---- Plot A ----
    plot_prototypes_and_chi(
        model,
        title="Static ISDO - Prototypes and  $\chi$ ",
        save_path="static_isdo_prototypes_chi.png",
        show=True,
    )

    # ---- Plot B ----

```

```

    plot_chi_and_test(
        model,
        test_state_0,
        title="Static ISDO -  $\chi$  and Test Input",
        save_path="static_isdo_chi_test.png",
        show=True,
    )

    plot_chi_and_test(
        model,
        test_state_1,
        title="Static ISDO -  $\chi$  and Test Input",
        save_path="static_isdo_chi_test.png",
        show=True,
    )

```

File: src/training/protocol_static/evaluate_static_isdo.py

```

import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.IQL.baselines.static_isdo_classifier import StaticISDOClassifier
from src.utils.paths import load_paths

BASE_ROOT, PATHS = load_paths()

EMBED_DIR = PATHS["embeddings"]
PROTO_DIR = PATHS["class_prototypes"]
K = int(PATHS["class_count"]["K"])

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

X_test = X[test_idx]
y_test = y[test_idx]

clf = StaticISDOClassifier(PROTO_DIR, K)
y_pred = clf.predict(X_test)

acc = accuracy_score(y_test, y_pred)
print(f"ISDO Accuracy (test): {acc:.4f}")

"""
ISDO Accuracy (test): 0.8840
"""

```


File: src/training/classical/make_embedding_split.py

```
import os
import numpy as np
from sklearn.model_selection import train_test_split

from src.utils.paths import load_paths
from src.utils.seed import set_seed
set_seed(42)

BASE_ROOT, PATHS = load_paths()
EMBED_DIR = PATHS["embeddings"]

X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

indices = np.arange(len(y))

train_idx, test_idx = train_test_split(
    indices,
    test_size=0.3,
    random_state=42,
    stratify=y
)

np.save(os.path.join(EMBED_DIR, "split_train_idx.npy"), train_idx)
np.save(os.path.join(EMBED_DIR, "split_test_idx.npy"), test_idx)

print("Saved split:")
print("Train:", len(train_idx))
print("Test :", len(test_idx))
```

File: src/training/classical/train_embedding_models.py

```
import os
import json
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, normalize
from sklearn.metrics import accuracy_score, roc_auc_score

from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.neighbors import KNeighborsClassifier

from src.utils.paths import load_paths
from src.utils.seed import set_seed
set_seed(42)
```

```

# -----
# Load paths
# -----
BASE_ROOT, PATHS = load_paths()

EMBED_DIR = PATHS["embeddings"]
LOG_DIR = PATHS["logs"]
os.makedirs(LOG_DIR, exist_ok=True)

# -----
# Load embeddings
# -----
X = np.load(os.path.join(EMBED_DIR, "val_embeddings.npy"))
y = np.load(os.path.join(EMBED_DIR, "val_labels.npy"))

train_idx = np.load(os.path.join(EMBED_DIR, "split_train_idx.npy"))
test_idx = np.load(os.path.join(EMBED_DIR, "split_test_idx.npy"))

print("Loaded embeddings:", X.shape)

# -----
# Preprocessing (DEPRECATED: Now handled in extract_embeddings.py)
# -----
# # 1) Standardize (important for linear models)
# scaler = StandardScaler()
# X_std = scaler.fit_transform(X)
#
# # 2) L2-normalize (important for similarity & quantum)
# X_l2 = normalize(X_std, norm="l2")

# -----
# Train / test split
# -----

# Using raw pre-normalized float64 embeddings for all models
Xtr = X[train_idx]
Xte = X[test_idx]
ytr = y[train_idx]
yte = y[test_idx]

results = {}

# =====
# 1 Logistic Regression (Linear separability)
# =====
print("\nTraining Logistic Regression...")
logreg = LogisticRegression(
    max_iter=1000,
    n_jobs=-1
)
logreg.fit(Xtr, ytr)

pred_lr = logreg.predict(Xte)

```

```

proba_lr = logreg.predict_proba(Xte)[: , 1]

results["LogisticRegression"] = {
    "accuracy": accuracy_score(yte, pred_lr),
    "auc": roc_auc_score(yte, proba_lr)
}

# =====
# ❷ Linear SVM (Max-margin)
# =====
print("Training Linear SVM...")
svm = LinearSVC()
svm.fit(Xtr, ytr)

pred_svm = svm.predict(Xte)

results["LinearSVM"] = {
    "accuracy": accuracy_score(yte, pred_svm),
    "auc": None # LinearSVC has no probability estimates
}

# =====
# ❸ k-NN (Distance-based similarity)
# =====
print("Training k-NN...")
knn = KNeighborsClassifier(
    n_neighbors=5,
    metric="euclidean"
)
knn.fit(Xtr, ytr)
print("Knn neighbors:", knn.n_neighbors)
pred_knn = knn.predict(Xte)
proba_knn = knn.predict_proba(Xte)[: , 1]

results["kNN"] = {
    "accuracy": accuracy_score(yte, pred_knn),
    "auc": roc_auc_score(yte, proba_knn)
}

# -----
# Save results
# -----
with open(os.path.join(LOG_DIR, "embedding_baseline_results.json"), "w") as f:
    json.dump(results, f, indent=2)

# -----
# Print summary
# -----
print("\n=== Embedding Baseline Results ===")
for model, metrics in results.items():
    print(
        f"{model:>18} | "
        f"Acc: {metrics['accuracy']:.4f} | "

```

```

        f"AUC: {metrics['auc']}"
    )

## output
"""
🌱 Global seed set to 42
Loaded embeddings: (5000, 32)

Training Logistic Regression...
Training Linear SVM...
Training k-NN...
Knn neighbors: 5

=== Embedding Baseline Results ===
LogisticRegression | Acc: 0.9047 | AUC: 0.9664224751066857
      LinearSVM    | Acc: 0.9053 | AUC: None
      kNN          | Acc: 0.9260 | AUC: 0.9711219772403983
"""

```

File: src/training/classical/extract_embeddings.py

```

import os
import torch
import numpy as np
from torch.utils.data import DataLoader, Subset
from tqdm import tqdm

from src.classical.cnn import PCamCNN
from src.data.pcam_loader import get_pcam_dataset
from src.data.transforms import get_eval_transforms
from src.utils.paths import load_paths
from src.utils.seed import set_seed

set_seed(42)

BASE_ROOT, PATHS = load_paths()
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

CHECKPOINT = os.path.join(PATHS["checkpoints"], "pcam_cnn_best.pt")
os.makedirs(PATHS["embeddings"], exist_ok=True)

model = PCamCNN(embedding_dim=32).to(DEVICE)
model.load_state_dict(torch.load(CHECKPOINT, map_location=DEVICE))
model.eval()

dataset = get_pcam_dataset(PATHS["dataset"], "val", get_eval_transforms())
subset = Subset(dataset, range(5000))
loader = DataLoader(subset, batch_size=128, num_workers=6, pin_memory=True)

embeds, labels, lable_polar = [], [], []

```

```

with torch.no_grad():
    for x, y in tqdm(loader):
        z = model(x.to(DEVICE), return_embedding=True)
        # Convert to float64 FIRST, then normalize for maximum precision
        z = z.to(torch.float64)
        z = torch.nn.functional.normalize(z, p=2, dim=1)

        embeds.append(z.cpu().numpy())
        labels.append(y.numpy().astype(np.float64))
        lable_polar.append((y.numpy()*2 - 1).astype(np.float64))

np.save(os.path.join(PATHS["embeddings"], "val_embeddings.npy"),
np.vstack(embeds).astype(np.float64))
np.save(os.path.join(PATHS["embeddings"], "val_labels.npy"),
np.concatenate(labels).astype(np.float64))
np.save(os.path.join(PATHS["embeddings"], "val_labels_polar.npy"),
np.concatenate(lable_polar).astype(np.float64))

```

File: src/training/classical/visualize_embeddings.py

```

import os
import numpy as np
import matplotlib.pyplot as plt
from sklearn.manifold import TSNE
from src.utils.paths import load_paths
from src.utils.seed import set_seed

set_seed(42)

_, PATHS = load_paths()

X = np.load(os.path.join(PATHS["embeddings"], "val_embeddings.npy"))
y = np.load(os.path.join(PATHS["embeddings"], "val_labels.npy"))

tsne = TSNE(n_components=2, perplexity=30, max_iter=1000, random_state=42)
X2 = tsne.fit_transform(X)

plt.figure(figsize=(7, 6))
plt.scatter(X2[y == 0, 0], X2[y == 0, 1], s=8, label="Benign")
plt.scatter(X2[y == 1, 0], X2[y == 1, 1], s=8, label="Malignant")
plt.legend()
plt.savefig(os.path.join(PATHS["figures"], "embedding_tsne.png"), dpi=300)
plt.show()

```

File: src/training/classical/train_cnn.py

```

import os
import torch

```

```

import torch.nn as nn
from torch.utils.data import DataLoader
from tqdm import tqdm
import json
import matplotlib.pyplot as plt

from src.classical.cnn import PCamCNN
from src.data.pcam_loader import get_pcam_dataset
from src.data.transforms import get_train_transforms, get_eval_transforms
from src.utils.paths import load_paths
from src.utils.seed import set_seed

set_seed(42)
#torch.backends.cudnn.benchmark = True

# -----
# Load paths
# -----
BASE_ROOT, PATHS = load_paths()
DATA_ROOT = PATHS["dataset"]

# -----
# Config
# -----
BATCH_SIZE = 64
EPOCHS = 30
LR = 1e-3
EMBEDDING_DIM = 32
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

os.makedirs(PATHS["checkpoints"], exist_ok=True)
os.makedirs(PATHS["logs"], exist_ok=True)
os.makedirs(PATHS["figures"], exist_ok=True)

# -----
# Training / Evaluation loops
# -----
def train_one_epoch(model, loader, criterion, optimizer):
    model.train()
    running_loss, correct, total = 0.0, 0, 0

    for images, labels in tqdm(loader, desc="Training", leave=False):
        images, labels = images.to(DEVICE), labels.to(DEVICE)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * images.size(0)
        correct += outputs.argmax(1).eq(labels).sum().item()
        total += labels.size(0)

    return running_loss / total, correct / total

```

```

@torch.no_grad()
def evaluate(model, loader, criterion):
    model.eval()
    running_loss, correct, total = 0.0, 0, 0

    for images, labels in tqdm(loader, desc="Validation", leave=False):
        images, labels = images.to(DEVICE), labels.to(DEVICE)
        outputs = model(images)
        loss = criterion(outputs, labels)

        running_loss += loss.item() * images.size(0)
        correct += outputs.argmax(1).eq(labels).sum().item()
        total += labels.size(0)

    return running_loss / total, correct / total

def main():
    print(f"🚀 Training on device: {DEVICE}")

    train_set = get_pcam_dataset(DATA_ROOT, "train",
get_train_transforms())
    val_set = get_pcam_dataset(DATA_ROOT, "val", get_eval_transforms())

    train_loader = DataLoader(train_set, BATCH_SIZE, shuffle=True,
num_workers=6, pin_memory=True)
    val_loader = DataLoader(val_set, BATCH_SIZE, shuffle=False,
num_workers=6, pin_memory=True)

    model = PCamCNN(embedding_dim=EMBEDDING_DIM).to(DEVICE)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=LR,
weight_decay=1e-4)
    scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
        optimizer, mode="max", factor=0.5, patience=2
    )

    best_val_acc, patience, wait = 0.0, 10, 0
    history = {k: [] for k in ["train_loss", "train_acc", "val_loss",
"val_acc"]}

    for epoch in range(1, EPOCHS + 1):
        print(f"\n📅 Epoch {epoch}/{EPOCHS}")

        tr_loss, tr_acc = train_one_epoch(model, train_loader, criterion,
optimizer)
        val_loss, val_acc = evaluate(model, val_loader, criterion)
        scheduler.step(val_acc)

        history["train_loss"].append(tr_loss)
        history["train_acc"].append(tr_acc)
        history["val_loss"].append(val_loss)

```

```

        history["val_acc"].append(val_acc)

    print(f"Train Acc {tr_acc:.4f} | Val Acc {val_acc:.4f}")

    if val_acc > best_val_acc:
        best_val_acc = val_acc
        torch.save(model.state_dict(),
os.path.join(PATHS["checkpoints"], "pcam_cnn_best.pt"))
        print("✓ Best validation accuracy reached : Saved checkpoint")
        wait = 0
    else:
        wait += 1

    if wait >= patience:
        print("■ Early stopping")
        break

    torch.save(model.state_dict(), os.path.join(PATHS["checkpoints"],
"pcam_cnn_final.pt"))
    print("✓ Final checkpoint saved")
    # Save logs
    with open(os.path.join(PATHS["logs"], "train_history.json"), "w") as f:
        json.dump(history, f, indent=2)

# Plots
epochs = range(1, len(history["train_loss"]) + 1)
plt.figure()
plt.plot(epochs, history["train_acc"], label="Train")
plt.plot(epochs, history["val_acc"], label="Val")
plt.legend()
plt.savefig(os.path.join(PATHS["figures"], "cnn_accuracy.png"))
plt.close()

plt.figure()
plt.plot(epochs, history["train_loss"], label="Train")
plt.plot(epochs, history["val_loss"], label="Val")
plt.legend()
plt.savefig(os.path.join(PATHS["figures"], "cnn_loss.png"))
plt.close()

if __name__ == "__main__":
    main()

```

File: src/training/classical/verify_embbeings.py

```

import os
import numpy as np
from src.utils.paths import load_paths

```



```
def verify_embeddings():
    BASE_ROOT, PATHS = load_paths()
    EMBED_DIR = PATHS["embeddings"]

    file_path = os.path.join(EMBED_DIR, "val_embeddings.npy")
    if not os.path.exists(file_path):
        print(f"File not found: {file_path}")
        return

    print(f"Verifying: {file_path}")
    X = np.load(file_path)
    print(f"Shape: {X.shape}, Dtype: {X.dtype}")

    # Calculate norm-squared for each sample
    norms_sq = np.sum(X**2, axis=1)

    max_val = np.max(norms_sq)
    min_val = np.min(norms_sq)
    mean_val = np.mean(norms_sq)

    print(f"Max norm squared: {max_val:.15f}")
    print(f"Min norm squared: {min_val:.15f}")
    print(f"Mean norm squared: {mean_val:.15f}")

    # Qiskit usually has a tolerance around 1e-8 or 1e-10
    tolerance = 1e-8
    violations = np.sum(np.abs(norms_sq - 1.0) > tolerance)

    print(f"Violations (> {tolerance} absolute diff from 1.0):
{violations}")

    if violations > 0:
        idx = np.argmax(np.abs(norms_sq - 1.0))
        print(f"Worst violation at index {idx}: {norms_sq[idx]:.15f}")

if __name__ == "__main__":
    verify_embeddings()
```

File: src/training/classical/visualize_pcam.py

```
import matplotlib.pyplot as plt
from src.data.pcam_loader import get_pcam_dataset
from src.utils.paths import load_paths
from src.utils.seed import set_seed

set_seed(42)

_, PATHS = load_paths()

dataset = get_pcam_dataset(PATHS["dataset"], "test")
```

```
plt.figure(figsize=(10, 5))
for i in range(2):
    img, label = dataset[i]
    plt.subplot(1, 2, i + 1)
    plt.imshow(img.permute(1, 2, 0))
    plt.title("Malignant" if label else "Benign")
    plt.axis("off")

plt.show()
```

File:

src/training/protocol_fixed_regime3b_responsible/test_regime3b_egim
e3b_responsible.py

```
import os
import numpy as np
from sklearn.metrics import accuracy_score

from src.utils.paths import load_paths
from src.utils.load_data import load_data
from src.utils.label_utils import ensure_polar
from src.IQL.learning.class_state import ClassState
from src.IQL.learning.memory_bank import MemoryBank
from src.IQL.backends.exact import ExactBackend
from src.IQL.regimes.regime3b_responsible import Regime3BResponsible

def main():
    print("\n🚀 Testing Regime-3B (Responsible-Set)\n")

    X_train, X_test, y_train, y_test = load_data("polar")

    # Initial polarized memory
    backend = ExactBackend()
    class_states = []

    for cls in [-1, +1]:
        idx = np.where(y_train == cls)[0][0]
        chi = X_train[idx].astype(np.complex128)
        chi /= np.linalg.norm(chi)
        class_states.append(
            ClassState(chi, backend=backend, label=cls)
        )

    memory_bank = MemoryBank(class_states)

    model = Regime3BResponsible(
        memory_bank=memory_bank,
        eta=0.1,
```

```

        alpha_correct=0.0,
        alpha_wrong=1.0,
        tau=0.1,
    )

    train_acc = model.fit(X_train, y_train)

    print("\n=== Training Summary ===")
    print(model.summary())
    print(f"Train Accuracy : {train_acc:.4f}")

    y_pred = model.predict(X_test)
    test_acc = accuracy_score(y_test, y_pred)

    print("\n=== Evaluation ===")
    print(f"Test Accuracy : {test_acc:.4f}")
    print(f"Memory Size : {len(memory_bank.class_states)}")

    print("\n✓ Regime-3B (Responsible-Set) test completed.\n")

if __name__ == "__main__":
    main()

```

File: src/classical/cnn.py

```

import torch
import torch.nn as nn
import torch.nn.functional as F

class PCamCNN(nn.Module):
    """
    Lightweight CNN for PCam feature extraction.
    Produces low-dimensional embeddings suitable for quantum encoding.
    """

    def __init__(self, embedding_dim: int = 32, num_classes: int = 2):
        super().__init__()

        # ----- Convolutional backbone -----
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2), # 48x48

            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),

```

```

        nn.MaxPool2d(2), # 24x24

        nn.Conv2d(64, 128, kernel_size=3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(inplace=True),

        nn.AdaptiveAvgPool2d((1, 1)) # 128 x 1 x 1
    )

    # ----- Embedding head -----
    self.embedding = nn.Linear(128, embedding_dim)

    # ----- Temporary classifier (used ONLY for CNN training) -----
    --
    self.classifier = nn.Linear(embedding_dim, num_classes)

def forward(self, x, return_embedding: bool = False):
    x = self.features(x)
    x = x.view(x.size(0), -1) # flatten

    embedding = self.embedding(x)
    embedding = F.relu(embedding)

    if return_embedding:
        return embedding

    logits = self.classifier(embedding)
    return logits

```

File: src/classical/**init**.py