
ISDO Training Regimes — Implementation Plan & Tracking Summary

Scope

These regimes **build on Direction 1 (ISDO)**, which is **closed**. They **do not modify**:

- the ISDO observable
- the ISDO circuit (B)

They only define **how class states are constructed and updated**.

Shared Foundations (Applies to All Regimes)

Fixed Components (Never Trained)

- **Quantum circuit:** ISDO-B (transition-based interference)
- **Measurement:** ancilla $\langle Z \rangle$
- **Decision rule:**

$$\hat{y}(x) = \text{sign}(\text{Re}\langle \chi | \psi(x) \rangle)$$

Learned Objects

- **Quantum class state(s) $|\chi\rangle$**
- These are the **only trainable parameters**

Data Pipeline

```
image
-> CNN
-> embedding ($\mathbb{R}^{32})
-> L2 normalize
-> $|\psi\rangle
-> ISDO($|\psi\rangle, |\chi\rangle)
```

Regime 1 — Static Prototype Aggregation (Baseline)

Purpose

- Non-iterative baseline

- Fast, stable, interpretable
 - Already implemented and validated
-

Definition

Given labeled embeddings:

- Class +1: $\{|\phi_k^+\rangle\}$
- Class -1: $\{|\phi_k^-\rangle\}$

Construct:

$$|\chi\rangle = \text{normalize} \left(\sum_k |\phi_k^+\rangle - \sum_k |\phi_k^-\rangle \right)$$

Implementation Steps

1. Collect embeddings per class
2. Optionally cluster (KMeans) -> prototypes
3. Sum positive prototypes
4. Subtract negative prototypes
5. Normalize -> class state $|\chi\rangle$

Inference

For each test sample:

$$s = \text{Re}\langle\chi|\psi\rangle \quad \Rightarrow \quad \hat{y} = \text{sign}(s)$$

Properties

Aspect	Status
Training	One-shot
Updates	None
Stability	Very high
Circuit depth	Fixed
QML validity	Yes
Novelty	Moderate

When to Use

- Baseline comparisons
 - Low-data regimes
 - Cold-start initialization for Regime 2 or 3
-

Regime 2 — Online / Incremental ISDO Training (Recommended Core)

Purpose

- Enable **learning over time**
- Adapt class state to data distribution
- Avoid gradients, shots, or parameterized circuits

This is the **quantum perceptron regime**.

Initialization (Cold Start)

Option A — Zero state:

$$|\chi_0\rangle = |0\rangle^{\otimes n}$$

Option B — Bootstrap (recommended):

$$|\chi_0\rangle = \text{normalize} \left(\sum_{i=1}^K y_i |\psi_i\rangle \right)$$

Online Update Rule

For each training sample $(|\psi\rangle, y)$, with $y \in \{+1, -1\}$:

1. Measure:

$$s = \text{Re}\langle\chi_t|\psi\rangle$$

2. If correctly classified:

`y * s >= 0 -> no update`

3. If misclassified:

$$|\chi_{t+1}\rangle = \text{normalize} (|\chi_t\rangle + \eta \cdot y |\psi\rangle)$$

Implementation Notes

- Update is **classical vector arithmetic**
 - Renormalize after every update
 - Learning rate η can be:
 - constant
 - decaying
 - adaptive
-

Properties

Aspect	Status
Training	Online
Updates	Deterministic
Gradients	None
Circuit	Fixed
Shot cost	Minimal
Interpretability	High
QML validity	Strong

When to Use

- Continual learning
 - Streaming data
 - Non-stationary datasets
 - As the **main learning regime**
-

Regime 3 — Multi-State Quantum Memory (Advanced / Nonlinear)

Purpose

- Increase expressivity
- Approximate nonlinear decision boundaries
- Avoid deep circuits or kernels

This is **ensemble learning in Hilbert space**.

Model Structure

Maintain a set of class states:

$$\mathcal{M} = \{|\chi^{(1)}\rangle, \dots, |\chi^{(M)}\rangle\}$$

Each state:

- initialized independently
- trained independently (Regime 2 rule)

Inference Rule

For input $|\psi\rangle$:

$$s_j = \text{Re}\langle\chi^{(j)}|\psi\rangle$$

Decision:

$$\hat{y} = \text{sign}\left(\max_j s_j\right)$$

Training Variants

Variant A — Winner-updates

- Only the best-matching ($|\chi^{(j)}\rangle$) is updated

Variant B — All-updates

- All memory states updated with different learning rates

Variant C — Pruned memory

- Periodically discard low-utility states
-

Properties

Aspect	Status
Expressivity	High
Circuit depth	Fixed
Parameters	Quantum states
Interpretability	Medium
QML novelty	Very high
Complexity	Moderate

When to Use

- Nonlinear class boundaries
 - Multimodal class distributions
 - Few-shot / continual learning
-

Relationship Between Regimes

Regime 1
↓ (initialization)
Regime 2
↓ (parallelization)
Regime 3

- Regime 1 feeds Regime 2
 - Regime 2 generalizes to Regime 3
 - No regime invalidates ISDO
-

Key Constraints (Must Not Violate)

- Do NOT train gate parameters
 - Do NOT introduce variational ansätze
 - Do NOT change the ISDO observable
 - Always renormalize class states
 - Updates must be linear combinations of embeddings
-

Final Status Summary

Regime	Purpose	Status
Regime 1	Static baseline	Implemented
Regime 2	Online learning	Ready to implement
Regime 3	Nonlinear extension	Defined

Recommended Next Action

Implement Regime 2 end-to-end on your existing embeddings and log:

- accuracy vs time

- number of updates
- stability of $|\chi\rangle$

Once Regime 2 is stable, Regime 3 becomes trivial.